

# Cut-Shortcut Pointer Analysis: Re-imagining Context-Sensitivity without Contexts

SHENGYUAN YANG, University of Wisconsin-Madison, USA

WENJIE MA, University of California, Berkeley, USA

THOMAS REPS, University of Wisconsin-Madison, USA

TIAN TAN, Nanjing University, China

YUE LI, Nanjing University, China

Over the past decades, context sensitivity has been considered as one of the most effective ideas for improving the precision of pointer analysis for object-oriented programming languages such as Java. By analyzing each method under distinct contexts, context sensitivity separates the static representations of different dynamic instantiations of variables and heap objects, thereby reducing spurious object flows caused by method calls. However, despite great precision benefits, context sensitivity brings heavy efficiency costs because each method is essentially cloned and analyzed under different contexts. To mitigate this issue, numerous selective context-sensitive approaches have been proposed, applying context sensitivity only to selected methods while analyzing others context-insensitively. However, the selective approaches do not fully remove the efficiency bottleneck because they still rely on the fundamental idea of context sensitivity—analyzing multiple method copies based on calling contexts—which limits their ability to scale to large programs.

In this work, we present a fundamentally different approach, called CUT-SHORTCUT, for fast and precise pointer analysis for Java. The core insight is that the primary effect of cloning methods under different contexts is to filter spurious object flows merged within callee methods; from the vantage point of a typical pointer flow graph (PFG), such effects can be simulated by suppressing the introduction/addition of imprecise flow edges (*Cut*) and adding *Shortcut* edges that directly connect source pointers to the target ones, bypassing method boundaries. We articulate this insight into a general principle that achieves the precision benefits of context sensitivity without explicitly using contexts. This principle is instantiated by identifying and handling three well-characterized program patterns to safely suppress imprecise flows and add precise shortcut flows in the PFG. Our approach is formalized through inference rules, and we formally prove its soundness. We further present a CFL-reachability formulation of CUT-SHORTCUT, both to relate our technique to well-understood formal models of pointer analysis, and to show that CUT-SHORTCUT has the same worst-case asymptotic complexity as context-insensitive analysis. To extend our approach to modern Java programs that introduce new sources of imprecision, we propose CUT-SHORTCUT<sup>S</sup>, an extension of CUT-SHORTCUT that lifts one of its patterns to handle precision loss caused by *Streams* and their interaction with *Lambda Expressions*. This extended pattern distinguishes object flows across different stream pipelines, thereby preserving precision without requiring context sensitivity. Our approaches are implemented on the state-of-the-art pointer analysis framework TAI-E. We evaluate CUT-SHORTCUT on 10 large, complex Java programs used in recent literature, and evaluate CUT-SHORTCUT<sup>S</sup> on three new benchmark suites that we created, containing in total 20 programs with heavy stream usage. The evaluation results are compelling: CUT-SHORTCUT not only achieves precision

---

Some of the material in the paper has previously appeared in <https://dl.acm.org/doi/10.1145/3591242>.

Authors' Contact Information: Shengyuan Yang, University of Wisconsin-Madison, USA, syang686@wisc.edu; Wenjie Ma, University of California, Berkeley, USA, windsey@berkeley.edu; Thomas Reps, University of Wisconsin-Madison, USA, reps@cs.wisc.edu; Tian Tan, Nanjing University, China, tiantan@nju.edu.cn; Yue Li, Nanjing University, China, yueli@nju.edu.cn.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM 1558-4593/2025/11-ART11

<https://doi.org/0000.0000>

comparable to (selective) context-sensitive analysis (where applicable), but also consistently matches or outperforms context-insensitive analysis in efficiency across all benchmarks—running faster in 7 out of 10 cases and equally fast in the remaining 3. Meanwhile, CUT-SHORTCUT<sup>S</sup> generates substantial precision improvement, while preserving scalability and efficiency for programs with heavy stream usage. To the best of our knowledge, this work is the first that has been able to achieve such a good efficiency-precision trade-off for such hard-to-analyze Java programs, including large-scale and feature-rich applications.

CCS Concepts: • Theory of computation → Program analysis.

Additional Key Words and Phrases: Static Analysis, Pointer Analysis, Alias Analysis, Context Sensitivity, Java

## 1 Introduction

Pointer analysis is a family of static-analysis techniques aimed at computing a set of abstract values that characterize what a pointer-valued variable in a program might point to during the execution of the program. The result of pointer analysis lays the groundwork for a wide range of static-analysis applications, such as bug detection [8, 9, 52], security analysis [4, 19], program optimization [71, 88], verification [16, 57], and understanding [43, 73]. Given its importance, over the past decades researchers have invested a substantial amount of effort in trying to build fast and precise algorithms for pointer analysis [67].

Andersen-style pointer analysis [2] has been considered as the standard context-insensitive pointer analysis for Java [67, 72]. We can see it as an iterative algorithm applied to an on-the-fly constructed pointer flow graph (PFG) [39, 41, 83]—where nodes represent program pointers (including variables and instance fields) and edges represent subset constraints between pointers’ points-to sets. During the analysis, abstracted objects are propagated along the edges of the PFG. Because the context-insensitive approach of Andersen’s algorithm does not distinguish between different call sites of the same method, incoming points-to sets are merged in callees; thus, it is extremely fast but with low precision.

To obtain higher precision, one can use a context-sensitive approach to pointer analysis. Basically, for each method call (say  $c$ ) to a callee (say  $m$ ), all elements in  $m$  (like variables, instance fields, and heap objects) will be cloned and analyzed under a context (say  $e$ ) related to certain program elements that distinguish  $c$  from other calls. Accordingly, different forms of context sensitivity have been investigated, such as call-site sensitivity (if  $e$  is a sequence of call sites) [65, 66], object sensitivity (if  $e$  is a sequence of allocation sites) [49, 50], and type sensitivity (if  $e$  is a sequence of types) [69]. Because the object flows merged in the callees are separated under different contexts, the spurious objects flows introduced by method calls can be filtered out, thereby increasing the precision. However, context sensitivity comes with heavy efficiency costs because the analysis has to compute and maintain a huge number of context-sensitive intermediate results. Consequently, context-sensitive algorithms—including 2obj,<sup>1</sup> the most widely used algorithm—do not terminate in hours for several complex programs in the standard DaCapo benchmarks [32, 36, 40].

To be able to handle large and complex programs, in recent years numerous *selective* context-sensitive approaches have been proposed [20, 22, 23, 26, 30, 32, 39–41, 46, 47, 55, 56, 70, 84]. Generally, they rely on a pre-analysis to select a set of methods that are critical to improving precision but do not jeopardize scalability when analyzed using a context-sensitive approach. Then they only apply context sensitivity to those selected methods while analyzing the remaining ones context-insensitively. Although selective context sensitivity significantly improves the scalability of pointer analysis, its efficiency has not yet been fully unlocked because the core methodology of context sensitivity remains unchanged: an algorithm still must replicate methods and analyze their

---

<sup>1</sup>2obj stands for “object sensitivity with a context length of two.”

elements separately under different contexts. When many such replications occur, the analysis can become prohibitively slow and lose scalability [41, 70].

In this work, we present CUT-SHORTCUT, a novel approach to fundamentally change the status quo where pointer analysis for Java primarily relies on context sensitivity to acquire higher precision.

*CUT-SHORTCUT*. The key observation is that, with a context-insensitive approach, imprecision occurs when object flows are merged in a method  $m$  and then the merged flow goes outside  $m$ . In contrast with replicating the method  $m$  into multiple copies to distinguish object flows from different call sites, the approach taken in CUT-SHORTCUT simulates the effect of context sensitivity by suppressing the introduction/addition of imprecise flow edges (*Cut*), and instead adding only more-precise flow edges from exact source pointers to the target ones, bypassing method  $m$  (*Shortcut*). This approach allows us to achieve the benefits of context sensitivity without replicating and analyzing  $m$  under different contexts. More specifically, CUT-SHORTCUT eliminates the need for a context-insensitive pre-analysis phase (as typically required in selective context-sensitive approaches) and performs a single-round, context-insensitive pointer analysis on an on-the-fly constructed pointer flow graph (PFG). The analysis proceeds just like standard context-insensitive pointer analysis, except that certain imprecise edges are suppressed, and precise shortcut edges are added to our PFG. In essence, the core pointer-analysis algorithm remains unchanged—it still propagates points-to information over the PFG until a fixpoint is reached. However, it now operates on a refined graph (compared with the graph constructed by a context-insensitive analysis), where imprecision has been preemptively avoided through targeted graph construction. Here we have two key questions to answer:

- *Which edges to avoid adding?* We suppress the addition of certain edges that bring merged object flows inside a method to somewhere outside (e.g., to the variables of the caller that receive values from the callee’s return variables). In other words, merging object flows within a method does not inherently reduce precision—it is only when these merged flows escape the method that precision loss occurs.
- *Where to add shortcut edges?* We identify the source nodes  $s$  that precede the merging of object flows, and introduce new edges from  $s$  to the corresponding target nodes of the edges whose addition has been suppressed, ensuring both soundness and precision of the analysis.

Although the basic idea seems natural, it is challenging both to identify *which edges to avoid adding* and *where to add shortcut edges*. For the former, we must identify edges that significantly harm precision and assess whether it is possible—and beneficial—to avoid adding them. For the latter, once certain PFG edges are not added, failing to add shortcut edges that fully capture the necessary object-flow sources can compromise soundness. To ensure soundness, rather than trying to tackle all kinds of precision loss, in this work, we articulate a guiding principle, then focus on well-characterized program patterns that align with this principle and contribute substantially to precision loss. Building on this framework, we introduce CUT-SHORTCUT, which systematically improves precision by targeting and handling three general program patterns.

*CUT-SHORTCUT*<sup>S</sup>. While several whole-program pointer-analysis techniques have been proposed for Java in recent years, most remain rooted in the Java 6 language core [22, 23, 28–30, 32, 37, 40, 41, 46, 47, 70, 76–78], leaving many hard-to-analyze features introduced in newer versions of Java unaddressed. In particular, the interaction between *Streams* and *Lambda Expressions* (both introduced in Java 8) often causes substantial precision loss, because object flows become conflated across distinct stream pipelines. Although recent efforts have begun to target modern Java features (e.g., work on module-aware context-sensitive pointer analysis for the Java Platform Module System [38]), existing precision-enhancing techniques—such as context sensitivity and selective

148 context sensitivity—fail to provide analyses that are both precise and scalable when applied to  
 149 programs with heavy use of streams.

150 Building on the core idea of CUT-SHORTCUT —suppressing imprecise flow edges and introducing  
 151 precise shortcut edges—we refine and extend our approach to the stream setting. Concretely, we  
 152 extend one of the existing patterns in CUT-SHORTCUT to also handle stream pipelines, enabling  
 153 accurate tracking of input-to-output relationships within each pipeline and distinguishing object  
 154 flows across pipelines without relying on context sensitivity. However, because real-world stream  
 155 usage is commonly intertwined with lambda expressions, the pattern is further extended to interact  
 156 with an existing static analysis of Java lambda expressions [18], thereby ensuring sound and precise  
 157 handling of functional-style stream operations. This integration enables precise handling of diverse  
 158 behaviors of Java streams. The resulting unified approach, CUT-SHORTCUT<sup>S</sup>, takes a significant  
 159 step forward in realizing highly efficient and precise pointer analysis for modern Java programs.

160 *Contributions.* In summary, the work described in this paper makes the following contributions:

- 161 • CUT-SHORTCUT relies on the insight that one can simulate the effect of a context-sensitive tech-  
 162 nique by modifying what edges are added to the PFG during a context-insensitive Andersen-  
 163 style pointer analysis. We instantiate this principle in CUT-SHORTCUT by exploiting three  
 164 well-characterized program patterns, designing algorithms to identify and handle them on-  
 165 the-fly with pointer analysis while constructing the PFG (Section 3).
- 166 • We formalize CUT-SHORTCUT as a set of inference rules (Section 4) and formally prove its  
 167 soundness (Section 5).
- 168 • We provide a CFL-reachability formulation to express our approach as a formal-language  
 169 reachability problem (Section 6). This formulation not only connects our approach to the  
 170 well-established formal-language-reachability-based formulations of pointer analysis, but also  
 171 offers a perspective on the efficiency advantages of CUT-SHORTCUT over context-sensitive  
 172 techniques.
- 173 • We present CUT-SHORTCUT<sup>S</sup>, an enhanced variant of CUT-SHORTCUT, which extends one of  
 174 the patterns used in CUT-SHORTCUT to accommodate modern Java features, specifically the  
 175 handling of *Streams* and their interaction with *Lambda Expressions* (Section 7).
- 176 • We implement CUT-SHORTCUT and CUT-SHORTCUT<sup>S</sup> on the state-of-the-art Java pointer-  
 177 analysis framework TAI-E, and evaluate their effectiveness (Section 8).
  - 178 – CUT-SHORTCUT exhibits an extremely promising efficiency and precision trade-off: it is  
 179 even faster than context-insensitive analysis on 7 out of 10 benchmarks, and matches  
 180 its efficiency on the remaining 3, while achieving precision comparable to (selective)  
 181 context-sensitive analyses—when those analyses can scale.
  - 182 – CUT-SHORTCUT<sup>S</sup> further improves precision for programs with heavy stream usage, as  
 183 shown on our newly proposed benchmark suites. By distinguishing object flows through  
 184 different stream pipelines, CUT-SHORTCUT<sup>S</sup> achieves a level of precision unattainable by  
 185 any existing (selective) context-sensitivity approach, while maintaining efficiency that is  
 186 consistently faster than, or on par with, context insensitivity.

188 This article refines and substantially extends our earlier work [48], presented at the 44th ACM  
 189 SIGPLAN Conference on Programming Language Design and Implementation (PLDI 2023). Whereas  
 190 the conference version comprised 26 pages, this article spans 73 pages, with the following major  
 191 extensions:

- 192 • Sections 2–5 reflect the core content of the conference version but have been carefully  
 193 rephrased and reorganized throughout to improve clarity and precision, offering a more com-  
 194 prehensive and updated presentation. In addition, Section 3.1 introduces a new comparison  
 195 with method-summary-based pointer analysis, highlighting key differences and advantages.

197 Sections 4.2–4.4 provide detailed explanations of the mechanisms for each pattern, which  
 198 were not fully presented in [48] due to space constraints, along with several refinements.  
 199 Finally, the soundness theorem in Section 5 is now supported by a complete and rigorous  
 200 proof, whereas it was only sketched in [48].

- 201 • Section 6 introduces a CFL-reachability formulation of CUT-SHORTCUT, illustrated with de-  
 202 tailed examples. This formulation demonstrates that while significantly improving precision,  
 203 CUT-SHORTCUT preserves the asymptotic complexity of context-insensitive analysis. It not  
 204 only connects our work to well-established language-reachability-based formal frameworks,  
 205 but also provides a principled explanation for why CUT-SHORTCUT achieves high efficiency  
 206 compared to context-sensitivity-based techniques.
- 207 • Section 7 introduces a new extension, CUT-SHORTCUT<sup>S</sup>, which abstracts stream pipelines  
 208 statically and tracks their input elements and output locations on the fly, enabling the sup-  
 209 pression of imprecise flow edges and the addition of precise ones within a single round of  
 210 pointer analysis. To realize this extension, we systematically classify stream operations into  
 211 categories and design dedicated handling techniques for each, including an interaction with  
 212 an existing analysis of Java lambda expressions [18], to mitigate precision loss caused by  
 213 functional-style stream operations.
- 214 • Section 8 presents updated experimental results for CUT-SHORTCUT and new evaluations  
 215 for CUT-SHORTCUT<sup>S</sup>. For CUT-SHORTCUT, we report (i) results based on additional precision  
 216 metrics, (ii) an additional comparison with ZIPPER<sup>e</sup>-guided 2-type sensitive analyses [41],  
 217 and (iii) the results of new experiments to isolate the contribution of each of the patterns  
 218 used in CUT-SHORTCUT. For CUT-SHORTCUT<sup>S</sup>, we evaluate its effectiveness on three new  
 219 benchmark suites that we created, which consist of Java programs with heavy stream usage,  
 220 designed to balance operation coverage, representativeness, and realism: (i) 10 synthetic  
 221 micro-benchmarks systematically covering diverse stream operations and their combinations  
 222 on small inputs; (ii) 5 enterprise-backend-inspired benchmarks that mimic common web-  
 223 backend data-processing and transmission workflows using streams, with an in-memory mock  
 224 database; and (iii) 5 real-world programs with intensive stream usage beyond the enterprise  
 225 application domain.

226 We provide an open-source implementation of our refined and extended version of CUT-SHORTCUT  
 227 (including CUT-SHORTCUT<sup>S</sup>) at [86]. All experimental results reported in this article are fully re-  
 228 producible using this artifact.

## 230 2 Motivation

231 We use the example in Figure 1 to illustrate the basic idea of CUT-SHORTCUT and show how it  
 232 differs from context insensitivity and context sensitivity for pointer analysis. Class Carton has a  
 233 field item and provides setItem() and getItem() to modify and retrieve the item value, respectively.  
 234 In main(), two Item objects ( $o_{16}, o_{21}$ ) are stored in the field item of two Carton objects ( $o_{15}, o_{20}$ ) via  
 235 method setItem() and retrieved later via getItem(). (We use  $o_i$  to denote the abstract heap object  
 236 allocated at line  $i$ .) After execution, result1 (result2) will only point to  $o_{16}$  ( $o_{21}$ ). Below we explain  
 237 how context insensitivity, context sensitivity, and our approach work for this example.

238 *Context Insensitivity.* Figure 1(a) shows the PFG fragment of the example. (An edge in the PFG  
 239 that connects  $p$  to  $q$  indicates that what is pointed to by  $p$  should also be pointed to by  $q$ . We use  
 240  $pt(p)$  to denote the points-to set that the algorithm computes for  $p$ .) Because a context-insensitive  
 241 approach does not distinguish the two calls to setItem(),  $pt(item1)$  and  $pt(item2)$  will be merged  
 242 into  $pt(item)$  (where item is the parameter at line 3), so that the store statement at line 4 causes  
 243  $o_{15}.item$  and  $o_{20}.item$  to have imprecise points-to sets. Moreover, the two calls of getItem() are

```

246 1 class Carton {           13 // usage code
247 2   Item item;           14 void main() {
248 3   void setItem(Item item){15 Carton c1 = new Carton();//o15
249 4     this.item = item;    16 Item item1 = new Item();//o16
250 5   }                     17 c1.setItem(item1);
251 6   Item getItem() {      18 Item result1 = c1.getItem(); o15.item
252 7     Item r = this.item; 19 Carton c2 = new Carton();//o20
253 8     return r;          20 Item item2 = new Item();//o21
254 9   }                     21 c2.setItem(item2);
255 10 }                     22 c2.setItem(item2);
256 11 class Item {}         23 Item result2 = c2.getItem();
257 12
258 24 }

```

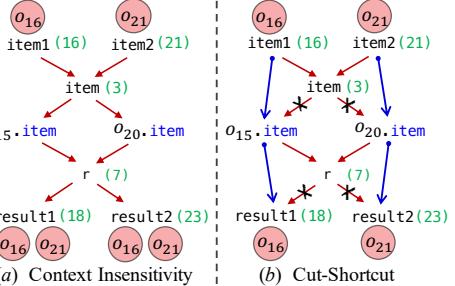


Fig. 1. An example to illustrate the basic idea of CUT-SHORTCUT.

handled in the same manner, so that the load statement at line 7 causes the merged objects ( $\{o_{16}, o_{21}\}$ ) to flow to  $r$ , thereby causing the points-to sets of  $result1$  and  $result2$  to be imprecise.

*Context Sensitivity.* The goal of context sensitivity is to separate the flows merged at  $item$  (line 3) and  $r$  (line 7), and let  $o_{16}$  only flow to  $result1$  and  $o_{21}$  only flow to  $result2$ . To distinguish the object flows from different call sites (namely, lines 17 and 22 for  $setItem()$ , and lines 18 and 23 for  $getItem()$ ), every program element in  $setItem()$  (i.e.,  $item$  and  $this.item$ ) and  $getItem()$  (i.e.,  $this.item$  and  $r$ ) should be qualified with a context and analyzed separately, which is equivalent to analyzing each of those methods twice under two different contexts. We can anticipate that the cost of computing and maintaining context-sensitive information would be very heavy in the real world, especially when the program is large or complex. Selective context sensitivity alleviates this problem, but does not fundamentally change it, because even a small set of selected methods in a complex program can threaten scalability and blow up the analysis when they are analyzed context-sensitively [41, 70].

*CUT-SHORTCUT.* Figure 1(b) depicts how CUT-SHORTCUT works, where the arrows with crosses indicate edges whose addition to the PFG are suppressed, and the blue arrows indicate the added shortcut edges. With a context-insensitive approach, the imprecision in  $pt(o_{15}.item)$  and  $pt(o_{20}.item)$  arises from edges  $item \rightarrow o_{15}.item$  and  $item \rightarrow o_{20}.item$  (because of the store statement at line 4), so the addition of these two edges are suppressed. Instead, we add shortcut edges from the precise source (i.e.,  $item1$  or  $item2$ ) directly to the target nodes. Likewise, for nodes  $result1$  and  $result2$ , the introduction of edges from  $r$  are suppressed, and shortcut edges  $o_{15}.item \rightarrow result1$  and  $o_{20}.item \rightarrow result2$  are added. In our new PFG, by applying the same algorithm as context insensitivity (propagating the points-to results along the PFG edges), we derive points-to results as precise as context sensitivity for this example.

### 3 The Cut-Shortcut Approach, Informally

We introduce the overall principle lying behind our approach (Section 3.1), on which the three program patterns, the field-access pattern (Section 3.2), the container-access pattern (Section 3.3) and the local-flow pattern (Section 3.4) are based.

#### 3.1 Overview

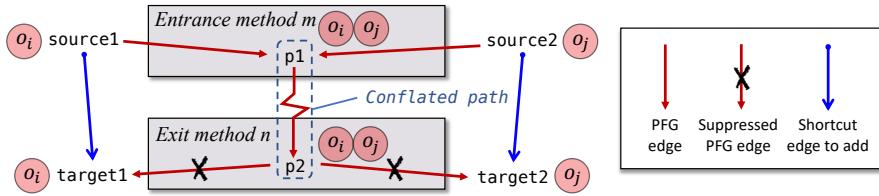
CUT-SHORTCUT addresses the precision loss inherent in context-insensitive pointer analysis by adhering to a general principle: suppress the addition of imprecise flow edges that would carry merged object flows out, and instead add precise shortcut edges as their substitutes in the pointer flow graph (PFG). To elucidate this principle, we first introduce some key terminology.

295     *Definition 3.1 (Local/Non-local pointers).* A pointer  $p$  is *local* to a method  $m$  if  $p$  is a variable  
 296     declared in  $m$ . Otherwise,  $p$  is *non-local* to  $m$  (i.e.,  $p$  is declared outside  $m$  or is an instance field.)

297     *Definition 3.2 (Conflated path, ENTRANCE and EXIT).* A path  $p$  in PFG is a *conflated path* if: (1) it  
 298     begins at a confluence point, where two (or more) pointers *non-local* to a method  $m$  merge into a  
 299     pointer  $s$  that is *local* to  $m$ , and (2) it ends at a divergence point, where a pointer  $e$  *local* to a method  
 300      $n$  leads to two (or more) pointers *non-local* to  $n$ . We say that  $s$  is the *start node* of  $p$ , and that  $e$  is  
 301     the *end node* of  $p$ . We also say that  $m$  is  $p$ 's ENTRANCE, and that  $n$  is  $p$ 's EXIT.

302     For example, the path `item` →  $o_{15}$ .`item` → `r` in Figure 1(a) is a conflated path. Here, `item` and `r`  
 303     are the start and end nodes, respectively, with `setItem()` and `getItem()` serving as its ENTRANCE  
 304     and EXIT, respectively. Similarly, the single-node paths `item` (line 3) and `r` (line 7) are also conflated  
 305     paths.

306     *Definition 3.3 (TARGET and SOURCE pointers).* Given a conflated path  $p$ , its TARGET pointer (TARGET  
 307     for short) is any successor of  $p$ 's end node in PFG. A TARGET receives the merged object flow leaving  
 308     EXIT. A SOURCE pointer (SOURCE for short) of a given TARGET  $t$  is any predecessor of the start node  
 309     of  $p$  whose points-to set carries (part of) the objects that flow to  $t$  (i.e., contributes to the points-set  
 310     of  $t$ ) during dynamic execution. A TARGET may have multiple SOURCES, and multiple TARGETS of  
 311     the same conflated path may share one or more SOURCES. In context-insensitive analysis, object  
 312     flows from distinct SOURCES of different TARGETS are merged inside ENTRANCE and propagated  
 313     together to all TARGETS.



316  
 317     Fig. 2. Illustration of the insight that is the basis of CUT-SHORTCUT, about which edges to suppress and where  
 318     319     320     321     322     323  
 324     325     326  
 327     328     329     330     331     332     333     334     335     336     337     338     339     340     341     342     343  
 324     Fig. 2. Illustration of the insight that is the basis of CUT-SHORTCUT, about which edges to suppress and where  
 325     326  
 327     328     329     330     331     332     333     334     335     336     337     338     339     340     341     342     343  
 327     Now we explain the insight that lies behind our principle using Figure 2.  $p_1$ ,  $p_2$ ,  $\text{target1}$ ,  $\text{target2}$ ,  
 328      $\text{source1}$ , and  $\text{source2}$  are PFG nodes that represent program pointers (i.e., variables or instance  
 329     fields), and  $o_i$  and  $o_j$  represent two abstract heap objects. Suppose that  $\text{source1}$  and  $\text{source2}$  are  
 330     *non-local* to method  $m$ , and  $\text{target1}$  and  $\text{target2}$  are *non-local* to method  $n$ . If there exists a PFG  
 331     path from  $p_1$  to  $p_2$ , this path qualifies as a conflated path by Definition 3.2. At runtime, suppose  
 332     that  $\text{target1}$  only points to  $o_i$ , which originates from  $\text{source1}$ , and that  $\text{target2}$  only points to  $o_j$ , which  
 333     originates from  $\text{source2}$ . With a context-insensitive approach, however, these object  
 334     flows are conflated, leading both  $\text{target1}$  and  $\text{target2}$  to point to both  $o_i$  and  $o_j$ . In many cases  
 335     of real programs, precise points-to results can be easily identified by human intuition with the  
 336     comprehension of semantic information that are ignored by traditional pointer analysis. CUT-  
 337     SHORTCUT emulates such a scheme: it identifies certain conflated paths where SOURCES for each  
 338     TARGET can be soundly identified (by leveraging certain semantic clues to better capture the dynamic  
 339     behavior of the program) so that we can effectively improve precision by suppressing the edges  
 340     connecting the end node of the conflated path to the TARGETS (e.g.,  $p_2 \rightarrow \text{target1}/\text{target2}$ ) and  
 341     adding shortcut edges to connect corresponding SOURCES for each TARGET, like  $\text{source1} \rightarrow \text{target1}$   
 342     and  $\text{source2} \rightarrow \text{target2}$ . Note that Figure 2 shows a simple case with only two TARGETS, each

344 having a single SOURCE. In real programs, a TARGET may have multiple SOURCES, all of which must  
345 be connected via shortcut edges to maintain soundness.

346 To enhance precision, we aim to discover as many conflated paths as possible and determine  
347 minimal SOURCE sets that fully cover each TARGET. However, generalizing this process is difficult  
348 due to the semantic diversity of real-world code [39, 41] and the challenge of ensuring soundness  
349 when the addition of some edges are suppressed in the PFG. To increase precision, while maintaining  
350 soundness, CUT-SHORTCUT defines three concrete program patterns prevalent in practice—field-  
351 access pattern, container-access pattern, and local-flow pattern—which instantiate our general  
352 principle and account for a large portion of precision loss. The following subsections detail how we  
353 identify conflated paths, their ENTRANCES and EXITS, and the corresponding TARGETS and SOURCES  
354 for each pattern, allowing us to safely suppress the addition of imprecise edges and introduce  
355 precise shortcut edges.

356 Finally, we emphasize that, although the PFG built by our approach is different from the original  
357 context-insensitive PFG, CUT-SHORTCUT does not physically modify an existing graph during the  
358 analysis. Rather, it builds the graph on-the-fly with guidance from pattern-specific rules. Thus, our  
359 pointer analysis remains a single-phase, monotonic process without any graph rewriting—imprecise  
360 edges are never added, and precise edges are introduced as needed during construction.

361 *Comparison with Selective Context Sensitivity.* CUT-SHORTCUT is fundamentally different from  
362 selective context sensitivity approaches [28, 30, 39–41, 70, 76], which select methods by different  
363 heuristics and apply contexts directly to them (equivalent to cloning the methods and analyzing  
364 them). The main cost of these analyses is from cloning the methods and analyzing them multiple  
365 times. In contrast, to avoid the overhead of cloning, CUT-SHORTCUT neither selects methods  
366 nor analyzes them context-sensitively; instead, it simulates the effect of context sensitivity by  
367 suppressing imprecise edges and adding precise edges in the PFG. (The challenge is to reduce  
368 imprecision while not compromising soundness.) No contexts are applied to any methods in  
369 CUT-SHORTCUT.

370 *Comparison with Method-Summary-Based Pointer Analysis.* Method summarization is a widely  
371 used technique in inter-procedural program analysis, and has been applied to pointer analysis [11,  
372 15, 54, 85]. A summary encodes how a method manipulates pointers—what objects its parameters  
373 may point to, how it updates pointer relationships internally, and what pointers it returns or makes  
374 accessible to callers. In terms of how these issues were handled in the motivating example in  
375 Section 2, CUT-SHORTCUT may seem to resemble method summarization by summarizing pointer  
376 flow at caller’s invocation site. However, as we detail in later sections, CUT-SHORTCUT differs from  
377 summary-based pointer analysis in several key aspects:

- 378 • *No Context Overhead:* Summary-based pointer analysis approaches typically rely on context  
379 sensitivity to improve precision, often reanalyzing or reusing specialized summaries for each  
380 context [15, 85]. CUT-SHORTCUT, however, never duplicates methods during analysis, and  
381 hence achieves high efficiency.
- 382 • *Granularity:* Summary-based approaches encapsulate the entire method as a single unit,  
383 whereas CUT-SHORTCUT does not generate whole-method summaries. Instead, it operates at a  
384 finer granularity by selectively rerouting imprecise object flows for a method (or a series of  
385 methods, depending on the specific pattern) while leaving unaffected parts intact.
- 386 • *Top-Down Manner:* Summary-based approaches compute method summaries upfront and  
387 instantiate them as constraints when the summarized method is called. In contrast, CUT-  
388 SHORTCUT applies standard top-down context-insensitive pointer analysis on a refined PFG;  
389 the analysis still proceeds from callers to callees in a top-down way (with some callees’  
390 out-flow being ignored).

392

- *Pattern-Based Selectivity*: Summary-based approaches treat all methods uniformly, generating summaries for every procedure regardless of necessity. CUT-SHORTCUT, however, selectively identifies specific program patterns that commonly cause precision loss, making it more lightweight and focused.
- *Cross-Invocation Flow Rerouting*: The principle of CUT-SHORTCUT enables re-routing of object flows across multiple call sites. This capability is exemplified in the container-access pattern, where the input source of a container instance is directly connected to its retrieval sites by shortcut edges. Such cross-invocation flow rerouting is difficult to achieve using traditional method-summarization techniques.

### 3.2 Field-Access Pattern

There are two kinds of field accesses in Java, field store and field load (in this pattern, we focus on instance fields). Java programs often wrap field accesses in methods (e.g., setter/getter), so the client code needs to access the fields by calling these methods. However, in a context-insensitive analysis, object flows involved in the accesses to fields of different objects will be merged inside these methods, leading to imprecision. In this pattern, we tackle such precision loss via the overarching principle described in Section 3.1. To aid understanding, we first introduce the basic cases for handling field stores (Section 3.2.1) and loads (Section 3.2.2) using our motivating example, provided in Figure 1. We then describe a more general case in Section 3.2.3 to show how our semantics-based scheme tracks the value flows from the caller (and the caller of the caller ...), where the values are finally stored in the target fields through a series of invocation parameters.

3.2.1 *Handling of Store*. We address the imprecision in the points-to results of instance fields. For a field-store statement  $s : x.f = y$ , if objects pointed to by both  $x$  and  $y$  come from the parameters (including the `this` variable) of the method  $m$  that contains  $s$ , then incoming object flows from arguments of different calls (of  $m$ ) are merged inside  $m$ , which may cause imprecision in  $pt(o_i.f)$  where  $o_i \in pt(x)$ . Line 4 in Figure 1 gives such an example: `this` and `item` are parameters of method `setItem()`. When analyzing `setItem()` context-insensitively, objects from call sites at lines 17 and 22 are merged in it, leading to an imprecise result, i.e.,  $pt(o_{15}.item) = pt(o_{20}.item) = \{o_{16}, o_{21}\}$ .

In CUT-SHORTCUT, our idea is to suppress the addition of certain store edges that propagate the merged object flows to instance fields (i.e., TARGETS), and find SOURCES at the call-sites that call the method containing those field stores. We search for store statements of the form  $s : x.f = y$ , in which both  $x$  and  $y$  are parameters of method  $m$  (that contains  $s$ ) and not redefined in  $m$ . (This condition guarantees that the values of both  $x$  and  $y$  all come from arguments of  $m$ 's call-sites.) If such an  $s$  is found, then multiple calls of  $m$  result in a conflated path (with only one node  $y$ , e.g., `item` in Figure 1(a)), with  $m$  being both ENTRANCE and EXIT. Accordingly, TARGETS are the instance fields being accessed by  $x.f$  and stored at  $s$  (e.g.,  $o_{15}.item$  and  $o_{20}.item$  stored at line 4). To prevent merged object flows from being propagated to the TARGETS, we suppress the addition of store edges to them in the PFG. For example, the addition of  $item \rightarrow o_{15}.item$  and  $item \rightarrow o_{20}.item$  are suppressed in Figure 1.

Because the values of both  $x$  and  $y$  come from arguments of  $m$ 's call-sites, we can match the SOURCE for TARGET ( $s$ ) at each call-site—namely, the SOURCE is the argument passed to  $y$  (e.g., the argument is `item1` at line 17 and  $y$  is the parameter `item` at line 3) and the TARGET is  $o_i.f$ , where  $o_i \in pt(a)$  and  $a$  is the argument passed to  $x$  (e.g., the argument  $a$  is `c1` at line 17 and  $x$  is the parameter `this`). For example, at line 17, `item1` is the SOURCE for TARGET  $o_{15}.item$  (as  $pt(c1) = \{o_{15}\}$ ), and likewise, `item2` is the SOURCE for TARGET  $o_{20}.item$  at line 22 (because  $pt(c2) = \{o_{20}\}$ ). Thus, we add shortcut edges  $item1 \rightarrow o_{15}.item$  and  $item2 \rightarrow o_{20}.item$ , which in this example achieves a precise result, i.e.,  $pt(o_{15}.item) = \{o_{16}\}$  and  $pt(o_{20}.item) = \{o_{21}\}$  (depicted in the top half of Figure 1(b)).

442    3.2.2 *Handling of Load*. We address the imprecision in the points-to results for left-hand side (LHS)  
 443    variables whose values are returned from the methods that load instance fields. For a field-load  
 444    statement of the form  $s : x = y.f$ , if objects pointed to by  $y$  come from a parameter (including the  
 445    this variable) of the method  $m$  that contains  $s$ , and  $x$  is the return variable of  $m$ , then the objects  
 446    loaded from  $y.f$  (where  $y$  points to the incoming objects from arguments of different calls of  $m$ ) are  
 447    merged inside  $m$ , which may cause imprecision in points-to sets for the LHS variables of  $m$ 's call  
 448    sites. Line 7 in Figure 1 gives such an example: this is a parameter of method `getItem()` and  $r$  is  
 449    its return variable. When analyzing `getItem()` context-insensitively, objects from the call sites at  
 450    lines 18 and 23 (i.e.,  $o_{15}$  and  $o_{20}$ ) are merged in it, and objects loaded from  $o_{15}$  and  $o_{20}$  (i.e.,  $o_{16}$  and  
 451     $o_{21}$ ) are also merged at  $r$  and propagated to `result1` and `result2`, leading to the imprecise result  
 452     $pt(result1) = pt(result2) = \{o_{16}, o_{21}\}$ .

453    In CUT-SHORTCUT, our idea is to suppress the addition of these return edges from  $m$  (that return  
 454    values loaded from fields), which propagate merged object flows to the LHS variables (i.e., TARGETS),  
 455    and find SOURCES at  $m$ 's call sites. We search for load statements of the form  $s : x = y.f$  in which (i)  
 456    the values of both  $x$  and  $y$  are directly related to the variables of the call sites of  $m$  (that contains  
 457     $s$ ), i.e.,  $y$  is a parameter of  $m$  and not redefined in  $m$  (this condition guarantees that the value of  
 458     $y$  all come from arguments of  $m$ 's call sites), and (ii)  $x$  is the return variable of  $m$ . If such  $s$  and  
 459     $m$  are found, then multiple calls of  $m$  result in a conflated path (with only one node  $x$ , e.g.,  $r$  in  
 460    Figure 1(a)), with  $m$  being both ENTRANCE and EXIT. Hence, the addition of return edges from  $m$  to  
 461    the LHS variable (i.e., the TARGET) of each call-site that calls  $m$  are suppressed to avoid propagation  
 462    of merged flows, and the SOURCE at each call-site is identified, i.e.,  $o_i.f$  where  $o_i \in pt(a)$  and  $a$  is  
 463    the argument passed to  $y$  (e.g.,  $a$  is  $c1$  at line 18 and  $y$  is `this`). For example, for field load at line  
 464    7 in Figure 1, the edges  $r \rightarrow result1$  and  $r \rightarrow result2$  are suppressed. For `result1`, we find its  
 465    SOURCE,  $o_{15}.item$ , at line 18 (because  $pt(c1) = \{o_{15}\}$ ), and add shortcut edge  $o_{15}.item \rightarrow result1$ .  
 466    Similarly, we add  $o_{20}.item \rightarrow result2$  at line 23. By this means, in this example we obtain a precise  
 467    result:  $pt(result1) = \{o_{16}\}$  and  $pt(result2) = \{o_{21}\}$  (see the bottom half of Figure 1(b)).

468    3.2.3 *Handling of Nested Calls for Field Accesses*. In real-world programs, field accesses become  
 469    complicated when nested method calls are involved. For instance, when a constructor is called, it  
 470    may call another constructor or setter method where the store statement is finally executed. In  
 471    addition, when inheritance is used, it becomes common for field stores to involve nested calls. For  
 472    example, in Figure 3, `A.set()` is a setter called by `A`'s constructor (line 3). There are two nested  
 473    calls to it,  $8 \rightarrow 3 \rightarrow A.set()$  and  $10 \rightarrow 3 \rightarrow A.set()$  (a call site is represented by its line number).

---

```

476 1 class A {           6 // usage code
477 2   T f;             7 T t1 = new T(); //o7
478 3   A(T t){this.set(t);} 8 A a1 = new A(t1); //o8
479 4   set(T p){this.f=p;} 9 T t2 = new T(); //o9
480 5 }                   10 A a2 = new A(t2); //o10

```

---

482    Fig. 3. An example of nested calls for field store.

483    manner to obtain better precision. Specifically, for a method that contains a field store, we analyze  
 484    its nested calls (from the method itself to its caller, caller's caller, and so on, if possible) to add  
 485    shortcut edges at proper call-sites. For the field store in `A.set()`, we trace back to the call-sites  
 486    at lines 8 and 10, add shortcut edges  $t1 \rightarrow o_8.f$  and  $t2 \rightarrow o_{10}.f$ , and derive the precise result  
 487     $pt(o_8.f) = \{o_7\}$  and  $pt(o_{10}.f) = \{o_9\}$ . For the details of our approach to handling nested calls for  
 488    field stores, see Section 4.2, where we formally define this pattern using a set of inference rules.

489    Here, if we add shortcut edges,  $t \rightarrow o_8.f$  and  $t \rightarrow o_{10}.f$ , at the call-site of `A.set()`  
 490    (line 3), we still lose precision because both  
 491    `this` and  $t$  are parameters of `A()`, and ob-  
 492    ject flows from call sites at lines 8 and  
 493    10 are merged here, leading to the impre-  
 494    cise result  $pt(o_8.f) = pt(o_{10}.f) = \{o_7, o_9\}$ .

495    To handle nested calls for field stores, we

496    extend our approach in a more general

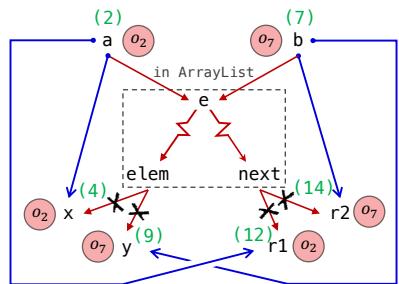
### 491 3.3 Container-Access Pattern

492 Containers (e.g., `ArrayList`) are pervasively used in Java programs. Because numerous objects flow  
 493 in and out of containers, how they are handled has a crucial effect on the precision of a pointer-  
 494 analysis algorithm. With a context-insensitive approach, objects stored in different containers  
 495 may flow to—and get merged—in the same container methods, and thus there may be horrendous  
 496 precision loss when the merged objects flow out.

497 To separate object flows in different containers, conventional approaches improve precision by  
 498 analyzing container methods context-sensitively. However, as pointed out by recent work [3, 14],  
 499 applying context sensitivity to container methods will bring heavy efficiency costs, especially  
 500 for large programs. As a result, to analyze containers precisely while reducing analysis overhead,  
 501 researchers [3, 14, 21] have proposed various techniques. He et al. [21] perform a pre-analysis  
 502 to identify context-dependent objects based on container-usage patterns, and then use a context-  
 503 debloating technique to speed up object sensitivity by forming contexts with context-dependent  
 504 objects only. Fegade and Wimmer [14] and Antoniadis et al. [3] propose to manually rewrite  
 505 container implementations for composing code models that are simpler than the original imple-  
 506 ments, but still preserve the side effects of containers needed for pointer analysis; then, the  
 507 rewritten container code is analyzed using a context-sensitive approach. In CUT-SHORTCUT, we also  
 508 need to handle containers for obtaining good precision. However, we do it differently by adopting  
 509 a scheme that uses the principle from Section 3.1, without relying on context sensitivity at all.  
 510

511 3.3.1 *Handling of Containers.* Figure 4 shows an `ArrayList` example of a container. (For now, just  
 512 focus on lines 1-9.) We create two `ArrayList`s,  $o_1$  (line 1) and  $o_6$  (line 6), and add two objects  $o_2$   
 513 (line 3) and  $o_7$  (line 8) to them, respectively. At lines 4 and 9, we retrieve the elements from  $o_1$  and  
 514  $o_6$  and assign them to  $x$  and  $y$  separately. Clearly,  $x$  ( $y$ ) only points to  $o_2$  ( $o_7$ ) at run time.

```
516 1 ArrayList l1 = new ArrayList(); // o1, ptH(l1) = {h1}
 2 Object a = new Object(); // o2
 3 l1.add(a); // call to ENTRANCE, a is SOURCE
 4 Object x = l1.get(0); // call to EXIT, x is TARGET
 5
 6 ArrayList l2 = new ArrayList(); // o6, ptH(l2) = {h2}
 7 Object b = new Object(); // o7
 8 l2.add(b); // call to ENTRANCE, b is SOURCE
 9 Object y = l2.get(0); // call to EXIT, y is TARGET
10
11 Iterator it1 = l1.iterator(); // ptH(it1) = {h1}
12 Object r1 = it1.next(); // call to EXIT, r1 is TARGET
13 Iterator it2 = l2.iterator(); // ptH(it2) = {h2}
14 Object r2 = it2.next(); // call to EXIT, r2 is TARGET
```



527 Fig. 4. An example of the container-access pattern.

528 Container implementations typically offer APIs for adding elements into, and retrieving elements  
 529 from, containers. When a program creates multiple instances of the same container class (e.g.,  
 530 two `ArrayList` container instances in Figure 4) and uses these APIs separately on each instance,  
 531 context-insensitive pointer analysis fails to distinguish the flows associated with different containers.  
 532 Specifically, the input objects added to each container instance are merged and propagated to all  
 533 retrieval sites, leading to significant imprecision. Consider Figure 4: with a context-insensitive  
 534 analysis, both  $o_2$  and  $o_7$  are merged into the internal variable  $e$  (the parameter of `add()`), flow  
 535 through a series of edges (representing the internal working of `ArrayList`, shown in the PFG in  
 536 Figure 4) to the return value  $\text{elem}$  of `get()`, and then further flow to both  $x$  and  $y$ . As a result, both  
 537  $pt(x)$  and  $pt(y)$  contain  $o_2$  and  $o_7$ , which is imprecise.

Following the over-arching principle of CUT-SHORTCUT, we identify these in-container conflated paths for each container class, which start at the parameters of addition APIs and end at the return values of retrieval APIs. These internal conflated paths (e.g.,  $e \rightsquigarrow \text{elem}$  in Figure 4) have their ENTRANCE being the addition method and EXIT being the retrieval method. Let  $\text{ENTRANCE}_{cs}$  and  $\text{EXIT}_{cs}$  denote call sites to the addition and retrieval APIs, respectively. At an  $\text{ENTRANCE}_{cs}$ , the relevant input argument (say  $\text{ENTRANCE}_{cs}^{arg}$ , e.g.,  $a$  at line 3) is treated as a SOURCE, while at an  $\text{EXIT}_{cs}$ , the LHS variable (say  $\text{Exit}_{cs}^{lhs}$ , e.g.,  $x$  at line 4) is treated as a TARGET.

The key challenge that remains is matching SOURCES and TARGETS —i.e., determining which pairs correspond to the same container instance. To handle this task uniformly—not only for simple cases like the example above, but also for more complex scenarios such as element retrieval via iterators or via Map’s `keySet()` and `values()`—we introduce a static abstraction for container instances, which we call *hosts*, and a mapping called the *pointer-host map*, denoted as  $pt_H$ . In our analysis, each dynamic container instance is abstracted statically by a host  $h$  generated at its allocation site, analogous to a standard allocation-site-based heap abstraction, but specialized for containers. The map  $pt_H$  then associates pointers with the set of hosts they may relate to. For example, because the `ArrayList` object  $o_1$  is assigned to 11, we have  $pt_H(11) = \{h_1\}$ , where  $h_1$  denotes the host abstracting the first `ArrayList` instance. Similarly, we have  $pt_H(12) = \{h_2\}$ . To match SOURCES and TARGETS, we compare the  $pt_H$  sets of the receiver variables at each pair of addition-method and retrieval-method call-sites (denoted by  $\text{ENTRANCE}_{cs}^{rv}$  and  $\text{EXIT}_{cs}^{rv}$ ). If these sets overlap—i.e., they can share a common host—the corresponding  $\text{ENTRANCE}_{cs}^{arg}$  and  $\text{Exit}_{cs}^{lhs}$  are considered a match, and we add a shortcut edge from the  $\text{ENTRANCE}_{cs}^{arg}$  (the SOURCE) to the  $\text{Exit}_{cs}^{lhs}$  (the TARGET). In our above example, both `add()` and `get()` are called on 11, whose  $pt_H$  includes  $h_1$ . Therefore, we add the shortcut edge  $a \rightarrow x$ . Similarly, we add  $b \rightarrow y$  for the second host. More straightforwardly, this process can be viewed from the perspective of each host  $h$  maintaining a *source set* and a *target set*, consisting of the SOURCES and TARGETS associated with  $h$  via ENTRANCE and EXIT call-sites. We then connect each SOURCE in  $h$ ’s *source set* to each TARGET in its *target set* via shortcut edges. As a result, for this example the analysis obtains precise points-to for the variables  $x$  and  $y$ , namely,  $pt(x) = \{o_2\}$  and  $pt(y) = \{o_7\}$ . In this context,  $pt_H$  serves a role analogous to a traditional points-to relation. We will revisit its other uses shortly.

**3.3.2 Handling of Container-Dependent Objects.** Lines 11-14 in Figure 4 show a common usage of containers, i.e., accessing elements via iterators. At static time, each iterator depends on an abstract container instance (i.e., a host), and the elements retrieved from the iterator are the ones stored in that container instance (i.e., the SOURCES of that host). We call the iterator objects the *container-dependent objects*.<sup>2</sup> With a context-insensitive approach, when elements are retrieved through `Iterator.next()` methods (at line 12 and 14),  $o_2$  and  $o_7$  will flow from  $e$  (the parameter of method `add()`) to `next` (the return value of method `next()`) through a series of in-container PFG edges, and both flow to  $r_1$  and  $r_2$ , leading to an imprecise result. To handle containers comprehensively for better precision, we need to take such usages into account.

To handle container-dependent objects, our principle remains the same as the one in Section 3.3.1: the addition of return edges from EXITS to the  $\text{Exit}_{cs}^{lhs}$  should be suppressed, and shortcut edges should be added to connect matched  $\text{ENTRANCE}_{cs}^{arg}$  and  $\text{Exit}_{cs}^{lhs}$  if  $pt_H(\text{ENTRANCE}_{cs}^{rv})$  and  $pt_H(\text{Exit}_{cs}^{rv})$  overlap. The difference is that we need to expand the set of ENTRANCES and EXITS, as well as the rules for deriving  $pt_H$ . In our example, elements of containers are retrieved by the method `next()` (lines 12 and 14); hence, the additional conflated path here is the path  $e \rightsquigarrow \text{next}$  (shown in

<sup>2</sup>Container-dependent objects include not only iterators, but also some other objects that depend on certain containers, such as the collection views of Map (e.g., the return value of `Map.keySet()`), an example of which is shown in Figure 12 in Section 4.3.

589 the PFG of Figure 4), with its ENTRANCE still being `add()`, but EXIT being `next()`. As for  $pt_H$ , we  
 590 need to expand this mapping relation by (1) expanding the  $pt_H$  for pointers that point to container-  
 591 dependent objects (we call these pointers  $p_{dep}$ , e.g., `it1` and `it2` are  $p_{dep}$  because they point to  
 592 Iterator objects), and (2) for each  $p_{dep}$ , adding the abstract container instance(s), i.e., `host(s)`, that  
 593 they depend on (say  $h$ ), into their  $pt_H$  (e.g., we want  $pt_H(it1) = \{h_1\}$  and  $pt_H(it2) = \{h_2\}$ ). Now the  
 594 question left is, how does one identify  $h$  for different  $p_{dep}$  pointers? To answer this question, we first  
 595 define a new set of methods called *host-propagation methods*. If a container-dependent object (e.g.,  
 596 an Iterator object) is created and returned by invoking a method  $m$  declared in the container class  
 597 (the class of a container-dependent object is typically an inner class of its corresponding container  
 598 class, and hence can access the container’s content). We call  $m$  a *host-propagation method*. In our  
 599 example, `iterator()` at line 11 is a host-propagation method because it is a method declared in the  
 600 `ArrayList` class, and invoking it creates/returns an Iterator object that is pointed to by  $p_{dep}$  (`it1` at  
 601 line 11). Thus, to map a  $p_{dep}$  to  $h$ , for each call-site of a host-propagation method, we propagate the  
 602 hosts at its receiver variable (e.g., `hosts` in  $pt_H(11)$ , at call-site line 11) to its LHS variable (e.g., `it1` at  
 603 line 11). As a result, for this example we have  $pt_H(it1) = \{h_1\}$  because  $pt_H(it1) \supseteq pt_H(11) = \{h_1\}$ .

604 To account for aliasing when building the  $pt_H$  relation (for example, when two pointers refer  
 605 to the same `ArrayList` heap object, then both should map to the host abstracting that object), its  
 606 derivation relies on the points-to information obtained during pointer analysis. (The detailed rules  
 607 are presented in Section 4.3.) In turn,  $pt_H$  guides the pointer analysis in adding precise *shortcut*  
 608 edges. Consequently, the construction of  $pt_H$  and the pointer analysis are performed in a mutually  
 609 recursive fashion, making its derivation inherently on-the-fly: our analysis proceeds in a single  
 610 round, where  $pt_H$  and  $pt$  are computed together until a fixed point is reached.

611 To sum up, by statically abstracting container instances as hosts and identifying their SOURCES  
 612 (input elements) and TARGETS (output locations) on-the-fly, we can add *shortcut* PFG edges that  
 613 directly connect its SOURCES and TARGETS within the pointer analysis. This approach allows us  
 614 to avoid relying on imprecise interprocedural return flows—where return values from different  
 615 container instances may be merged—to determine the points-to sets of container outputs. As a  
 616 result, our approach improves analysis precision without impairing scalability.

617 Our scheme requires specifying a small set of container-related APIs in advance—such as EN-  
 618 TRANCES, EXITS, and *host-propagation methods*—while all remaining behaviors are handled auto-  
 619 matically by the analysis. Additional mechanisms needed to achieve full precision are omitted here  
 620 for brevity and deferred to Section 4.3. For generality, we focus on JDK container classes, whose  
 621 APIs are stable and well-documented. Because the classification of relevant APIs is typically clear  
 622 from their method names, one author was able to complete the necessary specification in about  
 623 five hours. For custom or third-party containers, we ensure soundness by conservatively disabling  
 624 the suppression of corresponding PFG edges—a mechanism detailed in Appendix A. Compared to  
 625 past work [3, 14], which requires to rewrite code implementations for related container APIs, our  
 626 scheme is more general and simpler. Moreover, unlike prior approaches—which still rely on context  
 627 sensitivity for analyzing container methods [3, 14, 21]—our analysis is entirely context-insensitive.

### 629 3.4 Local-Flow Pattern

630 The local-flow pattern is a lightweight pattern that addresses the case of precision loss when the  
 631 values of a method’s parameters flow to the return variable via a series of local assignments. When  
 632 such a method is called from multiple call-sites, their argument values are merged into the method,  
 633 and propagated together to the LHS variable at each call-site, leading to imprecise points-to results.  
 634 We use the example in Figure 5 to illustrate this pattern. In method `select()`, the values of its  
 635 parameters `p1` and `p2` will flow to return variable `r`. It has two call-sites at lines 12 and 16. With a  
 636 context-insensitive approach, objects  $o_{10}, o_{11}, o_{14}$  and  $o_{15}$  will be passed to and merged in `select()`,

```

638 1 A select(A p1,A p2){ 9 // usage code
639 2   A r;           10 A a1 = new A(); //o10
640 3   if (...)        11 A a2 = new A(); //o11
641 4     r = p1;       12 A r1 = select(a1, a2);
642 5   else            13
643 6     r = p2;       14 A a3 = new A(); //o14
644 7   return r;      15 A a4 = new A(); //o15
8 }               16 A r2 = select(a3, a4);

```

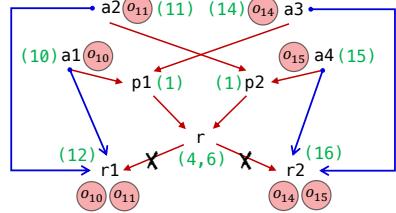


Fig. 5. An example of the local-flow pattern.

and then flow together to both  $r_1$  and  $r_2$  imprecisely: in any execution,  $r_1$  ( $r_2$ ) only point to  $o_{10}$  or  $o_{11}$  ( $o_{14}$  or  $o_{15}$ ).

Following the over-arching principle of CUT-SHORTCUT, for this pattern, multiple calls to such a method  $m$  (whose parameter values flow to the return variable) forms a conflated path (paths) that starts from the specific method parameter(s) and ends at the return value (e.g.,  $p_1 \rightarrow r$  and  $p_2 \rightarrow r$  are two conflated paths), with ENTRANCE and EXIT both being  $m$ . We focus on the cases that the return values must all come from the method parameters via (potentially a series of) local assignments (excluding other sources of return values, such as field loads and method calls), and the detection of such cases relies on an intraprocedural value-flow analysis. Hence, we can safely use arguments (which correspond to the parameters) as the SOURCES, and add shortcut edges from them to the TARGETS, i.e., the LHS variables of the call-sites. In this example, the addition of the return edges from  $r$  are suppressed, and we add shortcut edges  $a_1 \rightarrow r_1$  and  $a_2 \rightarrow r_1$  ( $a_3 \rightarrow r_2$  and  $a_4 \rightarrow r_2$ ) for the call-site at line 12 (16). As a result,  $r_1$  ( $r_2$ ) points to only  $\{o_{10}, o_{11}\}$  ( $\{o_{14}, o_{15}\}$ ) in CUT-SHORTCUT, which is precise. We formally define this analysis pattern in Section 4.4.

*Limitations of CUT-SHORTCUT.* Note that in CUT-SHORTCUT, the over-arching principle (Figure 2) is instantiated as three analysis patterns that are designed for the analysis of Java programs. Consequently, CUT-SHORTCUT cannot be directly applied to other programming languages; however, the high-level idea, namely suppressing the addition of imprecise flow edges and adding semantics-preserving precise flow edges in the PFG might inspire techniques for improving the precision of pointer analysis for other languages. For example, other programming languages may have precision-loss patterns similar to the patterns we proposed. Regarding the applicability of CUT-SHORTCUT to other abstractions like context sensitivity, our current approach is not designed to be pluggable into other (selective) context-sensitive analyses. Rather than selecting methods and analyzing them context-sensitively, CUT-SHORTCUT directly guides PFG construction without applying any contexts. However, it could be beneficial to have such a combination. For instance, the methods whose PFG edges are not affected by our approach, but considered by other selective context-sensitivity approaches [28, 39–41, 70], could be analyzed context-sensitively for possibly better precision. As for flow sensitivity, like almost all whole-program pointer analyses for Java in the literature [31, 46, 67, 72, 76], our approach is also flow-insensitive. It would be non-trivial to apply CUT-SHORTCUT to a flow-sensitive analysis, but it would be interesting to try it in the future.

## 4 Formalizing Cut-Shortcut

In this section, we formalize our CUT-SHORTCUT pointer-analysis core (Section 4.1) and its three patterns (Sections 4.2–4.4).

### 4.1 Pointer Analysis with CUT-SHORTCUT

In this section, we formalize how the CUT-SHORTCUT principle is integrated into a context-insensitive pointer analysis. We first give in Figure 6 definitions of the domains and notation

variables	$x, y \in \mathbb{V}$	pointers	$p \in \mathbb{P} = \mathbb{V} \cup (\mathbb{O} \times \mathbb{F})$
heap objects	$o_i, o_j \in \mathbb{O}$	points-to sets	$pt : \mathbb{P} \rightarrow \mathcal{P}(\mathbb{O})$
methods	$m \in \mathbb{M}$	PFG	$\mathcal{G} = (\mathcal{N} \subseteq \mathbb{P}, \mathcal{E} \subseteq \mathbb{P} \times \mathbb{P})$
fields	$f, g \in \mathbb{F}$	PFG edges	$a \longrightarrow b \in \mathcal{E}$
instruction labels	$i, j \in \mathbb{L}$	call-graph edges	$\rightarrow_{call} \subseteq \mathbb{L} \times \mathbb{M}$

$\arg_k^i$	the $k$ -th argument of call site $i$
$LHS^i$	the LHS variable (if exists) of call site $i$
$param_k^m$	the $k$ -th parameter of method $m$
$ret^m$	the return variable of method $m$
$def_x$	the set of statements that define variable $x$

Fig. 6. Domains and notations.

used. The top section of Figure 6 lists relevant domains and specifies core functions and relations, while the bottom section provides a set of useful abbreviations. Most of the notation is standard or self-explanatory. Here,  $pt(x)$  denotes the points-to set of pointer  $x$ , mapping it to a subset of heap objects (with  $\mathcal{P}(\mathbb{O})$  denoting the power set of  $\mathbb{O}$ ). A call-graph edge  $i \rightarrow_{call} m$  indicates that call-site  $i$  invokes method  $m$ . The call-graph is constructed on-the-fly during pointer analysis, because precise resolution of dynamic dispatch relies on the points-to information of receiver variables at call-sites. The notation  $param_k^m$  and  $\arg_k^i$  denote method parameters of method  $m$  and invocation arguments of call site at  $i$  (using instruction label  $i$  as its identifier), in the program, respectively. For the special cases where  $k=0$ ,  $param_0^m$  denotes the *this* variable of  $m$ , and  $\arg_0^i$  denotes the receiver variable of call-site  $i$ . To simplify notation, we overload the membership operator  $\in$  to also express containment within methods:  $i \in m$  means instruction  $i$  appears in method  $m$ , and  $x \in m$  indicates that  $x$  is a local variable of  $m$ .

*Pointer Analysis, in a PFG View.* Figure 7 presents our formalization of pointer analysis (for now, ignore rule [SHORTCUT] and [premises in black boxes], which are introduced specifically for CUT-SHORTCUT). Essentially, our formalism is equivalent to ones appearing in the existing literature [50, 67, 72] because they all express Andersen-style pointer analysis for Java. To incorporate the CUT-SHORTCUT approach, we explicitly express the *subset constraints* among pointers by the edges of the PFG (pointer flow graph) [72, 81]. Specifically, all nodes of the PFG are pointers, and by [PROPAGATE], an edge  $a \longrightarrow b$  in the PFG indicates that the objects pointed to by  $a$  should also be pointed to by  $b$  (i.e., a subset constraint). Other rules describe how to generate PFG edges for different kinds of statements: object allocation ([NEW]), where allocation-site-based heap abstraction is used, namely,  $o_i$  denotes all heap objects allocated by instruction  $i$ ), local assignment ([ASSIGN]), instance field load ([LOAD]) and store ([STORE]), and method invocation ([CALL], [PARAM], [RETURN]). The rules for static members and arrays are elided, because accessing static members is straightforward, and array accesses are analyzed with all elements collapsed into a single field of the array object.

*CUT-SHORTCUT.* CUT-SHORTCUT improves precision over context-insensitive pointer analysis by adhering to a single core principle: suppress the addition of PFG edges that introduce imprecision, and add shortcut edges that propagate precise object flows—captured before flow merging occurs. In the inference rules shown in Figure 7, the boxed premises encode the *cut* mechanism, while the rule [SHORTCUT] encodes the *shortcut* mechanism, as detailed below:

- $i \notin cutStore$  in [STORE]: Here,  $cutStore \subseteq \mathbb{L}$  is a set of store statements (identified by their labels). By rule [STORE], if a store statement  $j$  is recorded in  $cutStore$ , then the store edges that would normally be generated for  $j$  are suppressed and not added to the PFG.

$$\begin{array}{c}
\frac{i: x = \text{new } A()}{o_i \in pt(x)} \text{ [NEW]} \quad \frac{x = y}{y \rightarrow x} \text{ [ASSIGN]} \quad \frac{a \rightarrow b}{pt(a) \subseteq pt(b)} \text{ [PROP]} \quad \frac{a \rightarrow b}{a \rightarrow b} \text{ [SHORTCUT]} \\
\\
\frac{x = y.f \quad o_i \in pt(y)}{o_i.f \rightarrow x} \text{ [LOAD]} \quad \frac{i: x.f = y \quad o_i \in pt(x)}{y \rightarrow o_i.f} \text{ [STORE]} \\
\\
\frac{i:r = b.m(a_1, a_2, \dots, a_n) \quad o_j \in pt(b) \quad m' = \text{dispatch}(m, o_j)}{i \rightarrow_{call} m' \quad o_j \in pt(\text{this}^{m'})} \text{ [CALL]} \\
\\
\frac{i:r = b.m(a_1, a_2, \dots, a_n) \quad i \rightarrow_{call} m'}{\forall 1 \leq k \leq n : \arg_k^i \rightarrow \text{param}_k^{m'}} \text{ [PARAM]} \quad \frac{i \rightarrow_{call} m \quad \boxed{ret^m \notin cutRet}}{ret^m \rightarrow \text{LHS}^i} \text{ [RETURN]}
\end{array}$$

Fig. 7. Rules for pointer analysis, with CUT-SHORTCUT.

- $i \notin cutRet$  in [RETURN]:  $cutRet \subset \mathbb{V}$  is a set of return variables. If a return variable  $ret^m$  is in  $cutRet$ , then [RETURN] prevents adding return edges from  $ret^m$  to the left-hand-side (LHS) variable at its call-sites (i.e.,  $\text{LHS}^i$  for each  $i$  such that  $i \rightarrow_{call} m$ ).
- [SHORTCUT]:  $a \rightarrow b$  denotes a shortcut edge between pointers  $a$  and  $b$ . By [SHORTCUT], each such shortcut edge contributes a corresponding PFG edge  $a \rightarrow b$ , preserving intended precise flow.

The edges suppressed by CUT-SHORTCUT are of two types: store edges and return edges, which are controlled by the sets  $cutStore$  and  $cutRet$ , respectively. The current three patterns of CUT-SHORTCUT handle imprecision introduced by these two types of edges. Extensions to CUT-SHORTCUT, such as CUT-SHORTCUT<sup>S</sup> (discussed in Section 7), may introduce new patterns that require suppressing additional types of edges. We emphasize that the derivation of  $cutStore$  and  $cutRet$  is performed prior to the pointer analysis. This pre-computation ensures that, once the main analysis begins, it is already known which edges should be suppressed—guaranteeing monotonicity of the analysis (i.e., no edge removals occur during the fix-point iteration). Our main analysis itself is then a single-round process, where PFG construction, call-graph construction (i.e., inference of  $\rightarrow_{call}$ ), and the derivation of the *pointer-host map* ( $pt_H$ ) used in the container-access pattern (detailed in Section 4.3) are performed on-the-fly in a mutually recursive fashion with the context-insensitive pointer analysis.

Next, we describe how the sets  $cutStore$  and  $cutRet$  are computed, and how shortcut edges are inferred, for each of the three patterns in Sections 4.2–4.4.

## 4.2 Field-Access Pattern

In this section, we formalize the field-access pattern of CUT-SHORTCUT.

**4.2.1 Handling of Store.** Figure 8 gives the rules for deriving  $cutStore$  and shortcut edges for handling field stores in the field-access pattern, including support for nested calls of stores. Intuitively, we suppress the addition of (imprecise) PFG edges introduced by a store statement and instead add precise shortcut edges at the call site (of the method containing the store statement)—when the base and right-hand-side (RHS) variables of the store (e.g.,  $x$  and  $y$  in  $x.f = y$ ) derive their values solely from method parameters (including this). To facilitate shortcut-edge inference, we define an auxiliary set called  $tempStore$ , which contains triples of the form  $\langle base, f, from \rangle$  to represent (potential) store operations. To support handling nested calls of field stores (as discussed in Section 3.2.3), we construct and propagate *temp stores* along call chains—starting from the innermost callee (which contains the actual store statement) and continuing through its callers—until the value of  $base$  or  $from$  depends (transitively) on something other than the parameter of a method (e.g., a local allocation, or the return value from a method invocation).

$$\begin{array}{c}
 \frac{i: x.f = y \quad i \in m \quad \text{param}_{k_1}^m = x}{\text{param}_{k_2}^m = y \quad \text{def}_x = \emptyset \quad \text{def}_y = \emptyset} \quad [CUTSTORE] \quad \frac{i \rightarrow_{call} m \quad \text{param}_k^m = x \quad \text{def}_x = \emptyset}{x \leftarrow \arg_k^i} \quad [ARG2VAR] \\
 \hline
 i \in \text{cutStore} \qquad \qquad \qquad x \leftarrow \arg_{k_1}^j \quad y \leftarrow \arg_{k_2}^j \qquad \qquad \qquad x \leftarrow \arg_k^i
 \end{array}$$
  

$$\frac{\frac{\frac{(i: x.f = y \wedge i \in \text{cutStore}) \vee \langle x, f, y \rangle \in \text{tempStore}}{x \leftarrow \arg_{k_1}^j \quad y \leftarrow \arg_{k_2}^j} \quad [PROPSTORE]}{\langle \arg_{k_1}^j, f, \arg_{k_2}^j \rangle \in \text{tempStore}}}{\langle base, f, from \rangle \in \text{tempStore} \quad base, from \in m \quad o_j \in pt(base)} \quad [SHORTCUTSTORE] \\
 \frac{((\text{def}_base \neq \emptyset) \vee (\text{def}_from \neq \emptyset) \vee (\nexists k : \text{param}_k^m = base) \vee (\nexists k : \text{param}_k^m = from))}{from \longrightarrow o_j.f}$$

Fig. 8. Rules for handling field stores in the field access pattern.

*Cut.* We compute *cutStore* using rule [CUTSTORE]. For a store statement  $i: x.f = y$ , if both  $x$  and  $y$  receive their values from method parameters and are not redefined within the method, then  $i$  is added to *cutStore*, indicating that the corresponding store edge should not be added into the PFG. This inference is performed prior to the main pointer analysis and is straightforward to compute, because it does not rely on any intermediate analysis results.

*Shortcut.* To simplify the inference rules, we introduce the auxiliary notation  $\leftarrow$ , defined by [ARG2VAR]. The judgment  $x \leftarrow \arg_k^i$  states that variable  $x$  in method  $m$  receives its value from the  $k^{th}$  argument of call site  $i$  and is not redefined in  $m$ —implying that  $x$  can only point to objects originating from  $\arg_k^i$ . We define two rules, [PROPSTORE] and [SHORTCUTSTORE], to derive shortcut edges for field stores. Rule [PROPSTORE] describes the construction and propagation of *temp stores*—either from store statements marked in *cutStore* or from existing *temp stores*—along call chains. Rule [SHORTCUTSTORE] determines when such propagation should terminate: if the disjunction in its premise holds, then the current *temp store* cannot be propagated further, and shortcut edges should be generated instead.<sup>3</sup>

**4.2.2 Handling of Load.** Figure 9 shows the rules for handling field loads in field-access pattern. Specifically, if return values of a method  $m$  are loaded from the objects that are passed to  $m$  as arguments, then precision loss can occur when arguments from different call sites of  $m$  are merged inside  $m$ . For such cases, the PFG edges from return variables of  $m$  to LHS variables of its call sites should be suppressed. Similar to the handling of stores, we define a set *tempLoad* to help identify what shortcuts should be added; *tempLoad* contains triples of the form  $\langle to, base, f \rangle$ , which represents a load operation  $to = base.f$ .

$$\frac{\frac{\frac{i: x = y.f \quad i \in m \quad \text{ret}^m = x}{y = \text{param}_k^m \quad |\text{def}_x| = 1 \quad \text{def}_y = \emptyset} \quad [CUTLOAD]}{\text{ret}^m \in \text{cutRet} \quad \langle x, y, f \rangle \in \text{tempLoad}}}{\langle to, base, f \rangle \in \text{tempLoad} \quad base \leftarrow \arg_k^i \quad o_i \in pt(\arg_k^i)} \quad [SHORTCUTLOAD]$$

Fig. 9. Rules for handling field loads in the field access pattern.

<sup>3</sup> In principle, if a callee contains  $x$  field stores and is invoked at  $y$  distinct call sites, a naïve construction of shortcut-store edges could introduce up to  $x \times y$  additional edges. In our implementation, we impose a cap of at most 10 field-store statements per handled method. Methods exceeding this threshold (i.e., with more than 10 field stores that satisfy the premise of [CUTSTORE]) are excluded from edge suppression and shortcut-edge insertion. Consequently, the number of shortcut edges added is bounded by a small constant (10) times the number of call sites. Similar caps are applied to our handling of field loads and to the local-flow patterns discussed later.

*Cut.* [CutLoad] defines how to derive  $cutRet$ . For a load statement  $i : x = y.f$ , if (i)  $x$  is the return variable  $ret^m$  and is only defined by this load statement (i.e.,  $|def_x| = 1$ ), and (ii)  $y$  obtains its value from method parameter and is not re-defined, then we add  $ret^m$  to  $cutRet$ , indicating that the return edges from  $ret^m$  to the LHS variables of call sites of  $m$  are to be suppressed. In addition, we record the triple  $\langle x, y, f \rangle$  in  $tempLoad$ . This inference is performed prior to the main pointer analysis.

**839**      *Shortcut.* In [SHORTCUTLOAD], for each invocation of a method that contains a load statement  
**840**      recorded in  $\text{tempLoad}$ , the argument ( $\text{arg}_k^j$ ) flowing to the base of the actual load statement is used  
**841**      as the new base to generate shortcut edges from instance fields to the LHS variable ( $\text{LHS}^j$ ) of call  
**842**      site  $j$ .

Compared with our earlier conference paper [48], the handling of *field loads* has been refined to consider only a single layer of calls, excluding nested calls. This change was motivated by two observations. First, experimental results indicated that handling nested calls of field loads is costly in time and space, but does not yield significant precision improvements. Second, supporting such handling would require performing strong updates to remove existing pointer-flow edges, which would cause the analysis to be non-monotonic.

850 4.3 Container-Access Pattern

JDK offers a powerful but complex *collection framework*. To apply the CUT-SHORTCUT principle to reduce the precision loss caused by containers, we introduce a series of concepts that model the key behaviors of containers.

We first introduce the notation that we use to explain how the container-access pattern is handled in CUT-SHORTCUT. (A table of the notation used is given in Figure 10.) A type  $\tau \in \mathbb{T}$  denotes a Java class or interface, and  $\mathcal{T}$  maps pointers/objects to types. A kind  $c \in \mathbb{K}$  classifies container elements as Coll (collection elements), Mkey (map keys), or Mval (map values), which separates keys from values for analyzing maps.<sup>4</sup> Hosts  $\mathbb{H}$  are defined as  $\mathbb{L} \times \mathbb{K}$ , where  $\mathbb{L}$  are instruction labels and  $\mathbb{K}$  are kinds—a host  $h_c^i$  abstracts the container instance(s) allocated at line  $i$  that holds elements of kind  $c$  (subscripts and superscripts are omitted when clear from context).  $pt_H$  is the pointer-host mapping introduced in Section 3.3.1, mapping each pointer to a set of hosts.  $<$ : denotes the standard subtyping relation in Java.  $\Leftarrow$  and  $\Rightarrow$  relate hosts to their SOURCES and TARGETS (as introduced in Section 3.3.1):  $h \Leftarrow s$  means that objects pointed to by  $s$  are stored as elements in the container instance(s) abstracted by  $h$  (we say  $s$  is a SOURCE of  $h$ ); conversely,  $h \Rightarrow t$  means that pointer  $t$  is a TARGET of host  $h$ , i.e., it receives elements stored in the container(s) abstracted by  $h$ . Finally, we define  $\triangleleft$  as a preorder over hosts, referred to as *host dependency*. Intuitively,  $h' \triangleleft h$  means that the elements stored in the container instance(s) abstracted by  $h'$  are also regarded as elements of those abstracted by  $h$ . In other words, the SOURCES of  $h$  depend on those of  $h'$ . This relation facilitates the propagation of SOURCES across hosts, especially for *bulk-entrance methods* like `ArrayList.addAll`.

types	$\tau \in \mathbb{T}$	subtyping	$<: \subseteq \mathbb{T} \times \mathbb{T}$
kinds	$c \in \mathbb{K} = \{\text{Coll}, \text{Mkey}, \text{Mval}\}$	host sources	$\Leftarrow \subseteq \mathbb{P} \times \mathbb{H}$
hosts	$h \in \mathbb{H} = \mathbb{L} \times \mathbb{K}$	host targets	$\Rightarrow \subseteq \mathbb{P} \times \mathbb{H}$
pointer-host map	$pt_{\mathbb{H}} : \mathbb{P} \rightarrow \mathcal{P}(\mathbb{H})$	host dependency	$\triangleleft \subseteq \mathbb{H} \times \mathbb{H}$

Fig. 10. Additional notation for the container-access pattern.

We now define four sets that specify which methods are container-manipulation operations, and of what kind. In effect, these sets configure CUT-SHORTCUT to implement the container-access pattern for the indicated methods. (These configuration sets are typically populated by the analysis

<sup>881</sup> <sup>4</sup>We also support map entries via an additional kind, omitted here for brevity as it is not central to the core design.

$$\begin{array}{c}
883 \quad \frac{\langle m, \_ \rangle \in \text{CONTEXIT}}{\text{ret}^m \in \text{cutRet}} \text{ [CUTCONT]} \qquad \frac{h \Leftarrow s \quad h \Rightarrow t}{s \rightarrow\!\!\!\rightarrow t} \text{ [SHORTCUTCONT]} \\
884 \\
885 \\
886 \quad \frac{i: x = \text{new } A() \quad A <: \text{Collection}}{h_{\text{Coll}}^i \in pt_H(x)} \text{ [COLHOST]} \qquad \frac{i: x = \text{new } A() \quad A <: \text{Map}}{h_{\text{Mkey}}^i \in pt_H(x) \quad h_{\text{Mval}}^i \in pt_H(x)} \text{ [MAPHOST]} \\
887 \\
888 \\
889 \quad \frac{\langle m, k, c \rangle \in \text{CONTENTR}}{i \rightarrow \text{call } m \quad h_c^l \in pt_H(\arg_0^i)} \text{ [HOSTSOURCE]} \qquad \frac{\langle m, c \rangle \in \text{CONTEXIT}}{i \rightarrow \text{call } m \quad h_c^l \in pt_H(\arg_0^i)} \text{ [HOSTTARGET]} \\
890 \quad h_c^l \Leftarrow \arg_k^i \\
891 \\
892 \quad \frac{\langle m, c \rangle \in \text{CONTPROP}}{i \rightarrow \text{call } m \quad h_c^l \in pt_H(\arg_0^i)} \text{ [HOSTPROP-I]} \qquad \frac{a \rightarrow b \quad h \in pt_H(a)}{\neg(a = \text{ret}^m \wedge \langle m, \_ \rangle \in \text{CONTPROP})} \text{ [HOSTPROP-II]} \\
893 \quad h_c^l \in pt_H(\text{LHS}^i) \\
894 \\
895 \\
896 \quad \frac{i \rightarrow \text{call } m \quad h_c^l \in pt_H(\arg_0^i)}{\langle m, k, c \rangle \in \text{CONTBULKENTR} \quad h_c^{l'} \in pt_H(\arg_k^i)} \text{ [HOSTBULKESTR]} \qquad \frac{h' \triangleleft h \quad h' \Leftarrow s}{h \Leftarrow s} \text{ [HOSTDEP]} \\
897 \\
898 \quad h_c^{l'} \triangleleft h_c^l \\
899
\end{array}$$

Fig. 11. Rules for handling the container-access pattern.

designer, based on the JDK, but could be added to by a user so that user-defined container classes are treated according to the container-access pattern by CUT-SHORTCUT.)

- CONTEXIT: a set of pairs  $\langle m, c \rangle$ , where  $m$  is an EXIT method that returns elements of kind  $c$ .
- CONTENTR: a set of triples  $\langle m, k, c \rangle$ , where  $m$  is an ENTRANCE method, and its  $k^{th}$  parameter stands for an input element of kind  $c$ , to be stored as an element of the container.
- CONTPROP: a set of pairs  $\langle m, c \rangle$ , where a *host-propagation method*  $m$  (introduced in Section 3.3.2) propagates hosts of kind  $c$  from the receiver at call site  $i$  ( $\arg_0^i$ ) to its LHS variable ( $\text{LHS}^i$ ).
- CONTBULKENTR: a set of triples  $\langle m, k, c \rangle$ , where  $m$  is a *bulk-entrance method* that adds all elements of kind  $c$  from one container (the  $k^{th}$  parameter) into another container represented by the receiver. For example,  $\langle \text{ArrayList}.addAll(\text{int}, \text{Collection}), 1, \text{Coll} \rangle$  specifies that all elements (of kind Coll) from its second parameter (a collection) are added into the receiver.

Given the above sets, the container-access pattern is handled according to the inference rules listed in Figure 11.

*Cut.* By [CUTCONT], all return edges from the methods specified by CONTEXIT must be suppressed. The inferences with respect to this rule are performed prior to the main pointer analysis, ensuring monotonicity.

*Shortcut.* [SHORTCUTCONT] describes the addition of shortcut edges by connecting SOURCES and TARGETS associated with the same host. Specifically, whenever both  $h \Leftarrow s$  (pointer  $s$  is a SOURCE of host  $h$ ) and  $h \Rightarrow t$  (pointer  $t$  is a TARGET of host  $h$ ) hold, a shortcut edge  $s \rightarrow\!\!\!\rightarrow t$  is introduced.

The remaining rules compute  $pt_H$  and derive tuples in the relations  $\Leftarrow$  and  $\Rightarrow$ :

[COLHOST] and [MAPHOST] introduce hosts at container-allocation sites. When the allocated object is a subtype of Collection, a Coll host is added to the  $pt_H$  of the receiving variable. When it is a subtype of Map, two distinct hosts of kinds Mkey and Mval are created for the same object  $o_i$ , thereby distinguishing keys from values.

[HOSTSOURCE] derives the relation  $\Leftarrow$  at each call site  $i$  invoking a container ENTRANCE method  $m$ . If the receiver variable ( $\arg_0^i$ ) at  $i$  is associated with a host of kind  $c$  (specified in the triple  $\langle m, k, c \rangle$ ), then the  $k^{th}$  argument at the call site ( $\arg_k^i$ ) is marked as a SOURCE of that host. Similarly,

[HOSTTARGET] derives  $\Rightarrow$  for each call-site  $i$  invoking a container EXIT method  $m$ —for each host of kind  $c$  found in  $pt_H$  of the receiver variable, the LHS variable is marked as its TARGET.

[HOSTPROP-I] handles call sites that invoke *host-propagation methods*. If a method  $m$  is specified to propagate hosts of kind  $c$ , then these hosts are propagated from the receiver variable ( $arg_0^i$ ) to the LHS variable ( $LHS^i$ ) at each invocation  $i$  of  $m$ . We illustrate this mechanism with an example involving the keySet view of a map (see Figure 12). Line 1 allocates a Map, resulting in two hosts (of kinds Mkey and Mval) added to  $pt_H(x)$  via [MAPHOST]. Line 2 invokes an ENTRANCE method `put()`. By [HOSTSOURCE], the arguments  $k$  and  $v$  are associated as SOURCES of the respective hosts

```

1 Map x = new Map();           // o1, ptH(x) = {hMkey1, hMval1}
2 x.put(k, v);                // hMkey1 ≈ k, hMval1 ≈ v
3 Set s = x.keySet();         // ptH(s) = {hMkey1}
4 Iterator it = s.iterator(); // ptH(it) = {hMkey1}
5 Object r = it.next();        // hMkey1 ⇒ r

```

in  $pt_H(x)$  (due to the specified triples  $\langle \text{Map.put}, 1, \text{Mkey} \rangle$  and  $\langle \text{Map.put}, 2, \text{Mval} \rangle$  in CONTENTR). Line 3 propagates only the Mkey host to  $s$  because we specify  $\langle \text{Map.keySet}, \text{Mkey} \rangle \in \text{CONTPROP}$ . Line 4 further propagates the Mkey host from  $s$  to  $it$ , because iterator is specified to propagate any hosts (specified

Fig. 12. An example of Map container usage.

as  $\langle \text{Collection.iterator}, \_ \rangle$  in CONTPROP). Line 5 invokes the EXIT method `next()`, associating  $r$  as a TARGET of the MKey host. Finally, a shortcut edge  $k \rightarrow r$  is introduced by [SHORTCUTCONT], linking the SOURCE and TARGET of this host. This example also highlights the role of host kinds in handling maps precisely.

[HOSTPROP-II] derives  $pt_H$  by propagating hosts along the PFG, which is essential for handling aliasing. If  $pt_H(x) \neq \emptyset$  and  $y$  is an alias of  $x$ , then  $pt_H(y)$  should include the hosts associated with  $x$ , because later invocations of ENTRANCE or EXIT methods may use  $y$  as the receiver. To ensure this behavior, we allow hosts to propagate along PFG edges, analogous to how heap objects flow in standard pointer analysis, allowing  $pt_H(p)$  to be derived for every pointer  $p$  that may reference a container instance. However, such propagation is excluded for one case: return edges of a host-propagation method  $m$ . The reason propagation is suppressed is because (1) [HOSTPROP-I] already derives  $pt_H$  for  $LHS^i$  when  $i \rightarrow_{call} m$ , and (2)  $\text{ret}^m$  may merge multiple hosts in its  $pt_H$ , so blindly propagating them out of the called method  $m$  would cause imprecision.

[HOSTBULKETR] handles *bulk-entrance methods*. In Figure 13, line 5 calls `addAll()` to add elements from container  $o_1$  into  $o_3$ . For soundness, the analysis must recognize not only  $a$  but also  $b$  as SOURCES of host  $h_{\text{Coll}}^1$ . To ensure this behavior, we define a *host-dependency* preorder  $\triangleleft$  whose semantics is given by [HostDep]. [HOSTBULKETR] infers such dependencies between hosts of the receiver and argument containers at bulk-entrance call-sites.<sup>5</sup> Because the standard configuration of the set CONTBULKESTR contains the tuple  $\langle \text{Collection.addAll}, 1, \text{Coll} \rangle$ , the [HOSTBULKETR] rule causes all Coll hosts in  $pt_H(12)$  (the first, i.e.,  $k=1$  argument) to be related to the receiver's Coll hosts

```

1 ArrayList l1 = new ArrayList(); // o1, ptH(l1) = {hColl1}
2 l1.add(a);                      // hColl1 ≈ a
3 ArrayList l2 = new ArrayList(); // o3, ptH(l2) = {hColl3}
4 l2.add(b);                      // hColl3 ≈ b
5 l1.addAll(l2);                  // hColl1 ⊲ hColl3, hence hColl1 ≈ b
6 Object r = l1.get(0);           // hColl1 ⇒ r

```

in  $pt_H(11)$ . Therefore, in this case we obtain  $h_{\text{Coll}}^3 \triangleleft h_{\text{Coll}}^1$ , which then triggers [HostDep] to propagate SOURCES from  $h_{\text{Coll}}^3$  to  $h_{\text{Coll}}^1$ , thereby making  $b$  a SOURCE of the latter. Consequently, two shortcut edge  $a \rightarrow r$  and  $b \rightarrow r$  are derived by [SHORTCUTCONT] once  $r$  is identified as a TARGET of  $h_{\text{Coll}}^1$  at line 6.

Fig. 13. An example of bulk operation for containers.

<sup>5</sup>The  $\triangleleft$  preorder, by definition it is a transitive relation. We omit the rules that define the transitivity property of  $\triangleleft$ .

Compared to the earlier conference paper [48], we provide more detailed explanations of the rules, illustrated with examples, including the treatment of *bulk-entrance methods*, which were previously omitted due to space constraints. We also extend the discussion of conservative fallback mechanisms (see Appendix A), covering sound handling for custom containers and the propagation of hosts to this variables. Moreover, the mechanisms of the container-access pattern have been refined to reduce the amount of information that the analysis designer has to specify, and improving the efficiency of the analysis.

#### 4.4 Local-Flow Pattern

Figure 14 gives the rules to derive *cutRet* and shortcut edges for the local-flow pattern.

$$\begin{array}{c}
 \frac{\text{def}_{\text{param}_k^m} = \emptyset}{\langle m, k \rangle \rightarrowtail \text{param}_k^m} \quad [\text{PARAM2VAR}] \qquad \frac{\begin{array}{c} \mathbf{i: x = y'} \quad \langle m, k \rangle \rightarrowtail y' \\ (\forall j \in \text{def}_x : (\mathbf{j: x = y}) \wedge (\langle m, \_ \rangle \rightarrowtail y)) \end{array}}{\langle m, k \rangle \rightarrowtail x} \quad [\text{PARAM2VARREC}] \\
 \hline
 \frac{\langle m, \_ \rangle \rightarrowtail \text{ret}^m}{\text{ret}^m \in \text{cutRet}} \quad [\text{CUTLFLOW}] \qquad \frac{i \rightarrow_{\text{call}} m \quad \langle m, k \rangle \rightarrowtail \text{ret}^m}{\arg_k^i \rightarrowtail \text{LHS}^i} \quad [\text{SHORTCUTLFLOW}]
 \end{array}$$

Fig. 14. Rules for local-flow pattern.

To formalize the local-flow pattern, we introduce the notation  $\rightarrowtail$ , as defined by rules [\[PARAM2VAR\]](#) and [\[PARAM2VARREC\]](#). The judgment  $\langle m, k \rangle \rightarrowtail x$  states that  $x$  is a local variable in method  $m$  that obtains its value from the  $k^{th}$  parameter of  $m$ . Moreover, [\[PARAM2VAR\]](#) and [\[PARAM2VARREC\]](#) together ensure that when  $\langle m, k \rangle \rightarrowtail x$  holds, the values of  $x$  flow solely from  $m$ 's parameters—possibly from multiple parameters (e.g., both  $\langle m, k \rangle \rightarrowtail x$  and  $\langle m, k' \rangle \rightarrowtail x$  may hold at the same time)—via zero or more local assignments. Crucially, this property guarantees that  $x$  does not receive values from other sources, such as field loads or method calls. With this notation in place, the edges to be suppressed—and shortcut edges to add—for local-flow pattern can be defined cleanly and precisely:

*Cut.* By [\[CUTLFLOW\]](#), if we find that the values of return variable  $\text{ret}^m$  all come from  $m$ 's parameter(s), then we record  $\text{ret}^m$  in *cutRet*. The inference of [\[PARAM2VAR\]](#), [\[PARAM2VARREC\]](#) and [\[CUTLFLOW\]](#) is performed prior to the main pointer analysis.

*Shortcut.* By [\[SHORTCUTLFLOW\]](#), if  $\langle m, k \rangle \rightarrowtail \text{ret}^m$  holds, then for each call site  $i$  of method  $m$ , we add a shortcut edge from the  $k^{th}$  argument of  $i$  (i.e.,  $\arg_k^i$ ) to its LHS variable  $\text{LHS}^i$ .

## 5 Soundness

In this section, we formally prove the soundness of CUT-SHORTCUT, based on the inference rules defined in Sections 4.1-4.4.

We begin by introducing notation. Let  $\mathcal{A}$  denote a pointer-analysis instance over an input program, with subscripts identifying specific approaches (e.g.,  $\mathcal{A}_{ci}$  for context-insensitive analysis,  $\mathcal{A}_{csc}$  for CUT-SHORTCUT). For analysis  $\mathcal{A}_a$ , we denote its corresponding pointer-flow graph (PFG) as  $\mathcal{G}_a = (\mathcal{N}_a, \mathcal{E}_a)$ . The set  $\mathcal{E}^\times$  is defined as the collection of PFG edges included in  $\mathcal{E}_{ci}$ , but explicitly suppressed by the boxed premises of rules [\[STORE\]](#) and [\[RETURN\]](#) in Figure 7. Formally:

$$e \in \mathcal{E}^\times \iff [(e = (y, o_j, f) \wedge i : x, f = y \wedge o_j \in \text{pt}(x) \wedge i \in \text{cutStore}) \vee (e = (\text{ret}^m, \text{LHS}^i) \wedge i \rightarrow_{\text{call}} m \wedge \text{ret}^m \in \text{cutRet})] \quad (1)$$

Let  $T = \{t \mid (s, t) \in \mathcal{E}^\times\}$ , and define the set  $\text{shortcutSrc}(t) = \{s \mid s \rightarrowtail t\}$ , where  $\rightarrowtail$  is derived from the four inference rules for shortcut edges in Figures 8, 9, 11, and 14.

Intuitively, we prove soundness as follows: for any object  $o$  that reaches a target  $t \in T$  dynamically via a path  $P$  in  $\mathcal{G}_{ci}$ , suppose the ending edge of this path is suppressed in  $\mathcal{G}_{csc}$ . Then, along path

1030  $P$ , we show that there exists a corresponding source  $s \in \text{shortcutSrc}(t)$ , such that  $o$  can still reach  
 1031  $t$  via a new path in  $\mathcal{G}_{csc}$ , by replacing the sub-path from  $s$  to  $t$  in path  $P$  with the shortcut edge  
 1032 connecting  $s$  to  $t$ .

1033 To formalize reachability from objects to arbitrary pointers, we first address a subtlety in our  
 1034 PFG formulation. Following Sridharan et al. [72], when the pointer-analysis algorithm constructs  
 1035 the call graph on-the-fly, the value flow from a receiver to this is not explicitly represented by  
 1036 an edge. At call site  $i : r = b.m(\dots)$ , the callee is resolved dynamically via  $o_j \in \text{pt}(b)$ , and only  $o_j$  is  
 1037 added to  $\text{pt}(\text{this}^{m'})$  with  $m' = \text{dispatch}(m, o_j)$ . Thus, no edge from  $b$  to  $\text{this}^{m'}$  exists in the PFG.  
 1038 While this way of constructing the PFG avoids imprecise propagation of  $o_k \in \text{pt}(b)$  for  $k \neq j$ , it  
 1039 complicates reasoning about reachability on the PFG. To help present our proof, we define an  
 1040 auxiliary relation  $\mathcal{E}^\dagger \subseteq \mathbb{P} \times \mathbb{P}$  ( $\mathbb{P}$  is the pointer domain) such that  $(\arg_0^i, \text{this}^m) \in \mathcal{E}^\dagger \iff i \rightarrow_{call} m$ .  
 1041 This relation makes the receiver-to-this flow explicit.

1042 *Definition 5.1 (Pointer-Flow Path).* Given analysis  $\mathcal{A}_a$  with edge sets  $\mathcal{E}_a$  and  $\mathcal{E}_a^\dagger$ , a *pointer-flow*  
 1043 *path* for object  $o_i$  is a sequence  $[p_0, p_1, \dots, p_n]$ , where  $p_0$  is the variable receiving  $o_i$  at its allocation  
 1044 site and  $\forall i \geq 0, (p_i, p_{i+1}) \in \mathcal{E}_a \cup \mathcal{E}_a^\dagger$ . (When such a pointer-flow path exists, we say that “ $o_i$  reaches  
 1045  $p_n$  under  $\mathcal{A}_a$ .”) We say such a path exists dynamically if each edge corresponds to a runtime value  
 1046 flow, which we sometimes express as “ $o$  reaches  $p_n$  under dynamic execution.” Consider a sound but  
 1047 imprecise analysis  $\mathcal{A}_a$ : every dynamic pointer-flow path must be included in the over-approximated  
 1048 path space under  $\mathcal{A}_a$ . The converse does not always hold, due to the possibility of spurious flows  
 1049 in an imprecise analysis.

1051 We now prove the core soundness theorem and its supporting lemma, under Assumption 1 and 2.

1052 ASSUMPTION 1. *We assume the soundness of the context-insensitive analysis  $\mathcal{A}_{ci}$ .*

1054 ASSUMPTION 2. *We assume that the input sets `CONTEXIT`, `CONTENTR`, `CONTPROP`, and `CONT-`  
 1055 `BULKENTR` for the container-access pattern of  $\mathcal{A}_{csc}$  are complete with respect to the container classes  
 1056 we consider.*

1057 THEOREM 5.2 (SOUNDNESS OF  $\mathcal{A}_{csc}$ ). *For any object  $o$  and pointer  $p$ , if  $o$  reaches  $p$  under dynamic  
 1058 execution, then  $o$  also reaches  $p$  under  $\mathcal{A}_{csc}$ .*

1059 PROOF. We construct a family of analyses  $\mathcal{A}_{csc}^l$  ( $0 \leq l \leq |\mathcal{E}^\times|$ ), where each  $\mathcal{A}_{csc}^l$  denotes an  
 1060 analysis that suppresses only a subset of the edges in  $\mathcal{E}^\times$ . Specifically, consider  $\mathcal{E}^\times$  as a fixed set  
 1061 (computed once by running  $\mathcal{A}_{csc}$ ), and assign each  $e \in \mathcal{E}^\times$  a unique index from 1 to  $|\mathcal{E}^\times|$ . The  
 1062 analysis  $\mathcal{A}_{csc}^l$  suppresses only those edges with index less than or equal to  $l$ .<sup>6</sup>

1063 For fixed  $l$ , we define the predicate  $Q^l(p)$  over a pointer  $p$  as follows:

$$1064 Q^l(p) \triangleq \text{For any object } o \text{ that reaches } p \text{ dynamically, } o \text{ also reaches } p \text{ under } \mathcal{A}_{csc}^l.$$

1065 We now prove, by induction on  $l$ , that for all  $l$ ,  $\forall p. Q^l(p)$  holds.

- 1066 i) *Base Case.* If  $l = 0$ ,  $\mathcal{A}_{csc}^0$  degenerates to  $\mathcal{A}_{ci}$ . By assumption 1,  $\mathcal{A}_{ci}$  is sound, which directly  
 1067 gives us  $\forall p. Q^0(p)$ .
- 1068 ii) *Inductive Step.* Suppose that  $0 \leq l < |\mathcal{E}^\times|$ , and assume that  $\forall p. Q^l(p)$  holds. We need to prove  
 1069 that  $\forall p. Q^{l+1}(p)$ . Let the suppressed edge indexed by  $l + 1$  be the edge  $(u, v)$ . For an arbitrary  
 1070 pointer  $p$ , we distinguish three cases:
  - 1071 a) If  $p = v$ , we must prove  $Q^{l+1}(v)$ .

1072 <sup>6</sup>To define  $\mathcal{A}_{csc}^l$  more formally: we keep all inference rules unchanged, except for the suppression conditions in `[STORE]`  
 1073 and `[RETURN]` (Figure 7). We replace the boxed premise in `[STORE]` with  $\mathcal{L}(y, o_i.f) > l$ , and that in `[RETURN]` with  
 1074  $\mathcal{L}(\text{ret}^m, \text{LHS}^i) > l$ , where  $\mathcal{L} : \mathbb{P} \times \mathbb{P} \rightarrow \mathbb{N}$  maps each edge in  $\mathcal{E}^\times$  to its index, and assigns all other edges a constant greater  
 1075 than  $|\mathcal{E}^\times|$ , ensuring only edges with index  $\leq l$  are suppressed in the PFG of  $\mathcal{A}_{csc}^l$ .

- 1079 b) If  $p \neq v$  and  $pt(p)$  depends on  $pt(v)$  (directly or indirectly). Because  $Q^l(p)$  holds by  
 1080 the inductive hypothesis, if we can show  $Q^{l+1}(v)$ , then  $Q^{l+1}(p)$  follows.  
 1081 c) If  $pt(p)$  does not depend on  $pt(v)$ , then the suppression of  $(u, v)$  does not affect  $pt(p)$ .  
 1082 Because  $Q^l(p)$  holds,  $Q^{l+1}(p)$  also holds.

1083 Thus, it suffices to prove  $Q^{l+1}(v)$ . To that end, we invoke Lemma 5.3, which guarantees that  
 1084 for any object  $o$  that reaches  $v$  dynamically along a path ending in  $(u, v)$ , there exists some  
 1085  $s \in \text{shortcutSrc}(v)$  on that path. (Among possible candidates, we choose  $s$  closest to  $o$  to  
 1086 ensure  $s$  precedes  $v$  when cycles are involved.) Because the suppression of  $(u, v)$  does not  
 1087 affect whether  $o$  reaches  $s$ , and  $o$  reaches  $s$  under  $\mathcal{A}_{csc}^l$  (by the inductive hypothesis), the  
 1088 dynamic path from  $o$  to  $s$  plus the shortcut edge  $(s, v)$ —added by  $\mathcal{A}_{csc}^{l+1}$ —form a valid path  
 1089 from  $o$  to  $v$  under  $\mathcal{A}_{csc}^{l+1}$ . Therefore,  $Q^{l+1}(v)$  holds.

1090 By induction, for all  $l \in [0, |\mathcal{E}^\times|]$  and all pointers  $p$ , any object  $o$  that reaches  $p$  dynamically also  
 1091 reaches  $p$  under  $\mathcal{A}_{csc}^l$ . Setting  $l = |\mathcal{E}^\times|$  yields  $\mathcal{A}_{csc}^l = \mathcal{A}_{csc}$ , completing the proof.  $\square$   
 1092

1093 The key lemma used in the proof of Theorem 5.2 is stated and proved below.  
 1094

1095 LEMMA 5.3. *Assume that  $\forall p. Q^l(p)$  holds for a given  $l$ , where  $0 \leq l < |\mathcal{E}^\times|$ . Suppose that the sup-  
 1096 pressed edge indexed by  $l+1$  is  $(u, v)$ . Then, for any object that dynamically reaches  $v$  via a pointer-flow  
 1097 path  $P$  that ends with  $(u, v)$ , there exists some  $s \in \text{shortcutSrc}(v)$  that precedes  $v$  on  $P$ , under Assump-  
 1098 tion 2.*

1099 PROOF. We split the proof into cases based on the kind of pattern in CUT-SHORTCUT that is  
 1100 responsible for suppressing the edge  $(u, v)$ .  
 1101

- 1102 i) *Store-Edge Case.* Suppose that  $(u, v)$  corresponds to a field-store edge—i.e., the first disjunct  
 1103 in Equation (1)—and takes the form  $(y, o_j.f)$ , where  $x.f = y$  is a store statement recorded in  
 1104 *cutStore* and  $o_j \in pt(x)$ . By [CUTSTORE], both  $x$  and  $y$  are formal parameters of the enclosing  
 1105 method  $m$ . Therefore, the path  $P$  must traverse an argument  $\arg_{k_1}^i$  (at call site  $i \rightarrow_{call} m$ ),  
 1106 passed to  $y$ . If  $\arg_{k_1}^i$  is still a method parameter and not redefined, then by repeated applica-  
 1107 tions of [PROPSTORE], we can trace back along the call chain to find the outermost arguments  
 1108  $s_y$  for  $y$ , and  $s_x$  for  $x$ , respectively (choosing the one nearest to  $o$  in the case of a cycle). Then,  
 1109 by [SHORTCUTSTORE],<sup>7</sup> a shortcut edge is added from  $s_y$  to  $o_j.f$ , because  $o_j \in pt(s_x)$  under  $\mathcal{A}_{csc}^{l+1}$   
 1110 (guaranteed by  $Q^l(s_x)$  and the fact that suppressing  $(y, o_j.f)$  does not affect  $pt(s_x)$ ). Thus,  
 1111 we find  $s_y \in \text{shortcutSrc}(o_j.f)$ , which precedes  $v$  in  $P$ .
- 1112 ii) *Return-Edge Case.* Suppose that  $(u, v)$  is a return edge, i.e., the second disjunct in Equa-  
 1113 tion (1), of the form  $(\text{ret}^m, \text{LHS}^i)$ , where  $i \rightarrow_{call} m$  and  $\text{ret}^m \in \text{cutRet}$ . We consider three  
 1114 sub-cases: degenerates
- 1115 a) If  $\text{ret}^m \in \text{cutRet}$  is derived from [CUTLOAD] by the *load* part of the field-access pattern.  
 1116 Then  $\text{ret}^m = x$  where  $x = y.f$  and  $y$  is the  $k^{\text{th}}$  parameter of  $m$  and is not redefined. Hence,  
 1117  $y$  originates from  $\arg_k^i$  at the call site  $i \rightarrow_{call} m$ . By [SHORTCUTLOAD], a shortcut edge is  
 1118 added from  $o_j.f$  to  $\text{LHS}^i$  for any  $o_j \in pt(\arg_k^i)$  under  $\mathcal{A}_{csc}^{l+1}$  (guaranteed by the inductive  
 1119

1120 <sup>7</sup>Deriving this shortcut edge requires that the relevant call-graph edge  $i \rightarrow_{call} m$  (and any along the call chain) be derivable  
 1121 under  $\mathcal{A}_{csc}^{l+1}$ . This is ensured by the inductive hypothesis  $\forall p. Q^l(p)$ , which guarantees the soundness of receiver's points-to  
 1122 set under  $\mathcal{A}_{csc}^l$  (for each receiver involved in that call chain). Moreover, suppressing  $(y, o_j.f)$  does not affect the derivation  
 1123 of these specific call-graph edges, as dispatch precedes callee-body analysis. Crucially, this also justifies the need for  
 1124 induction on the index of suppressed edges: earlier suppressed edges may influence the derivability of later call-graph  
 1125 edges (i.e., if the call graph were built imprecisely due to an imprecise receiver). By proceeding incrementally, we ensure  
 1126 that every call-graph edge involved in shortcut edges derivation of each inductive step remains soundly constructed. This  
 1127 reasoning applies to all shortcut-edge-derivation cases, and we omit further repetition for brevity.

- 1128 hypothesis  $Q^l(\arg_k^i)$  and the fact that suppressing the specific return edge  $(\text{ret}^m, \text{LHS}^i)$   
 1129 does not affect  $\text{pt}(\arg_k^i)$ ). Therefore, we find  $o_j.f \in \text{shortcutSrc}(\text{LHS}^i)$  along path  $P$ .  
 1130 b) If  $\text{ret}^m \in \text{cutRet}$  is derived from [CutCont] by the container-access pattern, then  
 1131 container elements are retrieved at call site  $i$ , from a container instance  $o_\kappa$  (or a container-  
 1132 dependent object that depends on  $o_\kappa$ ). To ensure that a suitable shortcut edge can  
 1133 be identified, we first need to establish that the host abstraction  $h_c^\kappa$  for the container  
 1134 instance  $o_\kappa$  exists and is available at call site  $i$ —that is,  $h_c^\kappa \in \text{pt}_H(\arg_0^i)$ . By [ColHost]  
 1135 and [MapHost], such a host  $h_c^\kappa$  is created at the allocation site of  $o_\kappa$  and associated with  
 1136 the LHS variable  $x$  referring to it. This host then propagates to  $\arg_0^i$  through one of  
 1137 two mechanisms: (1) via [HostProp-I], if  $\arg_0^i$  refers to a container-dependent object, or  
 1138 (2) via [HostProp-II], if  $\arg_0^i$  is an alias of  $x$ . Once  $h_c^\kappa$  reaches  $\arg_0^i$ , [HostTarget] derives  
 1139  $h_c^\kappa \Rightarrow \text{LHS}^i$ . We then need to find a suitable  $s$  along path  $P$  that explains how the object  
 1140 entered the container as an input element, and  $s \in \text{shortcutSrc}(\text{LHS}^i)$  can be derived.  
 1141 There are two cases for such an  $s$ : (1) a call site  $j$  invokes an ENTRANCE, where  $s$  is an  
 1142 argument passed in (i.e.,  $s = \arg_k^j$ ); (2) a *bulk-entrance method* inserts all elements of  
 1143 a host  $h'$  (where  $h' \triangleleft h_c^\kappa$ ), and  $h' \Leftarrow s$ . In both cases, we find that a shortcut edge that  
 1144 connects  $s$  to  $\text{LHS}^i$  is indeed derivable. For the first case, similar to the reason why  $h_c^\kappa$   
 1145 is available at call site  $i$ ,  $h_c^\kappa$  is also be available at call site  $j$ , and we derive  $h_c^\kappa \Leftarrow s$  by  
 1146 [HostSource]. For the second case, we still derive  $h_c^\kappa \Leftarrow s$ , but by applying [HostBulkEtr]  
 1147 at the call site of the *bulk-entrance method*, followed by [HostDep], which propagates  
 1148 SOURCES of  $h'$  to  $h_c^\kappa$ . Therefore, [ShortCutCont] will always add a shortcut edge from  $s$   
 1149 to  $\text{LHS}^i$ , establishing that  $s \in \text{shortcutSrc}(\text{LHS}^i)$ , as desired.  
 1150 c) If  $\text{ret}^m$  comes from [CutLFlow] by the local-flow pattern, then  $\langle m, \_ \rangle \rightarrowtail \text{ret}^m$  indicates  
 1151 that  $\text{ret}^m$  depends only on parameters of  $m$ . The dynamic path  $P$  must thus come from  
 1152 some  $\arg_k^i$  passed into  $m$ . By [ShortCutLFlow], a shortcut edge is added from  $\arg_k^i$  to  
 1153  $\text{LHS}^i$ , and we have  $\arg_k^i \in \text{shortcutSrc}(\text{LHS}^i)$ .  $\square$   
 1154

## 1155 6 A CFL-Reachability Formulation of CUT-SHORTCUT

1156 In this section, we show how CUT-SHORTCUT can be formulated as a CFL-reachability problem.  
 1157 Although our implementation is not CFL-reachability-based, this formulation situates our technique  
 1158 within the broader landscape of formal-language-reachability problems. It also clarifies why CUT-  
 1159 SHORTCUT achieves such high efficiency compared with context-sensitivity-based techniques.  
 1160 Specifically, we establish two points: (1) CUT-SHORTCUT can be transformed from the standard  
 1161 *single Dyck- $k$* <sup>8</sup> CFL-reachability formulation of context-insensitive, field-sensitive pointer analysis,  
 1162 without introducing an *additional Dyck- $k$*  language to pair method entry and exits as required for  
 1163 full context sensitivity. Thus, CUT-SHORTCUT remains solvable under the same formal model as  
 1164 context-insensitive analysis, without assuming additional computational power. (2) Relative to the  
 1165 context-insensitive baseline, CUT-SHORTCUT incurs only linear growth in grammar and graph size,  
 1166 thereby preserving asymptotic complexity when solved by a black-box CFL-reachability solver. In  
 1167 contrast, context-sensitivity-based techniques—even regularized ones such as  $k$ -CFA-based context  
 1168 sensitivity—replicate program nodes and edges under different contexts and can cause significant  
 1169 blowup.  
 1170

1171 Section 6.1 reviews the standard CFL-reachability formulation for Java pointer analysis. Sections  
 1172 6.2–6.3 present a two-step transformation that yields the CFL-reachability formulation of CUT-  
 1173 SHORTCUT. Finally, Section 6.4 discusses what the CFL-reachability formulation of CUT-SHORTCUT  
 1174 says about its computational complexity.

1175 <sup>8</sup>In a Dyck- $k$  language,  $k$  denotes the number of distinct types of parentheses defined in the alphabet.

## 6.1 The $L_F$ -Based CFL-Reachability Formulation of Java Pointer Analysis

Many program-analysis problems can be expressed as formal-language reachability problems, including pointer analysis [59]. Suppose that graph  $G$  is a directed graph with edge labels drawn from an alphabet  $\Sigma$ , and  $L$  is a formal language over  $\Sigma$ . Then each path  $p$  in  $G$  can be seen as labeled with a string  $s(p) \in \Sigma^*$  obtained by concatenating—in order—the edge labels along the path. We say that  $p$  is an  $L$ -path if  $s(p) \in L$ . When such a path exists from node  $u$  to node  $v$ , we say that  $v$  is  $L$ -reachable from  $u$ , denoted by  $u \xrightarrow{L} v$ . An  $L$ -reachability problem is a CFL-reachability problem when  $L$  is a context-free language. In that case, for any non-terminal  $S$  in the grammar of  $L$ , if  $s(p)$  belongs to the language generated from  $S$ , we say that  $u \xrightarrow{S} v$ , and refer to  $p$  as an  $S$ -path.

For Java pointer analysis, Andersen-style (context-insensitive) analysis can be formulated as a CFL-reachability problem over a pointer assignment graph (PAG) [74]—a graph representation slightly different from the PFG used in our approach. In the PAG, nodes represent variables and heap objects, and edges capture various types of value assignments. The associated language is denoted by  $L_F$ , where the subscript  $F$  reflects the fact that field sensitivity is enforced by the grammar. The PAG edges for  $L_F$  are constructed according to the rules shown in Figure 15. Unlike the PFG, PAG nodes do not include instance fields; instead, field accesses are encoded as edge labels—specifically,  $\text{store}[f]$  and  $\text{load}[f]$  label the edges generated by rules [ $L_F$ -STORE] and [ $L_F$ -LOAD], respectively.

$$\begin{array}{c}
 \frac{i: x = \text{new } A()} {o_i \xrightarrow{\text{new}} x} [\text{L}_F\text{-NEW}] \quad \frac{x = y} {y \xrightarrow{\text{assign}} x} [\text{L}_F\text{-ASSIGN}] \\
 \\ 
 \frac{x.f = y} {y \xrightarrow{\text{store}[f]} x} [\text{L}_F\text{-STORE}] \quad \frac{x = y.f} {y \xrightarrow{\text{load}[f]} x} [\text{L}_F\text{-LOAD}] \\
 \\ 
 \frac{i:r = b.m(a_1, a_2, \dots, a_n) \quad i \rightarrow_{\text{call}} m'} {b \xrightarrow{\text{assign}} \text{this}^{m'} \quad \forall 1 \leq k \leq n : \arg_k^i \xrightarrow{\text{assign}} \text{param}_k^{m'} \quad \text{ret}^{m'} \xrightarrow{\text{assign}} r} [\text{L}_F\text{-CALL}]
 \end{array}$$

Fig. 15. Rules for building the pointer assignment graph (PAG) for language  $L_F$ .

One important distinction from the PFG-based pointer analysis rules introduced earlier in Figure 7 is that the traditional CFL-reachability formulation of Andersen-style pointer analysis assumes a pre-built call graph—needed to derive the premise  $i \rightarrow_{\text{call}} m'$  in [ $L_F$ -CALL]. Such a call graph can be computed separately using techniques like Class Hierarchy Analysis (CHA) [13]. While recent work has explored CFL formulations that support on-the-fly call-graph construction [25], we follow the traditional formulation for conciseness and clarity. The grammar of  $L_F$  is as follows:

$$\begin{array}{l}
 \text{flowsTo} \longrightarrow \text{new flows}^* \\
 \text{flows} \longrightarrow \text{assign} \mid \text{store}[f] \text{ alias load}[f] \\
 \text{alias} \longrightarrow \overline{\text{flowsto}} \text{ flowsTo} \\
 \overline{\text{flowsto}} \longrightarrow \overline{\text{flows}}^* \text{ new} \\
 \overline{\text{flows}} \longrightarrow \overline{\text{assign}} \mid \overline{\text{load}[f]} \text{ alias } \overline{\text{store}[f]}
 \end{array}$$

Note that the inverse edge of a PAG edge  $x \xrightarrow{l} y$  is defined as  $y \xrightarrow{\bar{l}} x$  (where  $\bar{l}$  denotes the label on the original edge). The generation of inverse edges is not shown explicitly in Figure 15, but can be understood as adding an inverse edge for each forward edge. We define  $u$  and  $v$  to be aliases if there exists an object  $o_i$  such that  $u \xrightarrow{\text{flowsTo}} o_i \xrightarrow{\text{flowsTo}} v$ .

1226 The solution of the  $L_F$ -reachability problem over the PFG is a mapping from variables to (abstract)  
 1227 heap objects: in the mapping, a variable  $v$  points to an object  $o_i$  iff there exists a *flowsTo*-path from  
 1228  $o_i$  to  $v$ .

1229 Because  $L_F$  enforces field sensitivity by matching  $\text{store}[f]$  and  $\text{load}[f]$  as balanced parentheses, it  
 1230 corresponds to a *Dyck- $k$*  language over the family of fields (where  $k$  is the number of distinct fields).  
 1231 This language is a standard context-free language of balanced parentheses, and can be recognized  
 1232 by a single-stack pushdown automaton. A formulation of context-sensitive pointer analysis based  
 1233 on formal-language theory typically introduces an additional context-free language  $L_C$  to enforce  
 1234 context sensitivity by matching method entries and exits [71]. This extension effectively adds a  
 1235 second *Dyck- $k'$*  language over method contexts (where  $k'$  is the number of distinct calling contexts).  
 1236 Solving a pointer analysis that is both context- and field-sensitive thus becomes an *interleaved*  
 1237 *Dyck-reachability problem*: determining paths that simultaneously satisfy both  $L_F$  and  $L_C$ , which  
 1238 requires recognizing potentially crossing pairs, as in the sequence “ $\text{store}[f]$  entry[c]  $\text{load}[f]$  exit[c]”  
 1239 (where c denotes a calling context). However, it has been shown that precise interleaved Dyck-  
 1240 reachability is undecidable [60]. In practice, over-approximations are obtained by regularizing  
 1241  $L_C$  [71] and/or  $L_F$  [25, 74]. In contrast, in the next subsections we show that CUT-SHORTCUT can  
 1242 be formulated as a CFL-reachability problem without introducing a second Dyck language over  
 1243 calling contexts. Instead, its semantics can be expressed within a *single Dyck- $k$*  CFL,  $L_{csc}$ , which can  
 1244 be seen as a modification of  $L_F$ . Thus, it can be solved by the same formal model that recognizes  
 1245  $L_F$ , without requiring additional computational power.

1246 Our formulation of  $L_{csc}$  is presented as a two-step transformation from the baseline  $L_F$ . Step One  
 1247 addresses the local-flow pattern, the field-access pattern, and the *cut* mechanism of the container-  
 1248 access pattern, which can be expressed entirely by refining edges in the PAG (i.e., modifying  
 1249 terminals of the language). This processing does not require points-to information, and can therefore  
 1250 be performed prior to the CFL-reachability solving procedure. Step Two, in contrast, handles the  
 1251 *shortcut* mechanism of the container-access pattern, which depends on querying alias information  
 1252 between variables that may reference container objects. Its formulation must therefore be carried  
 1253 out on-the-fly, alongside the pointer analysis. To capture this approach, we extend the  $L_F$  grammar  
 1254 with additional host-tracking nonterminals. Splitting the formulation of CUT-SHORTCUT in these  
 1255 two steps cleanly separates the points-to-independent pre-processing step from the grammar  
 1256 extension, making the correctness clearer and the complexity overhead easier to analyze.

## 1257 6.2 Formulating CUT-SHORTCUT, Step One

1258 Assuming a pre-built call graph, we can express the local-flow pattern, the field-access pattern,  
 1259 and the *cut* mechanism of the container-access pattern in CUT-SHORTCUT using the rules shown  
 1260 in Figure 16. These rules derive binary relations over  $\mathbb{V} \times \mathbb{V}$  that represent imprecise assignments  
 1261 (which identify PAG edges to be suppressed) and precise shortcut assignments. Crucially, this  
 1262 derivation is independent of any points-to information, and can thus be performed before solving  
 1263 CFL-reachability on the PAG.

1264 Given the binary relations derived in Figure 16, along with the edge relations defined for  $L_F$  (also  
 1265 on  $\mathbb{V} \times \mathbb{V}$ ) in Figure 15, we now define the new PAG edge relations used in  $L_{csc}$  by applying set  
 1266 subtraction and union as follows (we use subscript  $F$  to denote relations from  $L_F$ , and the relations  
 1267 with no subscripts, such as  $\text{cutStore}[f]$ , are as computed by the rules in Figure 16):

$$\begin{aligned}
 \text{new}_{csc} &:= \text{new}_F \\
 \text{assign}_{csc} &:= (\text{assign}_F \setminus \text{cutReturn}) \cup \text{shortcutAssign} \\
 \text{store}[f]_{csc} &:= (\text{store}[f]_F \setminus \text{cutStore}[f]) \cup \text{shortcutStore}[f] \\
 \text{load}[f]_{csc} &:= \text{load}[f]_F \cup \text{shortcutLoad}[f]
 \end{aligned} \tag{2}$$

## FIELD-ACCESS PATTERN

$$\frac{\begin{array}{c} \mathbf{i: x.f = y} \quad i \in m \quad \text{param}_{k_1}^m = x \\ \text{param}_{k_2}^m = y \quad \text{def}_x = \emptyset \quad \text{def}_y = \emptyset \end{array}}{y \text{ cutStore}[f] x} [\text{L}_{\text{csc}}\text{-CUTSTORE}]$$

$$\frac{\langle base, f, from \rangle \in tempStore \quad base, from \in m}{((def_{base} \neq \emptyset) \vee (def_{from} \neq \emptyset) \vee (\#k : param_k^m = base) \vee (\#k : param_k^m = from)) \quad [L_{csc}\text{-}SHORTCUTSTORE]} \quad from\ shortcutStore[f]\ base$$

$$\frac{i : x = y.f \quad i \in m \quad \text{ret}^m = x \quad y = \text{param}_k^m \\ |\text{def}_x| = 1 \quad \text{def}_y = \emptyset \quad j \rightarrow \text{call } m}{\text{ret}^m \text{ cutReturn LHS}^j \quad \langle x, y, f \rangle \in \text{tempLoad}} \quad [L_{csc}\text{-CUTLOAD}]$$

$$\frac{\langle to, base, f \rangle \in \text{tempLoad} \quad base \leftarrow \text{arg}_k^j}{\text{arg}_k^j \text{ shortcutLoad}[f] \text{ LHS}^j} \quad [L_{csc}\text{-SHORTCUTLOAD}]$$

## LOCAL-FLOW PATTERN

$$\frac{\langle m, \_ \rangle \rightarrowtail \text{ret}^m \quad i \rightarrow \text{call } m}{\text{ret}^m \text{ cutReturn LHS}^i} [L_{\text{csc}}\text{-CUTLFLOW}] \qquad \frac{i \rightarrow \text{call } m \quad \langle m, k \rangle \rightarrowtail \text{ret}^m}{\arg^i_k \text{ shortcutAssign LHS}^i} [L_{\text{csc}}\text{-SHORTCUTLFLOW}]$$

## CONTAINER-ACCESS PATTERN (CUT)

$$\frac{\langle m, \_ \rangle \in \text{CONTEXIT} \quad i \rightarrow_{call} m}{\text{ret}^m \text{ cutReturn } \text{LHS}^i} [L_{csc}\text{-CUTCONT}]$$

Fig. 16. Rules for deriving the binary relations for imprecise assignments and precise shortcut assignments on the PAG, for the field-access pattern, local-flow pattern, and the *cut* mechanism of the container-access pattern. For the field-access pattern,  $[L_{csc}\text{-CUTSTORE}]$  and  $[L_{csc}\text{-SHORTCUTSTORE}]$  are adapted from  $[\text{CUTSTORE}]$  and  $[\text{SHORTCUTSTORE}]$  in Figure 8, with modifications to operate on the PAG rather than on PFG edges; similarly,  $[L_{csc}\text{-CUTLOAD}]$  and  $[L_{csc}\text{-SHORTCUTLOAD}]$  are adapted from Figure 9. For the local-flow pattern,  $[L_{csc}\text{-CUTLFLOW}]$  and  $[L_{csc}\text{-SHORTCUTLFLOW}]$  are adapted from their counterparts in Figure 14. Other rules from Figures 8, 9 and 14 that do not involve PAG edges remain valid under the CFL-reachability formulation and are omitted here for brevity. Finally, for the container-access pattern,  $[L_{csc}\text{-CUTCONT}]$  corresponds to  $[\text{CUTCONT}]$  in Figure 11; the remaining rules for this pattern will be reformulated in the next step (Section 6.3).

Step One of the transformation from  $L_F$  to  $L_{csc}$  thus involves the following substeps:

- i) Generate the relations for the unmodified PAG via Figure 15
  - ii) Generate the relations for the three CUT-SHORTCUT patterns via Figure 16
  - iii) Combine the information computed in the previous substeps via rule set (2)
  - iv) Define a new *single Dyck- $k$*  language  $L_{csc}$  as follows, based on the rules defined and edges computed in the previous steps:

$$\begin{aligned}
\text{flowsTo}_{csc} &\longrightarrow \text{new}_{csc} \text{ flows}_{csc}^* \\
\text{flows}_{csc} &\longrightarrow \text{assign}_{csc} \mid \text{store}[f]_{csc} \text{ alias}_{csc} \text{ load}[f]_{csc} \mid \text{shortcutCont} \\
\text{alias}_{csc} &\longrightarrow \underline{\text{flows}_{csc}} \text{ flows}_{csc} \\
\underline{\text{flowsTo}_{csc}} &\longrightarrow \underline{\text{flows}_{csc}}^* \text{ new}_{csc} \\
\underline{\text{flows}_{csc}} &\longrightarrow \underline{\text{assign}_{csc}} \mid \underline{\text{load}[f]_{csc}} \text{ alias}_{csc} \underline{\text{store}[f]_{csc}} \mid \underline{\text{shortcutCont}}
\end{aligned} \tag{3}$$

The definition of non-terminal *shortcutCont* is deferred to Step Two (Section 6.3).

The points-to relations of interest are the solution to the  $\text{flowsTo}_{\text{esc}}$ -reachability problem in the modified PAG, with respect to grammar (3) augmented by a grammar defined in Section 6.3.

Similar to  $L_F$ , in which *flowsTo*-paths represent standard points-to relations. In  $L_{csc}$ , a path  $o \text{ } flowsTo}_{csc} v$  indicates that variable  $v$  may point to heap object  $o$ . The structure of  $L_{csc}$  closely mirrors that of  $L_F$ , except that: (1) each edge relation from  $L_F$  has been replaced with its *csc*-refined version; and (2) a new non-terminal, *shortcutCont*, has been added to account for the *shortcut* mechanism of the container-access pattern, which we will elaborate on in Section 6.3.

```

1324 1 class Carton {
1325 2   Item item;
1326 3   void setItem(Item item){
1327 4     this.item = item;
1328 5   }
1329 6   Item getItem() {
1330 7     Item r = this.item;
1331 8     return r;
1332 9   }
1333 10 }
1334 11 class Item {}
1335 12
1336 13 // usage code
1337 14 void main() {
1338 15   Carton c1 = new Carton(); // o15
1339 16   Item item1 = new Item(); // o16
1340 17   c1.setItem(item1);
1341 18   Item result1 = c1.getItem();
1342 19
1343 20   Carton c2 = new Carton(); // o20
1344 21   Item item2 = new Item(); // o21
1345 22   c2.setItem(item2);
1346 23   Item result2 = c2.getItem();
1347 24 }

```

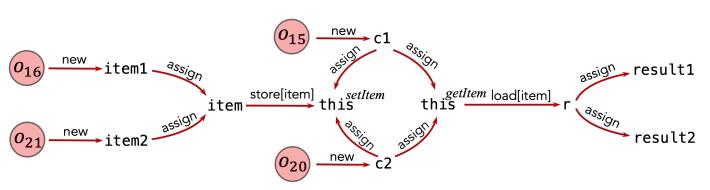
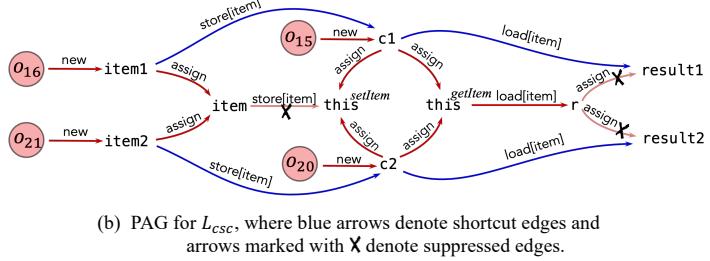
(a) PAG for  $L_F$ , constructed for the example shown on the left.(b) PAG for  $L_{csc}$ , where blue arrows denote shortcut edges and arrows marked with X denote suppressed edges.

Fig. 17. An example illustrating the CFL-reachability formulation of CUT-SHORTCUT. Subscripts of edge labels (i.e.,  $F$  in (a),  $csc$  in (b)) are omitted for brevity.

To concretize this formulation, Figure 17 revisits the same motivating example from Section 2 and demonstrates how CUT-SHORTCUT is captured under the CFL-reachability framework up to this point. In Figure 17(a), the  $store[item]$  edge from  $item$  to  $this^{setItem}$  brings a merged flow from both  $item1$  and  $item2$  into  $this^{setItem}$ . Similarly, the  $assign$  edges from  $r$  to  $result1$  and  $result2$  bring merged flows from  $this^{getItem}$  to both  $result$  variables. As a result, there are two  $L_F$ -paths ending at  $result1$ :

$$\begin{aligned}
 & o_{16} \xrightarrow{\text{new}} item1 \xrightarrow{\text{assign}} item \xrightarrow{\text{store[item]}} this^{setItem} \xrightarrow{\text{alias}} this^{getItem} \xrightarrow{\text{load[item]}} r \xrightarrow{\text{assign}} result1 \\
 & o_{21} \xrightarrow{\text{new}} item2 \xrightarrow{\text{assign}} item \xrightarrow{\text{store[item]}} this^{setItem} \xrightarrow{\text{alias}} this^{getItem} \xrightarrow{\text{load[item]}} r \xrightarrow{\text{assign}} result1
 \end{aligned}$$

where the *alias*-path between  $this^{setItem}$  and  $this^{getItem}$  is given by:

$$this^{setItem} \xrightarrow{\text{assign}} c1 \xrightarrow{\text{new}} o_{15} \xrightarrow{\text{new}} c1 \xrightarrow{\text{assign}} this^{getItem}$$

These paths imply that  $result1$  points to both  $o_{16}$  and  $o_{21}$ —an imprecise result. A similar imprecision occurs for  $result2$ .

In contrast, Figure 17(b) shows the PAG for the CFL formulation, where three imprecise edges are suppressed (crossed arrows), and four shortcut edges (blue arrows) are added, relative to Figure 17(a). As a result, only one  $L_{csc}$ -path leads to  $result1$ :

$$o_{16} \xrightarrow{\text{new}} item1 \xrightarrow{\text{store[item]}} c1 \xrightarrow{\text{load[item]}} result1$$

Thus,  $result1$  precisely points to  $o_{16}$ . Similarly, only one  $L_{csc}$ -path leads to  $result2$ , giving a precise analysis result. Note that all three suppressed edges in Figure 17(b) are necessary. For example, if the  $store[item]$  edge from  $item$  to  $this^{setItem}$  was not suppressed, an imprecise *flowsTo*-path from  $o_{21}$  to  $result1$  would be present, namely,

$$o_{21} \xrightarrow{\text{new}} item2 \xrightarrow{\text{store[item]}} this^{setItem} \xrightarrow{\text{alias}} c1 \xrightarrow{\text{load[item]}} result1$$

1373 If the assign edge from  $r$  to  $\text{result1}$  was not suppressed, then the PAG would contain the path

$$o_{21} \xrightarrow{\text{new}} \text{item2} \xrightarrow{\text{store[item]}} c_2 \xrightarrow{\text{assign}} \text{this} \xrightarrow{\text{getItem}} r \xrightarrow{\text{load[item]}} r \xrightarrow{\text{assign}} \text{result1}$$

1376 Similarly, not suppressing the assign edge to  $\text{result2}$  causes  $\text{result2}$  to imprecisely point to  $o_{16}$ .

### 1378 6.3 Formulating CUT-SHORTCUT, Step Two

1379 In this step, we formalize the *shortcut* mechanism of the container-access pattern. Because this  
 1380 mechanism depends on aliasing information for variables that may reference container objects,  
 1381 it must be integrated into the pointer analysis itself and executed in an on-the-fly manner. In a  
 1382 pointer flow graph (PFG)-based formulation, this on-the-fly behavior is realized by propagating  
 1383 hosts—our abstraction of container instances—along PFG edges during pointer analysis, in a manner  
 1384 similar to heap-object propagation (as discussed in Section 4.3). To express this behavior in terms  
 1385 of CFL-reachability, we introduce a new non-terminal *hostFlows*, which captures the paths that a  
 1386 host node may take to reach a variable—analogous to how the non-terminal *flows* in  $L_F$  captures  
 1387 the paths that an object may take to reach a variable. Non-terminal *hostFlows* is used to define a  
 1388 central non-terminal *shortcutCont*, which characterizes shortcut paths that objects may take to  
 1389 reach variables. Before presenting the full grammar for them, we first introduce the additional PAG  
 1390 edge-construction rules shown in Figure 18, which are used in the grammar of these non-terminals.  
 1391

$$\begin{array}{c} \frac{i : x = \text{new } A() \quad A <: \text{Collection} \vee \text{Map}}{h_i \xrightarrow{\text{newHost}} x} [L_{\text{csc}}\text{-NewHost}] \\[10pt] \frac{i \rightarrow \text{call } m \quad \langle m, k, \_ \rangle \in \text{CONTENTR}}{\arg_k^i \xrightarrow{\text{source}} \arg_0^i} [L_{\text{csc}}\text{-HOSTSOURCE}] \\[10pt] \frac{i \rightarrow \text{call } m \quad \langle m, \_ \rangle \in \text{CONTPROP}}{\arg_0^i \xrightarrow{\text{propagate}} \text{LHS}^i} [L_{\text{csc}}\text{-HOSTPROP}] \\[10pt] \frac{i \rightarrow \text{call } m \quad \langle m, k, \_ \rangle \in \text{CONTBULKENTR}}{\arg_k^i \xrightarrow{\text{hostDepend}} \arg_0^i} [L_{\text{csc}}\text{-HOSTDEP}] \\[10pt] \frac{i \rightarrow \text{call } m \quad \langle m, \_ \rangle \in \text{CONTEXIT}}{\arg_0^i \xrightarrow{\text{target}} \text{LHS}^i} [L_{\text{csc}}\text{-HOSTTARGET}] \\[10pt] \frac{i \rightarrow \text{call } m \quad \langle m, \_ \rangle \in \text{CONTPROP}}{\text{ret}^m \text{ PropExcept LHS}^i} [L_{\text{csc}}\text{-PROPEX}] \end{array}$$

1401 Fig. 18. Additional PAG edge-construction rules used to define the grammar of the non-terminals involved  
 1402 in the *shortcut* mechanism of the container-access pattern.  $[L_{\text{csc}}\text{-NewHost}]$  is adapted from  $[\text{ColHost}]$  and  
 1403  $[\text{MapHost}]$  in Figure 11; likewise,  $[L_{\text{csc}}\text{-HOSTDEP}]$  corresponds to  $[\text{HostBULKETR}]$ ;  $[L_{\text{csc}}\text{-HOSTSOURCE}]$  to  $[\text{HostSOURCE}]$ ;  
 1404  $[L_{\text{csc}}\text{-HOSTTARGET}]$  to  $[\text{HostTARGET}]$ ; and  $[L_{\text{csc}}\text{-HOSTPROP}]$  to  $[\text{HostPROP-I}]$ . Finally,  $[L_{\text{csc}}\text{-PROPEX}]$  captures the exceptional  
 1405 case excluded by the premise of  $[\text{HostPROP-II}]$ .

1406 The rules in Figure 18 encode<sup>9</sup> those in Figure 11. The first five rules directly introduce additional  
 1407 PAG edges. The last one,  $[L_{\text{csc}}\text{-ProEx}]$ , is used to define a new PAG edge type:  $\text{hostAssign} := \text{assign}_{\text{csc}} \setminus$   
 1408  $\text{PropExcept}$ , which consists of  $\text{assign}$  edges, excluding those captured by the  $\text{PropExcept}$  relation.  
 1409 With these new edges in place, we define the following non-terminals:

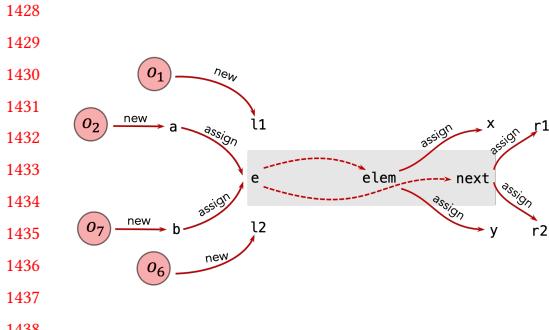
$$\begin{array}{l} \text{hostFlows} \longrightarrow \text{hostAssign} \mid \text{store}[f]_{\text{csc}} \text{ alias}_{\text{csc}} \text{ load}[f]_{\text{csc}} \mid \text{propagate} \mid \text{shortcutCont} \\ \text{hostFlows} \longrightarrow \overline{\text{hostAssign}} \mid \overline{\text{load}[f]_{\text{csc}}} \text{ alias}_{\text{csc}} \overline{\text{store}[f]_{\text{csc}}} \mid \overline{\text{propagate}} \mid \overline{\text{shortcutCont}} \\ \text{sourceOfHost} \longrightarrow \text{source} (\overline{\text{hostFlows}} \mid \text{hostDep})^* \overline{\text{newHost}} \\ \text{sourceOfHost} \longrightarrow \overline{\text{newHost}} (\text{hostDep} \mid \text{hostFlows})^* \overline{\text{source}} \\ \text{targetOfHost} \longrightarrow \overline{\text{target}} \overline{\text{hostFlows}}^* \overline{\text{newHost}} \\ \text{targetOfHost} \longrightarrow \overline{\text{newHost}} \overline{\text{hostFlows}}^* \overline{\text{target}} \end{array} \tag{4}$$

1419 <sup>9</sup>For clarity, we do not distinguish the three host kinds in this CFL formulation, although doing so is straightforward: each  
 1420 rule in Figure 18 could be instantiated once per kind. We omit this distinction here to avoid unnecessary repetition.

```

1422 1 ArrayList l1 = new ArrayList(); //O1, h1
1423 2 Object a = new Object(); //O2
1424 3 l1.add(a); //⟨add, l1, Coll⟩ ∈ CONTENTER
1425 4 Object x = l1.get(0); //⟨get, Coll⟩ ∈ CONTEXIT
1426 5
1427 6 ArrayList l2 = new ArrayList(); //O6, h6
1428 7 Object b = new Object(); //O7
1429 8 l2.add(b); //⟨add, l2, Coll⟩ ∈ CONTENTER
1430 9 Object y = l2.get(0); //⟨get, Coll⟩ ∈ CONTEXIT

```

(a) PAG fragment of  $L_F$  for the above code example.

```

10 // retrieve from container-dependent objects
11 Iterator it1 = l1.iterator(); //⟨iterator, l1⟩ ∈ CONTPROP
12 Object r1 = it1.next(); //⟨next, l1⟩ ∈ CONTEXIT
13 Iterator it2 = l2.iterator(); //⟨iterator, l2⟩ ∈ CONTPROP
14 Object r2 = it2.next(); //⟨next, l2⟩ ∈ CONTEXIT

```

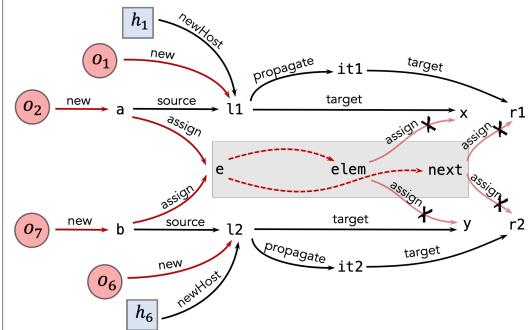
(b) PAG fragment of  $L_{csc}$ , where black arrows denote new edges, and arrows marked with X denote suppressed edges.

Fig. 19. An example to illustrate the CFL-reachability formulation of the container-access pattern. Black arrows in (b) represent additional PAG edges introduced via the rules in Figure 18. Dashed arrows indicate the presence of *flows* paths in (a) and *flows<sub>csc</sub>* paths in (b). Gray regions mark boundaries of JDK library internals. The subscripts *F* (in (a)) and *csc* (in (b)) on edge labels are omitted. For brevity, *AssignHost* edges in (b) are not shown, because they are not relevant to this example.

Intuitively, *hostFlows* captures the paths that a host node can take to reach host-related variables, such as a list or its iterator. The non-terminal *sourceOfHost* defines valid paths from a SOURCE of host  $h$  to the host node that represents  $h$ , while *targetOfHost* defines paths from a TARGET of  $h$  to  $h$ . A subtle point: in *sourceOfHost*-paths, an edge labeled *hostDep* is allowed, but such edges do not appear in *targetOfHost*-paths. This asymmetry reflects the preorder  $\triangleleft$  defined in Section 4.3, capturing SOURCE dependencies between hosts. Intuitively, the *hostDep* edges allow the SOURCES of host  $h$  to "jump" into the *hostFlows* of another host  $h'$  when  $h \triangleleft h'$ .

Finally, the *shortcutCont* non-terminal is defined by concatenating a *sourceOfHost* path with a *targetOfHost* path, matching the SOURCES and TARGETS of the same host:

$$\begin{aligned} \text{shortcutCont} &\longrightarrow \text{sourceOfHost } \text{targetOfHost} \\ \text{shortcutCont} &\longrightarrow \text{targetOfHost } \text{sourceOfHost} \end{aligned} \quad (5)$$

In Figure 19, we revisit the example introduced earlier in Section 3.3 to illustrate the CFL-reachability formulation of the container-access pattern. In Figure 19 (a), under a context-insensitive pointer analysis, there exist two  $L_F$ -paths leading to  $x$ . (Note that the *flows*-path from  $e$  to  $elem$  models the internal behavior of the *ArrayList* class.)

$$o_2 \xrightarrow{\text{new}} a \xrightarrow{\text{assign}} e \xrightarrow{\text{flows}} elem \xrightarrow{\text{assign}} x \quad \text{and} \quad o_7 \xrightarrow{\text{new}} b \xrightarrow{\text{assign}} e \xrightarrow{\text{flows}} elem \xrightarrow{\text{assign}} x$$

These paths give an imprecise result:  $x$  is considered to point to both  $o_2$  and  $o_7$ . Similarly,  $y$  is imprecisely inferred to point to both objects. In addition, for variable  $r1$ , retrieved via the iterator of  $l1$ , the following two  $L_F$ -paths exist:

$$o_2 \xrightarrow{\text{new}} a \xrightarrow{\text{assign}} e \xrightarrow{\text{flows}} next \xrightarrow{\text{assign}} r1 \quad \text{and} \quad o_7 \xrightarrow{\text{new}} b \xrightarrow{\text{assign}} e \xrightarrow{\text{flows}} next \xrightarrow{\text{assign}} r1$$

which similarly yield imprecise results for  $r_1$  (and likewise for  $r_2$ ). In contrast, under CUT-SHORTCUT, as shown in Figure 19 (b), imprecise edges are suppressed and additional host-handling edges are introduced. Now only one  $L_{csc}$ -path exists for  $x$ :

$$o_2 \xrightarrow{\text{new}} a \xrightarrow{\text{shortcutCont}} x$$

where the *shortcutCont* path from  $a$  to  $x$  expands to

$$a \xrightarrow{\text{sourceOfHost}} h_1 \xrightarrow{\text{targetOfHost}} x$$

which in turn expands to

$$a \xrightarrow{\text{source}} 11 \xrightarrow{\overline{\text{newHost}}} h_1 \xrightarrow{\text{newHost}} 11 \xrightarrow{\text{target}} x$$

Likewise, only one  $L_{csc}$ -path leads to  $y$ , yielding precise results for the points-to sets of both  $x$  and  $y$ . For iterators, the only  $L_{csc}$ -path for  $r_1$  expands to

$$o_2 \xrightarrow{\text{new}} a \xrightarrow{\text{source}} 11 \xrightarrow{\overline{\text{newHost}}} h_1 \xrightarrow{\text{newHost}} 11 \xrightarrow{\text{propagate}} \text{it1} \xrightarrow{\text{target}} r_1$$

This path, together with the absence of any  $L_{csc}$ -path to  $r_1$  from any other variable, gives the precise result that  $r_1$  points only to  $o_2$ . Similarly, the only  $L_{csc}$ -path for  $r_2$  yields the precise result that  $r_2$  points to  $o_7$ .

## 6.4 Complexity Discussion

From the perspective of CFL-reachability, the core insight of CUT-SHORTCUT remains consistent: compared to  $L_F$ , certain edges are suppressed in the PAG for  $L_{csc}$ —ensuring that some imprecise  $L_F$ -paths are no longer present—while new edges and non-terminals are introduced to enable the *shortcut* mechanism.

We emphasize that because the desired language  $L_{csc}$  for our language-reachability problem is expressed directly *using a context-free grammar* (i.e., grammars (3), (4) and (5)), it follows that  $L_{csc}$  is a context-free language. In contrast, if it were an interleaved-Dyck language, we would not be able to express it with a context-free grammar. Crucially,  $L_{csc}$  avoids introducing a second Dyck- $k$  language for matching method entries and exits, so that reachability queries under  $L_{csc}$  maintain comparable complexity to those under  $L_F$ , despite its improved precision.

Concretely, we will first show in Lemma 6.1 that transforming  $L_F$  to  $L_{csc}$  incurs only linear growth in grammar and graph size, and prove in Theorem 6.2 that  $L_{csc}$  preserves the asymptotic complexity of  $L_F$  under a black-box CFL solver. Finally, we compare this overhead with  $k$ -CFA [65], a common regularization for context-sensitive analysis: while  $k$ -CFA bounds the context length, it can still suffer from a potential cost blow up, where the blow-up factor is exponential in  $k$ .

**LEMMA 6.1 (TRANSFORMATION OVERHEAD).** *Moving from  $L_F$  to  $L_{csc}$  introduces only linear growth in the size of the underlying graph (nodes and edges) and the grammar.*

**PROOF.** Let  $|\mathbb{F}|$  denote the number of fields and  $|\mathbb{H}|$  the number of hosts. We analyze the effect on graph and grammar through Step One and Step Two of our transformation:

- (1) *Graph Nodes.* Step One adds no new nodes to the underlying PAG. Step Two introduces one additional node per host (i.e., each container allocation site), contributing  $|\mathbb{H}|$  nodes. If the original PAG has  $n$  nodes, then the new PAG has  $n + |\mathbb{H}|$  nodes. Because  $|\mathbb{H}|$  is itself bounded by the number of allocation sites, hence by  $O(n)$ , the total number of nodes is still  $O(n)$ .

- 1520 (2) *Graph Edges.* Step One refines the PAG of  $L_F$  by suppressing some edges and introducing  
 1521 shortcut edges. As noted in the footnote of Section 4.2.1, we impose a cap on the handling  
 1522 of the field-access and local-flow patterns, ensuring that the number of shortcut edges  
 1523 introduced is bounded by a small constant times the number of call sites in the program.  
 1524 Because the number of call sites is already  $O(m)$ , where  $m$  denotes the number of edges in  
 1525 the PAG of  $L_F$ , Step One contributes only  $O(m)$  edges overall. Step Two introduces a small  
 1526 number of new edge kinds (e.g., newHost, hostDep, etc.), each generated one-for-one from  
 1527 either allocation sites or container-API call sites (cf. Figure 18). Because both allocation  
 1528 sites and call sites are bounded by  $O(m)$ , Step Two likewise adds only linearly many edges.  
 1529 Hence, the total number of edges in the transformed PAG remains  $O(m)$ .
- 1530 (3) *Grammar.* The baseline grammar  $\mathcal{G}_F$  has size  $O(|\mathbb{F}|)$ , because it contains one family of  
 1531 productions parameterized by fields. Step One rewrites  $\mathcal{G}_F$  by adding subscripts to ex-  
 1532 isting nonterminals and introducing a placeholder nonterminal *shortcutCont*. Step Two  
 1533 extends the grammar with a constant number of new nonterminals (*hostFlows*, *sourceOfHost*,  
 1534 *targetOfHost*, *shortcutCont*) and a fixed schema of productions—these reuse existing termi-  
 1535 nals and introduce one additional  $\Theta(|\mathbb{F}|)$  family of field-parameterized productions inside  
 1536 *hostFlows*, of the form *store[f] alias<sub>csc</sub> load[f]* (and its symmetric variant). Instantiating such  
 1537 a schema for each  $f \in \mathbb{F}$  adds at most  $2|\mathbb{F}|$  concrete productions. Formally, the grammar size  
 1538 grows from  $|\mathcal{G}_F| = O(|\mathbb{F}|)$  to  $|\mathcal{G}_{csc}| = O(|\mathbb{F}|) + c + 2 \cdot O(|\mathbb{F}|) = O(|\mathbb{F}|)$  where  $c$  is a constant.  
 1539 Crucially, no cross-product over fields and hosts arises, so there is no  $O(|\mathbb{F}| \cdot |\mathbb{H}|)$  blow-up.

1540 To sum up: if the original  $L_F$  instance has graph size  $(n, m)$  and grammar size  $|\mathcal{G}_F|$ , then the  
 1541 transformed  $L_{csc}$  instance has size  $(O(n), O(m))$  and  $O(|\mathcal{G}_F|)$  respectively—at most linear growth.  
 1542  $\square$

1543 **THEOREM 6.2.**  *$L_{csc}$  preserves the asymptotic complexity of the baseline  $L_F$ .*

1544 **PROOF.** Let  $T(n, m, |\mathcal{G}|)$  and  $S(n, m, |\mathcal{G}|)$  denote the abstract time and space cost functions of a  
 1545 single-stack black-box CFL-reachability solver on a graph with  $n$  nodes,  $m$  edges, and grammar  
 1546 size  $|\mathcal{G}|$ . By Lemma 6.1, if the original  $L_F$  instance has graph node, graph edge, and grammar size  
 1547  $(n, m, |\mathcal{G}|)$ , then the transformed  $L_{csc}$  instance has size  $(O(n), O(m), O(|\mathcal{G}|))$ . Hence the cost of  
 1548 solving  $L_{csc}$  can be expressed as

1549 
$$\text{Time}_{csc} = O(|\text{PAG}_F|) + T(O(n), O(m), O(|\mathcal{G}|)), \quad \text{Space}_{csc} = S(O(n), O(m), O(|\mathcal{G}|)).$$

1550 The additive term  $O(|\text{PAG}_F|)$  in  $\text{Time}_{csc}$  corresponds to the cost of refining the PAG in Step One,  
 1551 which is proportional to the size of the original PAG size (i.e.,  $L_F$ ).  $\square$

1552 *Compare with k-CFA-based context-sensitivity.* We now compare this complexity with context-  
 1553 sensitivity-based techniques. As noted earlier, fully context- and field-sensitive pointer analysis  
 1554 is an interleaved-Dyck-reachability problem, and hence undecidable. A common regularization  
 1555 is to use *k*-CFA [65], which bounds method contexts by call strings of length  $k$ . This approach  
 1556 transforms the underlying PAG into a product graph: each node is paired with a call string of length  
 1557  $\leq k$ , and each original edge  $(u, v)$  is lifted to edges between  $(u, c_1)$  and  $(v, c_2)$ , where  $c_1$  and  $c_2$  are  
 1558 contexts (which may coincide or differ depending on the edge kind). The resulting graph is thus a  
 1559 blown-up version of the original PAG. Let  $\gamma$  be the number of call sites and  $\mathbb{C}$  the set of allowed  
 1560 contexts. A standard bound is

1561 
$$|\mathbb{C}| \leq 1 + \gamma + \gamma^2 + \cdots + \gamma^k = O(\gamma^k).$$

1562 If the original  $L_F$  instance has size  $(n, m, |\mathcal{G}|)$ , then the product graph under *k*-CFA has at most  
 1563  $n \cdot O(\gamma^k)$  nodes and  $O(m \cdot \gamma^k)$  edges, while the grammar size remains unchanged. Hence, for the same

1569 abstract single-stack CFL solver that handles  $L_F$  and  $L_{csc}$ , solving the  $k$ -CFA-based context-sensitive  
 1570 and field-sensitive pointer analysis requires

$$1571 \text{Time}_{k\text{-CFA}} = T(n \cdot O(\gamma^k), O(m \cdot \gamma^k), |\mathcal{G}|), \quad \text{Space}_{k\text{-CFA}} = S(n \cdot O(\gamma^k), O(m \cdot \gamma^k), |\mathcal{G}|).$$

1573 Thus,  $k$ -CFA can be solved by the same black-box CFL solver that solves  $L_F$  and  $L_{csc}$ , but at the cost  
 1574 of a potential blowup in the size of the analysis graph that is exponential in  $k$ . In practice,  $k$ -CFA is  
 1575 inefficient/less-scalable [36] than context insensitivity or CUT-SHORTCUT.

## 1576 7 CUT-SHORTCUT<sup>S</sup>

1577 In this section, we present CUT-SHORTCUT<sup>S</sup>, an extension of CUT-SHORTCUT that addresses imprecision caused by pervasive use of Java streams, by modeling various stream pipeline operations  
 1578 within pointer analysis.

### 1581 7.1 Motivation

1583 *Challenge Posed by Java Streams.* Pointer analysis for Java has seen substantial research  
 1584 progress in improving precision, efficiency, and scalability. However, most existing whole-program  
 1585 pointer analyses are built around core Java features and evaluated on Java 6 benchmarks [22, 23, 28–  
 1586 30, 32, 37, 40, 41, 46, 47, 70, 76–78].

1587 However, according to a report by New Relic [1], the most-used Java version in industry as of  
 1588 2023 is Java 11, with Java 8 being a close second. In practice, applying existing pointer-analysis  
 1589 approaches to Java 8 (or higher-version) programs will face significant challenges [18]: the precision,  
 1590 scalability, or even soundness may not be preserved when hard-to-analyze language features are  
 1591 used. In the past few years, effort has been made towards more sound and precise handling of those  
 1592 features, including Java reflection [42, 68], invoke-dynamic along with Java lambda expressions [18],  
 1593 and native codes [87]. That said, Java streams, a widely used feature introduced in Java 8, are still  
 1594 an obstacle for precise whole-program pointer analysis.

1595 Java streams are a pervasively used language feature in modern Java programs, and expose a set  
 1596 of APIs for developers to process collections of data using functional-style operations, to generate  
 1597 elegant code. In programs, streams are used as pipelines, which typically contain a source, zero  
 1598 or multiple intermediate operations, and one terminal operation. If a typical context-insensitive  
 1599 Andersen-style analysis [2] is used, the wide use of stream pipelines can pose a significant challenge  
 1600 to the precision, because the analysis will mix together elements (i.e., heap objects) that pass through  
 1601 different pipelines.

1602 *7.1.2 A Motivating Example.* Figure 20(a) shows a motivating example. The class Stud has two  
 1603 String fields: name and id. In main(), four Stud objects are created and assigned to variables a1 to  
 1604 a4; hence we observe  $pt(a_i) = \{o_i\}$  for each  $a_i$ . Figure 20(b) illustrates how these objects propagate  
 1605 in the first stream pipeline at runtime. Pointers a1 and a2 are passed as arguments for method  
 1606 invocation of Stream.of, which generates a new instance of Stream type that takes a1 and a2 as  
 1607 input elements. A Stream.filter operation then drops students whose name.length() <= 4, so only  
 1608 o1 survives. Then, an invocation of Stream.findFirst would find the first object in that pipeline,  
 1609 and wrap it into an Optional type (a wrapper type in Java to represent potentially null objects),  
 1610 which is finally unwrapped through Optional.get, with o1 assigned to variable r1. Therefore, r1  
 1611 would only point to o1 but not o2 after dynamic execution. Similarly, r2 would only point to o3.

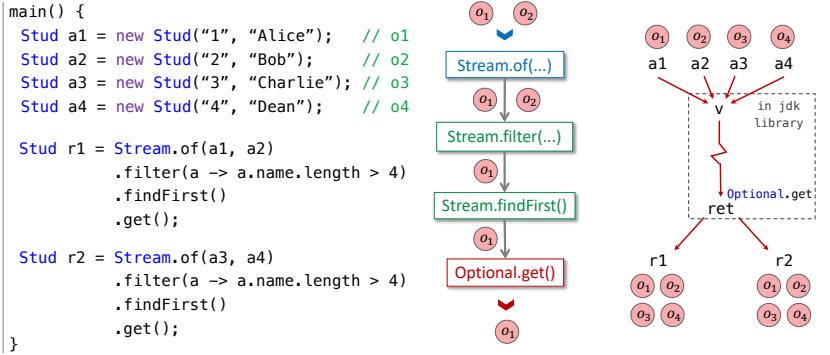
1612 *Static Over-Approximation Goal.* The exact behavior of Stream.filter is determined by the  
 1613 lambda expression<sup>10</sup> used as the argument at this call site. The semantics of an arbitrary lambda

1615 <sup>10</sup>In Java, each lambda expression is a reference to an object with a type annotated as FunctionalInterface, indicating that  
 1616 it is a “functional object.” More details of Java lambda expressions will be discussed in Section 7.4.

```

1618 Stud {
1619     String name, id;
1620     Stud(String name,
1621           String id) {
1622         this.name = name;
1623         this.id = id;
1624     }
1625     String getName() {
1626         return this.name;
1627     }
1628     String getID() {
1629         return this.id;
1630     }
1631 }
```

(a) A stream use example.



(b) Dynamic behavior of the first pipeline instance.

(c) PFG fragment under context insensitivity

Fig. 20. A motivating example for Java streams.

expression is difficult to model statically, because it could function as any  $\text{Stud} \mapsto \text{Boolean}$  function in this case. However, we can still aim for a sound over-approximation: a reasonable static analysis result would be  $pt(r1) = \{o_1, o_2\}$  and  $pt(r2) = \{o_3, o_4\}$ , meaning that the analysis can tell the difference between the objects passed through different stream pipelines.

*Context Insensitivity.* Figure 20(c) shows the PFG fragment under context insensitivity. The arguments of the first `Stream.of` call-site, namely  $a_1$  and  $a_2$ , are passed to the parameter  $v$  of the callee (which is a JDK library method). Due to context-insensitivity, the analysis does not distinguish between the two `Stream.of` invocations; hence  $a_3$  and  $a_4$  are also passed to  $v$ , resulting in the merging of object flows from both call-sites. The merged flow then propagates through the library’s internal methods and returns to both  $r_1$  and  $r_2$ , yielding  $pt(r_1) = pt(r_2) = \{o_1, o_2, o_3, o_4\}$ .

*Limitations of Context Sensitivity.* While deep context sensitivity can, in theory, recover the desired precision, it comes at exponential cost. In practice, even for our motivating example—which consists of only two stream pipelines—selective context-sensitive analyses that target stream-related methods fail to yield precise results. For instance, selective 5call-4h (5 layers of call-site sensitivity with 4 layers of context-sensitive heap abstraction) is unable to distinguish the two pipelines, and increasing the context depth further leads to out-of-memory errors (exceeding 100GB). Similarly, with object sensitivity, selective 3obj-2h (3 layers of object sensitivity with 2 layers of context-sensitive heap abstraction) also fails to distinguish the pipelines, and increasing the context depth provides no additional precision benefits for this program.

*Our Approach.* Instead of relying on context sensitivity via method cloning, we adopt the CUT-SHORTCUT philosophy—suppressing imprecise PFG edges and introducing precise shortcut edges. We extend our container-access pattern to handle Java streams by modeling stream pipelines as containers. The key insight is that each stream pipeline’s outputs can only derive from its own inputs. In our container-access pattern, we abstract container instances as hosts and associate each with its input elements (SOURCES) and output locations (TARGETS), then suppress imprecise return edges and introduce shortcut edges from SOURCES to TARGETS of the same host. This idea naturally extends to streams: pipelines are abstracted as a new kind of host, with SOURCES and TARGETS identified based on modeled stream operations. This uniform treatment lets us reuse some of the existing analysis rules, rather than introducing a completely separate mechanism for streams. It also cleanly accounts for conversions between streams and containers: by treating pipelines as a new kind of host, distinct from those for collections or maps, such conversions become

1667 transformations between host kinds. At the same time, streams introduce challenges beyond those  
 1668 of standard containers. Pipeline operations such as `map` and `flatMap` transform element values rather  
 1669 than merely transferring references, and have no direct analog in collections. These higher-order  
 1670 operations apply user-supplied callbacks (lambdas/method references) to elements, which are not  
 1671 handled “for free” by container abstractions and require careful modeling. Thus, while our reuse of  
 1672 the container-access pattern provides a clean starting point, capturing the distinct semantics of  
 1673 streams operations ensures that this extension is a non-trivial generalization.

1674 In the following subsections, we define the scope and classification of stream operations (Section  
 1675 7.2) and describe our modeling strategies in detail (Sections 7.3–7.5).

## 1677 7.2 Classify Stream Operations

1678 The Java stream APIs consist of a set of publicly accessible classes and methods that developers  
 1679 can use to process elements in a pipeline style. Figure 21 provides an overview of the stream APIs,  
 1680 focusing specifically on operations relevant to pointer analysis. To the best of our knowledge,  
 1681 we are the first to systematize Java stream operations from the perspective of pointer analysis,  
 1682 highlighting and classifying the different ways in which these operations affect object flow.  
 1683

1684 As shown, Java defines four interfaces to abstract streams over different types. `IntStream`,  
 1685 `LongStream`, and `DoubleStream` represent streams over primitive types (`int`, `long`, and `double`, respectively), while the `Stream` interface handles reference types. Because pointer analysis is concerned  
 1686 only with heap objects—which are reference types—we restrict our analysis to streams over reference  
 1687 types. It is worth noting that some primitive-type stream interfaces provide operations to convert to  
 1688 reference-type streams, such as `IntStream.mapToObj`; hence such operations must also be modeled  
 1689 in our analysis. In addition to the `Stream` interface, other classes may interact with reference-type  
 1690 streams. One example is the `Optional` class. As shown in the code snippet in Figure 20(a), calling  
 1691 `Stream.findFirst` does not return an element in that stream pipeline directly; instead, it returns  
 1692 an `Optional` object that wraps the element, which is later retrievable via `Optional.get`. Another  
 1693 example is the `Collection` interface, which provides the method `Collection.stream` to create a  
 1694 stream instance whose input elements are the elements of a `Collection` instance. In general, aside  
 1695 from methods declared in the `Stream` and `Optional` classes, any method in the JDK library that  
 1696 creates and returns an object of type `Stream` or `Optional` should be considered in our handling,  
 1697 because it may participate in the life-cycle of a stream pipeline.  
 1698

1699 We refer to all such methods as stream pipeline-related operations (stream operations, for short),  
 1700 and classify them into three main categories, corresponding to the three stages of a stream pipeline  
 1701 life-cycle: *creation*, *process*, and *output*. Representative examples of each category are shown in  
 1702 Figure 21, along with their method signatures.<sup>11</sup> Beyond life-cycle stages, stream operations can  
 1703 also be classified based on how they are parameterized, which indicates how they behave—and in  
 1704 turn determines how our analysis deals with them. Accordingly, each of the three main categories—  
 1705 *creation*, *process*, and *output*—is further divided into two or three subcategories, to reflect such  
 1706 behavioral characteristics. This classification is visually represented in Figure 21, where the color  
 1707 indicates the life-cycle stage of each operation, and the depth of color encodes its behavioral  
 1708 characteristics. The three behavioral subcategories are defined as follows:

1709  
 1710 <sup>11</sup>Note that our classification is based on how operations affect object flow, and thus differs slightly from the Java docu-  
 1711 mentation, which uses the terms *source*, *intermediate* operation, and *terminal* operation. For example, `Stream.findFirst` is  
 1712 classified as a *terminal* operation in the official documentation, but we treat it as a *process* operation, instead of an *output*  
 1713 operation. This approach is warranted because the returned value is not a stream element but an `Optional` object that wraps  
 1714 it, requiring further unwrapping. Consequently, our analysis handles such operations similarly to *process* operations like  
 1715 `Stream.filter` or `Stream.sorted`, while those unwrapping operations, like `Optional.get`, are classified as *output* operations.

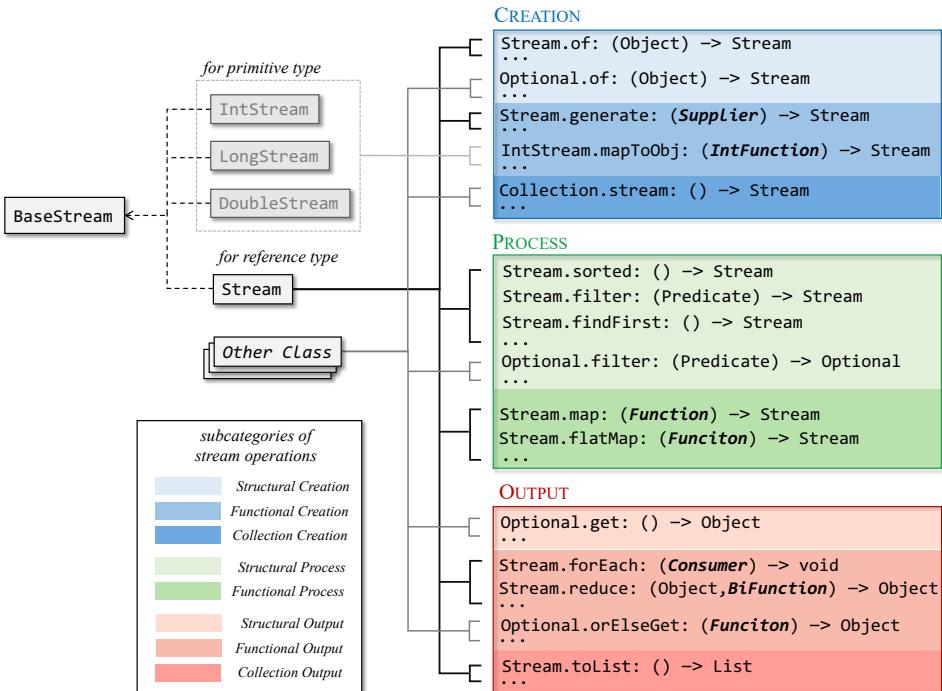


Fig. 21. Overview of pointer-analysis-related stream-pipeline operations

- (a) *Structural*. These operations have fixed semantics and do not take user-specified behavior (e.g., lambda expressions) as arguments. Examples include `Stream.of` (a *structural creation* operation), `Stream.findFirst` (a *structural process* operation), and `Optional.get` (a *structural output* operation). They are well-characterized and have predictable effect, and hence can be handled uniformly, either by specifying them in an input set already defined by the container-access pattern or by introducing a single new analysis rule. Details are given in Section 7.3.
- (b) *Functional*. These operations take functions (i.e., lambda expressions/method references) that encode how stream elements are processed or transformed as invocation arguments.<sup>12</sup> For example, `Stream.map` (a *functional process* operation) takes a `Function` argument (which represents the function to be applied to each pipeline element), and produces a new `Stream` object as its return value. The elements in the resulting `Stream` object depend on how the functional argument behaves. Similarly, `Stream.generate` and `Stream.forEach` are *functional creation* operation and *functional output* operation, respectively. Each *functional* operation has its unique semantics (regarding how the functional arguments may interact with stream elements) even within the same sub-category. For instance, `Stream.map` and `Stream.flatMap` are both *functional process* operations, but they affect stream elements in different ways. Consequently, we define a separate rule for each *functional* operation to model their different semantics. Section 7.4 details our handling of these operations, which includes integrating with an existing static analysis for Java's functional features [18].

<sup>12</sup>An exception to this category is `Stream.filter`, which, while technically functional (taking a function as its argument), only removes elements from the stream rather than producing new ones. Because our analysis focuses on tracking and distinguishing stream elements (not deletions), we classify it as a *structural process* operation, and handle it accordingly.

(c) *Collection-Interacting*. These operations bridge the Java stream and collection domains. For example, `Collection.stream` creates a `Stream` object from a `Collection` instance (hence a *collection creation* operation), and `Stream.toList` collects stream elements into a new `List` instance (hence a *collection output* operation). It is worth noting that *process* operations do not directly interact with collections—hence, there is no corresponding subcategory. The key to handling these operations is to model the transformations between streams and collections, which can be naturally expressed using the existing mechanisms of our container-access pattern. Detail of this handling is provided in Section 7.5.

The complete list of stream operations modeled in our analysis is provided in Appendix B.

### 7.3 Modeling Stream Pipelines

In this section, we extend the container-access pattern to statically abstract stream pipeline instances and handle *structural* pipeline operations uniformly. This extension enables the identification of input elements (SOURCES) and output locations (TARGETS) for each abstract stream instance on-the-fly, allowing the reuse of existing rules to support the suppression of imprecise PFG edges and addition of precise shortcut edges, for achieving our precision goal in analyzing Java streams.

**7.3.1 Abstracting Stream Pipelines.** We begin by introducing how abstract stream instances are represented in our analysis. Recall from Sections 3.3 and 4.3 that a *host* abstracts container instances (objects of type `Collection` or `Map`), and is defined on  $\mathbb{L} \times \mathbb{K}$ . Hosts are created per allocation site, such that all dynamic container objects allocated at the same site are mapped to the same (or the same two, for `Map`) host(s). To model streams, we introduce a new host kind `Strm` into the domain  $\mathbb{K}$  to represent abstract stream pipeline instances. However, unlike the container abstraction, stream pipelines are not abstracted per allocation site. This difference in treatment is motivated by the fact that many stream *creation* operations delegate `Stream` object creation to a shared factory method hidden within the JDK library, and statically we will observe only a single allocation site (inside that factory method) for all such creations. Consequently, different *creation* operations—and multiple invocations of the same *creation* operation—appear to return the same abstract heap object, making them indistinguishable.

To more accurately capture the intuition of what constitutes a dynamic stream pipeline, we adopt a finer-grained abstraction: *one Strm host per call-site of a creation operation*. Specifically, the instruction label of each call-site for a *creation* operation is used to uniquely identify a `Strm` host. Figure 22(a) shows the intermediate three-address code representation of the two stream pipelines in our motivating example. (Three-address code is a commonly used IR for static analysis, and our analysis is also based on it.) In this IR, the two call-sites labeled 11 and 12 correspond to separate invocations of `Stream.of`, and are abstracted as distinct `Strm` hosts:  $h_{\text{Strm}}^{l_1}$  and  $h_{\text{Strm}}^{l_2}$ , respectively.

**7.3.2 Modeling Structural Operations.** With `Strm` hosts representing abstract stream instances, we now explain how to model structural operations using the extended container-access pattern:

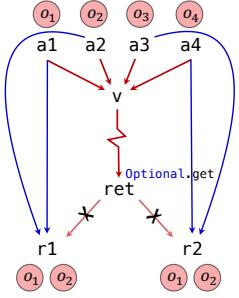
- **Structural Creation.** For stream creation, we introduce rule [S-CREATE] (Figure 23) and define an input set `STRUCRE`, where each pair  $\langle m, k \rangle$  specifies a *creation* operation  $m$  and its  $k$ -th argument as the source of input elements. At each call-site  $i$  of such a method, we: (i) create a `Strm` host  $h_{\text{Strm}}^i$  and derive  $h_{\text{Strm}}^i \in pt_H(\text{LHS}^i)$ , to associate it with the left-hand-side variable receiving the created `Stream` object, and (ii) record  $\text{arg}_k^i$  as a SOURCE of this host, denoted  $h_{\text{Strm}}^i \Leftarrow \text{arg}_k^i$ . For example, in Figure 22(a), for lines  $l_1$  and  $l_2$  (highlighted with blue), we have  $h_{\text{Strm}}^{l_1} \in pt_H(t1)$ , with  $a1$  and  $a2$  being its SOURCES. Similarly,  $h_{\text{Strm}}^{l_2} \in pt_H(t4)$ , with  $a3$  and  $a4$  being its SOURCES.
- **Structural Process.** A *structural process* operation  $m$ , such as `Stream.filter` or `Stream.findFirst` in Figure 22(a), highlighted in green, is registered in `CONTPROP` via  $\langle m, \text{Strm} \rangle$ , enabling rule

```

1814
1815 STRUCRE Stream t1 = Stream.of(a1, a2); // line l1,  $pt_H(t1) = \{h_{\text{Strm}}^{l_1}\}$ ,  $h_{\text{Strm}}^{l_1} \Leftarrow a1/a2$ 
1816 STRUPRO Stream t2 = t1.filter(lambda); //  $pt_H(t2) = \{h_{\text{Strm}}^{l_1}\}$ 
1817 STRUPRO Optional t3 = t2.findFirst(); //  $pt_H(t3) = \{h_{\text{Strm}}^{l_1}\}$ 
1818 STRUOUT Stud r1 = t3.get(); //  $h_{\text{Strm}}^{l_1} \Rightarrow r1$ 
1819 STRUCRE Stream t4 = Stream.of(a3, a4); // line l2,  $pt_H(t4) = \{h_{\text{Strm}}^{l_2}\}$ ,  $h_{\text{Strm}}^{l_2} \Leftarrow a3/a4$ 
1820 STRUPRO Stream t5 = t4.filter(lambda); //  $pt_H(t5) = \{h_{\text{Strm}}^{l_2}\}$ 
1821 STRUPRO Optional t6 = t5.findFirst(); //  $pt_H(t6) = \{h_{\text{Strm}}^{l_2}\}$ 
1822 STRUOUT Stud r2 = t6.get(); //  $h_{\text{Strm}}^{l_2} \Rightarrow r2$ 
1823

```

(a) IR for the two stream pipelines in motivating example



(b) PFG fragment under CUT-SHORTCUTS

Fig. 22. Example for handling structural operations.

[HOSTPROP-I] (in Figure 11) to propagate the associated Strm host from the receiver to the LHS. This models stream chaining accurately while retaining host identity. For the first pipeline in Figure 22(a), because  $h_{\text{Strm}}^{l_1} \in pt_H(t1)$ , we derive  $h_{\text{Strm}}^{l_1} \in pt_H(t2)$  and subsequently derive  $h_{\text{Strm}}^{l_1} \in pt_H(t3)$ . Similarly, for the second pipeline, we derive  $h_{\text{Strm}}^{l_2} \in pt_H(t5)$  and  $h_{\text{Strm}}^{l_2} \in pt_H(t6)$ .

- *Structural Output.* A *structural output* operation  $m$ , like `Optional.get` in Figure 22(a), highlighted in red, is registered in `CONT_EXIT` with  $\langle m, \text{Strm} \rangle$ , enabling (1) suppression of imprecise return edges by rule [CutCont] (in Figure 11), and (2) identification of TARGETS for Strm hosts via rule [HOSTTARGET] (in Figure 11). For example, the return variable of `Optional.get` is included in `cutRet` by [CutCont], and imprecise return edges from it to  $r1$  and  $r2$  are suppressed by the boxed premise of [RETURN] in Figure 7, as depicted in Figure 22(b). In addition, we derive  $h_{\text{Strm}}^{l_1} \Rightarrow r1$  and  $h_{\text{Strm}}^{l_2} \Rightarrow r2$  by [HOSTTARGET].

$$\frac{i \rightarrow_{call} m \quad \langle m, k \rangle \in \text{STRUCRE}}{h_{\text{Strm}}^i \in pt_H(\text{LHS}^i) \quad h_{\text{Strm}}^i \Leftarrow \arg_k^i} \text{ [S-CREATE]}$$

Fig. 23. The rule for *structural creation* stream operations.

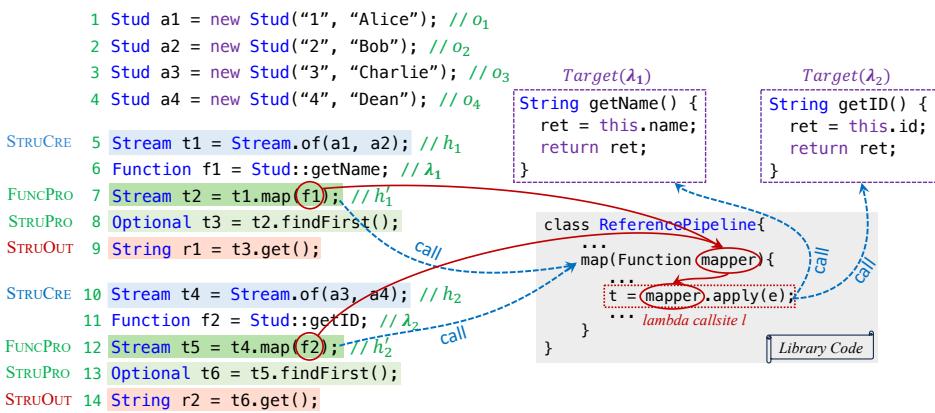
Finally, once both SOURCES and TARGETS are identified for each Strm host, rule [SHORTCUTCONT] (in Figure 11) is applied to add shortcut edges that connect sources to targets of the same host (e.g., from  $a1/a2$  to  $r1$ , and  $a3/a4$  to  $r2$ ), as illustrated in Figure 22(b).

*Summary of Extensions.* To support precisely abstracting stream pipelines and handling *structural* operations, we extend the container-access pattern in the following ways: (1) we extend the domain of host kinds  $\mathbb{K}$  to include a new kind `Strm` for abstract stream instances; (2) a rule [S-CREATE] is introduced to generate `Strm` hosts and locate their SOURCES; (3) for each *structural process* operation  $m$ , we specify  $\langle m, \text{Strm} \rangle \in \text{CONT_PROP}$ ; (4) for each *structural output* operation  $m$ , we specify  $\langle m, \text{Strm} \rangle \in \text{CONT_EXIT}$ ; and (5) the rule [HOSTPROP-II] in Figure 11 is updated by modifying the negation in its premise to  $\neg(a = \text{ret}^m \wedge \langle m, \_ \rangle \in \text{CONT_PROP} \cup \text{STRUCRE})$ . The modification of [HOSTPROP-II] prevents a `Strm` host that is generated within a *structural creation* operation from being propagated along the return edge.<sup>13</sup> Additional mechanisms for soundly handling potential unmodeled operations introduced in newer JDK versions are discussed in Appendix D.

<sup>13</sup>This blocking is necessary because many *structural creation* internally delegate stream creation to other such operations. Without this propagation blocking, multiple call-sites of the same *creation* operation (who calls another *structural creation* to create a stream instance) could result in the same `Strm` host (generated at the inner *structural creation* call-site) propagated out to all outer call-sites, undermining our precision goal.

## 1863 7.4 Handling Functional Operations

1864 In this section, we describe how to precisely handle *functional* stream operations, i.e., operations  
 1865 that take function-like values (such as lambdas or method references) as arguments. Such handling  
 1866 requires extending the container-access pattern by integrating it with an existing Java lambda  
 1867 analysis [18] to model a variety of operations. We begin by reviewing the necessary background on  
 1868 Java lambdas and method references, illustrated with a motivating stream example. We then outline  
 1869 how the existing lambda analysis models the creation and invocation of Java lambdas (or method  
 1870 references), and demonstrate why its results alone are not precise enough. Finally, we explain how  
 1871 the our stream handling interacts with this lambda analysis to obtain more precise results.  
 1872



1866 Fig. 24. Example of using *functional* operations in stream pipelines.  
 1867

1888 **7.4.1 Lambdas and Method References.** Java 8 introduced lambdas and method references to support  
 1889 functional-style programming, especially in conjunction with stream APIs. *Functional* stream  
 1890 operations (e.g., `Stream.map(Function)`) accept arguments of *functional interface* types—annotated  
 1891 interfaces that declare a special abstract method. These interfaces act as typed “function pointers,”  
 1892 allowing users to pass executable behavior as arguments. For example, in Figure 24, lines 7 and 12  
 1893 call the *functional* operation `Stream.map(Function)`, passing in arguments `f1` and `f2`, both of type  
 1894 `Function<Student, String>`.<sup>14</sup> A “function pointer” in Java is therefore a reference to a “functional  
 1895 object” (i.e., an instance of a functional interface) created using either a lambda expression (e.g.,  
 1896 `a -> a.getName()`) or a method reference (e.g., `Stud::getName`). Both forms result in a functional  
 1897 object whose *target method* (also called its *implementation method*) encodes the behavior to execute.  
 1898 In our example, lines 6 and 11 use method references `Stud::getName` and `Stud::getID` to create two  
 1899 functional objects (whose targets are the methods `Stud.getName` and `Stud.getID`, respectively), and  
 1900 then assign the two functional objects to pointers `f1` and `f2`, respectively.

1901 Internally, both lambdas and method references rely on the Java invokedynamic framework,  
 1902 using programmable dynamic linking to generate functional objects. Each functional object is  
 1903 associated with (1) a method handle to its *target method*, and (2) any captured environment values,  
 1904 if applicable. The key difference between lambdas and method references lies in how these two  
 1905 pieces of metadata are encoded in their syntax. In our example in Figure 24, only method references  
 1906 are used—chosen for their simplicity, because in this case they do not capture environment values  
 1907 and their target methods are easily identified.

1908 <sup>14</sup>The `Function` type is a standard functional interface provided in JDK, representing functions that take one argument and  
 1909 return a value. (The types of the argument and the return value are specified with generic typing.) Java provides other such  
 1910 interfaces, like `Consumer` (takes one argument, returns nothing) and `BiFunction` (takes two arguments, returns one).

When executing a *functional* stream operation, the function-pointer argument is passed into the internal stream-library code, where it is invoked multiple times—once for each element in the stream. In Figure 24, both `f1` and `f2` are passed to the parameter `mapper`, which is later invoked through the interface method `apply`—the special abstract method defined in the `Function` interface. This call site, labeled `l`, is thus identified as a *lambda call site*, and its actual callee is dynamically resolved to the *target method* of the functional object referred to by `mapper`. As a result, call site `l` resolves to `Stud.getName` for the first stream pipeline and to `Stud.getID` for the second. At runtime, the argument `e` at the lambda call site refers to each input element of the stream (i.e., the `Stud` objects), and serves as the receiver of the call (i.e., is bound to this in the target method). For instance, in the first pipeline, the lambda call site `l` is invoked twice—once on `o1` and once on `o2`—with `getName()` executed on each object. The return value of each invocation at the lambda call site is assigned to the left-hand-side variable `t`, which is then propagated through the stream pipeline as the mapped element. After execution, `r1` points to "Alice" (the result of invoking `getName()` on `a1`), and `r2` points to "3" (the result of invoking `ID()` on `a3`).

**7.4.2 Lambda Analysis.** Standard pointer analysis (with on-the-fly call-graph construction) is insufficient for resolving the dynamic targets of lambda call-sites. Consequently, a dedicated analysis is required to statically model functional call-graph edges. The lambda analysis proposed in [18] provides a comprehensive static modeling of Java’s invokedynamic-based functional features, including both lambda expressions and method references. In their work, method references are treated as a syntactically constrained subset of lambdas, allowing both constructs to be handled uniformly via a single rule set. For consistency, we refer to this unified approach as the *lambda analysis* throughout the article. This analysis interacts with the core pointer analysis in two key phases: (1) *Functional object creation*—where the analysis models the allocation of functional-interface instances and records key metadata, including its *target method* and any captured environment variables; and (2) *Lambda call-site resolution*—where the analysis resolves invocations on functional objects to their target methods, and models argument passing and return behavior.

$$\begin{array}{c}
 \frac{i : r = b.m(\dots) \quad \lambda \in pt(b) \quad \text{Target}(\lambda) = m}{i \rightarrow_{\lambda} m} [\text{L-CALL}] \qquad \frac{i \rightarrow_{\lambda} m}{\text{ret}^m \longrightarrow \text{LHS}^i} [\text{L-RET}] \\
 \\ 
 \frac{i \rightarrow_{\lambda} m \quad m \notin \text{Static} \quad \#cap^{\lambda} = 0}{\forall k \geq 1 : \text{Larg}_k^{i,\lambda} \longrightarrow \text{param}_{k-1}^m} [\text{L-InsLARG}] \qquad \frac{i \rightarrow_{\lambda} m \quad i \notin \text{STMDCALLSITE}}{\forall k \geq 1 : \arg_k^i \longrightarrow \text{Larg}_k^{i,\lambda}} [\text{ARG2LARG}]
 \end{array}$$

Fig. 25. Selected core rules of the lambda analysis. The boxed premise in **[ARG2LARG]** is added when integrated with our stream handling.

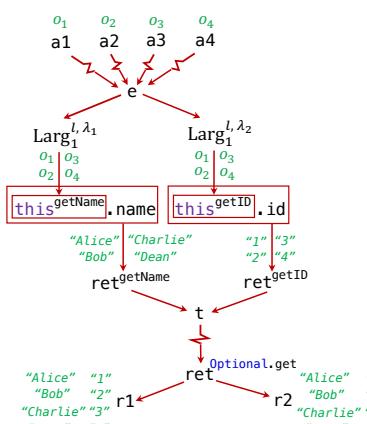
Figure 25 presents the inference rules relevant to our motivating example.<sup>15</sup> To align with our formalism, we rephrase the inference rules from [18] using the notational conventions adopted in this paper. We first introduce several new notations:  $\lambda \in \Lambda$  denotes a functional object generated by a lambda or method reference;  $\text{Target}(\lambda)$  maps a functional object  $\lambda$  to its target method; and  $\#cap_k^{\lambda}$  denotes the number of variables captured from environment by  $\lambda$ . These three quantities are predicates derived during the functional-object creation phase (whose rules are omitted). Although our analysis only interacts with the lambda analysis during its invocation phase (i.e., lambda call-site resolution), it is important to emphasize that both analyses are implemented on-the-fly and

<sup>15</sup>For brevity, the rules for functional objects creation are omitted, as they rely on lower-level modeling of invokedynamic, which is orthogonal to our analysis. We also omit rules related to captured variables and static target methods, as our illustrative examples focus on instance targets with no captured variables. The handling of dynamic dispatch of instance target methods is also not shown for simplicity. The complete rules of these complex scenarios are presented in Appendix B.

1961 operate in a mutually recursive manner with the pointer-analysis core—there is no predetermined  
 1962 order. In addition, we introduce the notation  $\text{Larg}_k^{i,\lambda}$  to denote a mock variable that abstracts the  
 1963  $k^{\text{th}}$  ( $k$  starts from 1) actual argument for a lambda call site  $i$  invoked on  $\lambda$ . The use of such mock  
 1964 variables simplifies the presentation of the interaction between the container-access pattern and  
 1965 the lambda analysis.

1966 We now explain how these rules interact on-the-fly with pointer analysis to build the pointer  
 1967 flow graph (PFG) illustrated in Figure 26, based on the code example in Figure 24, and highlight  
 1968 where imprecision arises along the way. The rule [L-CALL] constructs call-graph edges for lambda  
 1969 call sites (*lambda-call-edges* in short). Given an invoke instruction of the form  $i:r=b.m(\dots)$ , if the  
 1970 base pointer  $b$  points to a functional object  $\lambda$  whose target method is  $m$ , a lambda-call-edge  $i \rightarrow_{\lambda} m$   
 1971 is created. In our example (Figure 24), the functional objects  $\lambda_1$  and  $\lambda_2$  flow to the pointer mapper  
 1972 at the lambda call site  $l$  (inside the library code). Thus, two lambda-call-edges,  $l \rightarrow_{\lambda_1} \text{getName}$  and  
 1973  $l \rightarrow_{\lambda_2} \text{getID}$ , are generated.

1974 The rule [ARG2LARG] specifies how the mock variables  $\text{Larg}_k^{i,\lambda}$  receive their value flow in the  
 1975 lambda analysis. For now, we can set aside the boxed premise  $m \notin \text{STMLCALLSITE}$ —this condition



1991 Fig. 26. Selected PFG for the example  
 1992 in Figure 24 (without integration with  
 1993 stream handling).

1994 and no variables are captured by  $\lambda$ . Specifically, it adds PFG edges from each actual argument  $\text{Larg}_k^{i,\lambda}$   
 1995 to the corresponding formal parameter  $\text{param}_{k-1}^m$ . Note the index shift by  $-1$ : this adjustment is  
 1996 necessary to correctly pass the first actual argument to  $\text{this}^m$  of an instance method  $m$ , denoted as  
 1997  $\text{param}_0^m$  in the rule. For example, the first argument  $\text{Larg}_1^{l,\lambda_1}$  of the lambda-call-edge  $l \rightarrow_{\lambda_1} \text{getName}$   
 1998 flows to  $\text{this}^{\text{getName}}$  (similarly,  $\text{Larg}_1^{l,\lambda_2}$  flows to  $\text{this}^{\text{getID}}$ ), as depicted by the corresponding PFG  
 1999 edges in Figure 26.

2000 Finally, rule [L-RET] propagates the return value of the target method (i.e.,  $\text{ret}^m$ ) back to the  
 2001 left-hand-side variable (i.e.,  $\text{LHS}^i$ ) of the lambda call site. In our example, during pointer analysis  
 2002  $\text{ret}^{\text{getName}}$  points to the name field loaded from each receiver object (four Stud instances), and  
 2003 similarly,  $\text{ret}^{\text{getID}}$  points to the corresponding id fields. These returned objects flow to the LHS  
 2004 variable  $t$  at call site  $l$ , subsequently propagating to the returned value of the *structural output*  
 2005 operation  $\text{Optional.get}()$ , and eventually reaching both  $r1$  and  $r2$ . As a result, the points-to sets of  
 2006  $r1$  and  $r2$  both contain eight string objects (four student names and their four corresponding IDs).  
 2007

is introduced later to integrate the lambda analysis with our stream modeling. In essence, the lambda analysis identifies the  $k^{\text{th}}$  argument of the lambda call-site  $i$  (i.e.,  $\text{arg}_k^i$ ) as the value source for  $\text{Larg}_k^{i,\lambda}$  and adds edges in PFG to connect them accordingly. In our example, both  $\text{Larg}_1^{l,\lambda_1}$  and  $\text{Larg}_1^{l,\lambda_2}$  receive their values from  $\text{arg}_1^l$ . Because  $\text{arg}_1^l$  corresponds to the argument  $e$  at line  $l$ , which receives elements from both stream pipelines, the points-to set of  $e$  contains all objects referenced by pointers  $a1-a4$  (i.e.,  $o1-o4$ ). Consequently, these objects merge and flow into both  $\text{Larg}_1^{l,\lambda_1}$  and  $\text{Larg}_1^{l,\lambda_2}$ , as is shown in the PFG in Figure 26. Note that this merging causes imprecision: despite the lambda call site  $l$  being invoked separately on functional objects  $\lambda_1$  and  $\lambda_2$  (each corresponding to a distinct stream pipeline), the analysis merges their input elements in the argument pointer  $e$ .

The rule [L-INSLARG] subsequently propagates mock actual arguments to parameters for each lambda-call-edge  $i \rightarrow_{\lambda} m$ , provided that the target method  $m$  is an instance method

In general, input elements of distinct streams are merged into a single lambda call-site argument within the implementation of *functional* stream operations, resulting in imprecise propagation to the parameters (or this) of target methods of functional objects used in different stream pipelines.

**7.4.3 Integration with Stream Handling.** The central idea of integrating our extended container-access pattern with the lambda analysis is to replace the overly conservative argument flow (originally from lambda-call-site arguments, which merges elements from multiple stream pipelines) by precise shortcut edges derived from the SOURCES of Strm hosts. Consequently, we add the boxed premise  $i \notin \text{STMLCALLSITE}$  to the rule [ARG2LARG] in Figure 25. Here, STMLCALLSITE denotes the set of lambda call-sites that are part of any *functional* stream operation implementation (e.g., the call site  $l$  in Figure 24).<sup>16</sup> By adding this premise, only lambda invocations that are unrelated to any *functional* stream operation will have their mock actual arguments  $\text{Larg}_k^{i,\lambda}$  populated via the standard flow from the call-site arguments  $\text{arg}_k^i$ —in such cases, no precise source (e.g., identified stream SOURCES) is available to provide a more precise abstraction. On the other hand, for lambda call-sites involved in *functional* stream operations, we instead introduce explicit shortcut edges to improve precision. As discussed previously (Section 7.2), each *functional* stream operation has distinct semantics regarding how it transforms or consumes stream elements. Therefore, separate inference rules are required for each operation. The full set of rules is provided in Appendix C.

$$\frac{i \rightarrow_{\text{call}} m \quad m = \text{Stream.map}(\text{Function}) \quad \lambda \in pt(\text{arg}_0^i) \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i) \quad \text{Target}(\lambda) = m' \quad h_{\text{Strm}}^j \Rightarrow \text{Larg}_1^{*,\lambda} \quad h_{\text{Strm}}^i \in pt_H(\text{LHS}^i) \quad h_{\text{Strm}}^i \Leftarrow \text{ret}^{m'}}{h_{\text{Strm}}^j \Rightarrow \text{Larg}_1^{i,\lambda} \quad h_{\text{Strm}}^i \in pt_H(\text{LHS}^i) \quad h_{\text{Strm}}^i \Leftarrow \text{ret}^{m'}} \text{ [STMMap]}$$

Fig. 27. Rule for handling Stream.map.

The specific rule for handling the Stream.map operation is shown in Figure 27. First, we will create hosts of kind Strm for *structural-creation* operations at lines 5 (create  $h_{\text{Strm}}^5$ , henceforth denoted by  $h_1$  for simplicity) and line 10 (create  $h_{\text{Strm}}^{10}$ , henceforth denoted by  $h_2$ ) of Figure 24, and identify  $h_1$ 's SOURCES as  $a_1$  and  $a_2$ , and  $h_2$ 's SOURCES as  $a_3$  and  $a_4$ . At lines 7 and 12, where Stream.map is invoked, the rule [STMMap] applies: first, for each functional object  $\lambda$  referenced by the first argument  $\text{arg}_1^i$  (e.g.,  $\lambda_1$  for  $f_1$  at line 7, and  $\lambda_2$  for  $f_2$  at line 12), and for each host  $h$  of kind Strm associated with the receiver variable  $\text{arg}_0^i$  (e.g.,  $h_1$  for  $t_1$ ,  $h_2$  for  $t_4$ ), the first mock actual argument of any lambda invocation invoked on this  $\lambda$  is recorded as a TARGET of  $h$  (e.g.,  $h_1 \Rightarrow \text{Larg}_1^{i,\lambda_1}$  and  $h_2 \Rightarrow \text{Larg}_1^{i,\lambda_2}$ ). Second, a new Strm host  $h_{\text{Strm}}^i$  is generated to represent the resulting stream pipeline after the mapping operation. This host is associated with the LHS variable (e.g.,  $h'_1$  with  $t_2$  and  $h'_2$  with  $t_5$ ), and takes the return value of the lambda's target method as its SOURCE (e.g.,  $h'_1 \Leftarrow \text{ret}^{\text{getName}}$ ,  $h'_2 \Leftarrow \text{ret}^{\text{getID}}$ ).

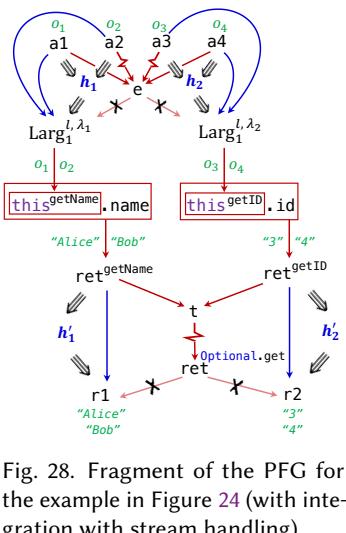


Fig. 28. Fragment of the PFG for the example in Figure 24 (with integration with stream handling).

(which connect SOURCES and TARGETS for each host). The bottom half of Figure 28 demonstrates the precise routing of lambda target return values to their corresponding stream outputs ( $r_1$  and

<sup>16</sup>This set is predetermined by the analysis designer via a lightweight local pre-analysis of the stream library code

2059 r2). Ultimately, our approach provides much more precise points-to sets results (r1 points to only  
 2060 "Alice", "Bob", and r2 points only to "3", "4"), compared to the merged, imprecise results shown  
 2061 in Figure 26.

2062 We emphasize that [STMMap] is specifically tailored to the Stream.map operation, reflecting its  
 2063 semantics of transforming a stream by applying a function to each element. Similar customized rules  
 2064 are required for other *functional* stream operations, which are detailed in Appendix B. Additionally,  
 2065 a conservative fallback mechanism for ensuring soundness is discussed in Appendix D.

## 2066 2067 7.5 Handling Collection-Interacting Operations

2068 In this section, we introduce how the transformations between Java streams and collections can  
 2069 be naturally modeled within our container-access pattern. Figure 29 illustrates an example of  
 2070 such transformation: a HashSet object is first allocated and then transformed into a Stream by  
 2071 invoking Collection.stream. (HashSet implements the Collection interface in Java.) The generated  
 2072 stream pipeline then performs a sorting operation via Stream.sorted, and subsequently stores the  
 2073 sorted elements into a List object. This kind of stream-to-collection interaction accounts for a large  
 2074 portion of real-world stream-usage patterns in Java applications [53]. According to the classification  
 2075 introduced in Section 7.2, the three stream operations in this example correspond to a *collection*  
 2076 *creation*, a *structural process*, and a *collection output*, respectively, each visually distinguished by  
 2077 subcategory-specific colors.

```
2078
2079     l1 Set set = new HashSet(); // ptH(set) = {hl1coll}
2080     ...
2081     COLLCRE l2 Stream t1 = set1.stream(); // ptH(t1) = {hl2strm}, hl1coll ⊲ hl2strm
2082     STRUPRO l3 Stream t2 = t1.sorted(); // ptH(t2) = {hl2strm}
2083     COLLOUT l4 List list = t2.toList(); // ptH(list) = {hl4coll}, hl2strm ⊲ hl4coll
2084     ...
2085
```

Fig. 29. An example of the interaction between streams and collections.

2086 The inference rules in Figure 30 formalize this interaction. Rule [C-CREATE] handles *collection*  
 2087 *creation* operations, which transform a container into a stream pipeline; rule [C-OUTPUT] handles the  
 2088 reverse transformation—i.e., converting a stream pipeline back into a container. In [C-CREATE], when  
 2089 a call site  $i$  invokes a method  $m \in \text{COLLCRE}$  (the set of all *collection creation* stream operations),  
 2090 we (1) generate a host of kind Strm at line  $i$  (i.e.,  $h_{\text{Strm}}^i$ ), associating it with the LHS variable by  
 2091 adding it to  $\text{pt}_H(\text{LHS}^i)$ ; and (2) identify the host(s) of kind Coll (i.e.,  $h_{\text{Coll}}^i$ ) that are associated with  
 2092 the receiver variable  $\text{arg}_0^i$ , and add a preorder  $\triangleleft$  (recall the semantics of  $\triangleleft$  given by rule [HOSTDEO-I]  
 2093 and [HOSTDEO-II] in Figure 11) between it and the newly generated Strm host (i.e.,  $h_{\text{Coll}}^i \triangleleft h_{\text{Strm}}^i$ ),  
 2094 indicating that all SOURCES of the former are also SOURCES of the later. Conversely, in [C-OUTPUT], for  
 2095 a call site  $i$  that invokes a method  $m \in \text{COLLOUT}$  (the set of all *collection output* stream operations),  
 2096 we (1) generate a new host of kind Coll (i.e.,  $h_{\text{Coll}}^i$ ) to abstract the container instance that receives  
 2097 the output elements of the stream pipeline, associating it with  $\text{LHS}^i$ , and (2) locate the Strm host(s)  
 2098  $h_{\text{Strm}}^i$  associated with the receiver variable  $\text{arg}_0^i$ , and derive  $h_{\text{Strm}}^i \triangleleft h_{\text{Coll}}^i$ , indicating the SOURCES  
 2099 are propagated from the Strm host to the newly generated Coll host.

$$\frac{i \rightarrow \text{call } m \quad m \in \text{COLLCRE} \quad h_{\text{Coll}}^i \in \text{pt}_H(\text{arg}_0^i)}{h_{\text{Strm}}^i \in \text{pt}_H(\text{LHS}^i) \quad h_{\text{Coll}}^i \triangleleft h_{\text{Strm}}^i} \text{ [C-CREATE]}$$

$$\frac{i \rightarrow \text{call } m \quad m \in \text{COLLOUT} \quad h_{\text{Strm}}^i \in \text{pt}_H(\text{arg}_0^i)}{h_{\text{Coll}}^i \in \text{pt}_H(\text{LHS}^i) \quad h_{\text{Strm}}^i \triangleleft h_{\text{Coll}}^i} \text{ [C-OUTPUT]}$$

Fig. 30. Rules for handling the interaction between streams and collections.

We now illustrate how the rules are applied to each line for the code snippet in Figure 29. At line  $l_1$ , a Coll host is created by [ColHost] in Figure 11, and associated with variable set. At line  $l_2$ , a *collection creation* stream operation is invoked, resulting in the creation of a Strm host  $h_{\text{Strm}}^{l_2}$ , which is associated with the LHS variable  $t1$ . Additionally,  $h_{\text{Coll}}^{l_1} \triangleleft h_{\text{Strm}}^{l_2}$  is derived. Line  $l_3$  corresponds to a *structural process* operation. Rule [HostProp-I] (in Figure 11) is applied to propagate the Strm host from  $t1$  to  $t2$ . Finally, line  $l_4$  invokes a *collection output* stream operation. According to [C-OUTPUT], a new Coll host is created to abstract the resulting collection instance and is associated with the LHS variable  $t4$  (i.e.,  $h_{\text{Coll}}^{l_4} \in pt_H(t4)$ ), with  $h_{\text{Strm}}^{l_2} \triangleleft h_{\text{Coll}}^{l_4}$ .

*Transformations through Collectors.* Among all stream operations, one API that requires special attention is Stream.collect(Collector), which performs a reduction using a Collector instance. The JDK class java.util.stream.Collectors provides a variety of factory methods to generate default collectors for different purposes—for example, Collectors.toSet() produces a collector that accumulates stream elements into a Set, while Collectors.groupingBy(Function) generates a collector that groups stream elements by a classification function and returns the result in a Map. In our analysis, we model such transformations precisely only when the collector is constructed via Collectors.toSet, Collectors.toList, or Collectors.toCollection. These collectors are widely used and have straightforward semantics—namely, accumulating all stream elements into a newly generated Collection instance—which aligns naturally with our handling in [C-OutPut] (by specifying Stream.collect in COLLOUT). In contrast, more complex collectors such as Collectors.toMap, Collectors.groupingBy or custom collectors exhibit dynamic behaviors that are difficult to reason about statically. Furthermore, empirical studies of stream-usage patterns [35, 62, 63] indicate that these complex collectors account for fewer than 10% of all collector usages in real-world programs. To balance soundness and practicality, we handle these cases conservatively, with detailed rules deferred to Appendix D.

## 8 Evaluation

In this section, we investigate the following research questions for evaluation.

**RQ1.** How does CUT-SHORTCUT compare to context-insensitive pointer analysis?

**RQ2.** How does CUT-SHORTCUT compare to conventional context-sensitive pointer analysis?

**RQ3.** How does CUT-SHORTCUT compare to state-of-the-art selective context-sensitive pointer analysis techniques that have similar goal (i.e., high efficiency with good precision)?

**RQ4.** How does each pattern of CUT-SHORTCUT affect the precision and efficiency?

**RQ5.** Is CUT-SHORTCUT<sup>S</sup> effective in analyzing programs with heavy stream usage?

*Implementation.* We implemented CUT-SHORTCUT and CUT-SHORTCUT<sup>S</sup> on top of the state-of-the-art Java static analysis framework TAI-E [75] (version 0.5.1, the latest stable release at the time this article was written). TAI-E is an imperative, highly extensible framework equipped with a powerful pointer analysis system that achieves high efficiency, precision, and soundness.<sup>17</sup> The core implementation of CUT-SHORTCUT is about 2,000 lines of Java code. We release and will maintain

<sup>17</sup>In our earlier conference paper [48], we also implemented the previous version of CUT-SHORTCUT on another pointer analysis framework, Doop [7] and reported experimental results on it. However, we no longer maintain that version and have not ported the new features (the optimizations in CUT-SHORTCUT and the CUT-SHORTCUT<sup>S</sup> extension) to it, because: (1) Doop uses a declarative Datalog-based solver that is significantly slower than TAI-E, which makes compared analyses far less scalable—for example, 2obj exhausting 128GB of memory or failing to terminate within 2 hours across all 10 benchmarks; and (2) its declarative nature makes extensive extensions difficult to debug. For these reasons, in this article we focus our evaluation on TAI-E, which provides a more scalable and efficient foundation, and thereby enables fair comparisons among CUT-SHORTCUT, CUT-SHORTCUT<sup>S</sup>, and state-of-the-art (selective) context-sensitive analyses.

2157 Table 1. Program sizes of benchmark programs, measured by the number of classes, methods, variable  
 2158 pointers, call sites, and statements in the intermediate representation (IR) after whole-program world building  
 2159 in the TAI-E analysis framework.

Program	Application (without JDK dependencies)					Whole-Program (with JDK dependencies)				
	#class	#method	#var-pointer	#callsite	#statement	#class	#method	#var-pointer	#callsite	#statement
soot	3,419	34,294	245,506	170,763	447,826	8,932	82,479	519,049	346,422	1,206,198
freecol	1,703	13,737	116,663	85,610	270,173	10,156	89,844	592,183	385,803	1,618,682
briss	1,728	15,315	137,974	92,441	408,859	8,649	78,970	518,207	343,183	1,525,911
gruntspudd	1,460	9,540	72,440	50,702	148,100	9,233	77,584	487,043	319,432	1,345,960
columba	4,777	38,499	238,449	162,386	523,377	12,388	106,410	650,498	430,980	1,709,493
jython	1,185	10,820	71,886	43,452	163,756	6,805	59,653	349,729	221,220	931,206
jedit	453	3,508	27,483	20,566	74,207	6,852	62,061	366,253	264,397	1,067,931
eclipse	2,529	25,145	199,838	130,898	518,812	8,086	73,400	474,402	306,804	1,279,765
hsqldb	492	5,696	53,757	37,843	124,664	6,159	55,794	339,936	222,872	911,043
findbugs	1,767	12,225	79,684	55,658	167,365	8,290	69,169	422,000	270,284	1,093,506

2172 the source code at [86]. All experiments were run on an Intel i7-13700K machine with 128 GB of  
 2173 running memory.  
 2174

2175 We use the same experimental settings for RQ1–RQ4, introduced below. For RQ5, different  
 2176 compared analyses, benchmarks, and precision metrics are employed, which are described in  
 2177 Section 8.5.

2178 *Compared Analyses.* To investigate RQ1–RQ4, we compare CUT-SHORTCUT to three types of  
 2179 pointer analyses: context insensitivity (CI), mainstream context sensitivity, and state-of-the-art  
 2180 selective context sensitivity. CI is the de facto fastest pointer analysis for Java. For conventional con-  
 2181 text sensitivity, we select the commonly-used 2-object-sensitive (2obj) [49, 50] and 2-type-sensitive  
 2182 (2type) [69] analyses. 2obj is highly-precise, and 2type often achieves much better scalability  
 2183 than 2obj while yielding comparable precision. For selective context-sensitivity, we consider a  
 2184 state-of-the-art approach ZIPPER<sup>e</sup> [41], which exhibits a better efficient and precision trade-off  
 2185 than many other selective approaches. We use ZIPPER<sup>e</sup>-2obj and ZIPPER<sup>e</sup>-2type to denote the  
 2186 selective 2-object-sensitive and 2-type-sensitive analyses guided by ZIPPER<sup>e</sup>. ZIPPER [39]—often  
 2187 evaluated as state-of-the-art in recent literature, with comparable precision to ZIPPER<sup>e</sup> but much  
 2188 less efficient—is not considered here because it fails to scale to most of the benchmark programs  
 2189 used in our experiments.  
 2190

2191 *Benchmarks.* To investigate RQ1–RQ4, we consider all the large and complex Java programs used  
 2192 for evaluation in recent related literature [28, 29, 32, 39–41, 46, 76], except for those that cannot  
 2193 be executed—because we cannot perform recall experiments on them for evaluating soundness.  
 2194 We analyze the programs with a large Java library, JDK1.6, which is commonly used in recent  
 2195 work [22, 23, 30, 31, 46]. In Table 1, we report program sizes in terms of key program elements.  
 2196 Specifically, we record the number of classes, methods, variable pointers, call sites, and total  
 2197 statements. Here, a “statement” refers to an instruction in the intermediate representation (IR) used  
 2198 by the TAI-E framework. All counts are obtained from the IR that the analysis framework builds for  
 2199 the entire program, rather than from the source code, to more accurately reflect the relationship  
 2200 between number of program elements and analysis performance. Table 1 reports both the size of  
 2201 the application part (excluding JDK dependencies but including other external libraries) and the  
 2202 size of the whole program (including the relevant portions of the JDK on which the application  
 2203 depends). All analyses are performed in a whole-program setting, where the pointer analysis fully  
 2204 analyzes the relevant JDK portions.  
 2205

*Precision Metrics.* To measure precision for RQ1-RQ4, we use the total number of objects pointed to by program variables, i.e.,  $\sum_v |pt(v)|$  (abbreviated as #var-pts), along with four independently useful clients that are widely-used as precision metrics in the pointer-analysis literature [28, 30, 32, 39–41, 46, 69, 70, 78]: a cast-resolution analysis (metric: the number of casts that may fail—#fail-cast), a method-reachability analysis (metric: the number of reachable methods—#reach-mtd), a devirtualization analysis (metric: the number of virtual call sites that cannot be disambiguated into monomorphic calls—#poly-call) and a call-graph-building analysis (metric: the number of call graph edges—#call-edge). In the rest of this section, we report all average speed-ups and average percent precision improvement using geometric means, as recommended by [17], because geometric means provide a more reliable summary of normalized numbers and help avoid misleading conclusions.

### 2217 8.1 RQ1: CUT-SHORTCUT vs. Context Insensitivity (CI)

2218 In this section, we examine how CUT-SHORTCUT fares against the fastest pointer analysis, CI. But  
 2219 in order to trust the reliability of CUT-SHORTCUT’s efficiency benefit, we must first confirm its  
 2220 soundness. In addition to the theoretical soundness proofs in Section 5.2, below we further validate  
 2221 its soundness by a recall experiment.

2222 *Soundness (Recall).* We execute all the evaluated programs with their default tests (e.g., for  
 2223 DaCapo benchmarks [6]) or the ones we input (e.g., for GUI programs, we click to interact with  
 2224 them and the results are recorded and will be provided in our artifact), then dynamically record their  
 2225 reachable methods and call-graph edges during execution, and examine how many of them can be  
 2226 recalled (over-approximated) by CUT-SHORTCUT and other analyses in our evaluation. (It is hard to  
 2227 instrument programs to gather other dynamic information, such as dynamic points-to relations.)  
 2228 The results show that CUT-SHORTCUT can recall all true reachable methods and call-graph edges  
 2229 discovered by other theoretically sound analyses<sup>18</sup>. Thus, in addition to the theoretical-soundness  
 2230 proof, we can trust that the remaining efficiency and precision results are reliable. The detailed  
 2231 results from the recall experiment are provided in Appendix E.  
 2232

2233 *Efficiency.* Figure 31 shows graphically the elapsed time of all analyses across the 10 benchmarks,  
 2234 and Table 2 presents the efficiency and precision results in full detail. As expected, CI is substantially  
 2235 faster than (selective) context-sensitive analyses. However, CUT-SHORTCUT achieves even better  
 2236 performance: it is faster than CI on 7 out of 10 benchmarks, and as fast as CI on the remaining 3.  
 2237 CUT-SHORTCUT’s speed advantage is more pronounced as program size increases. For example,  
 2238 on the largest benchmark columba (which has the highest number of program elements in whole-  
 2239 program setting, as reported in Table 1), CUT-SHORTCUT outperforms CI by more than 20 seconds  
 2240 (26% faster). On relatively smaller benchmarks (hsqldb, jython, findbugs), CUT-SHORTCUT performs  
 2241 on par with CI. Overall, CUT-SHORTCUT is on average 1.18× faster than CI across 10 benchmarks.

2242 The speed of CUT-SHORTCUT arises from both its methodology and its precision improvements.  
 2243 First, unlike context sensitivity—which improves precision by cloning methods and heap objects  
 2244 under different contexts—Cut-SHORTCUT enhances precision by suppressing imprecise PFG edges  
 2245 and adding precise ones, which generally introduces little overhead (see Section 6.2 for a formal  
 2246 discussion). Second, in contrast to existing selective context-sensitivity approaches, CUT-SHORTCUT  
 2247 requires no pre-analysis phase (a more detailed comparison is provided in Section 8.3); instead, it  
 2248 runs the context-insensitive pointer analysis algorithm on an on-the-fly constructed PFG. Third,

2249 <sup>18</sup>For a pointer analysis to be practical on real-world Java programs, it must handle dynamic features. However, some of  
 2250 these dynamic features, like reflection, cannot be soundly analyzed in general at static time [42], and therefore no existing  
 2251 pointer analysis is fully sound in their presence. The goal of our recall experiment is to demonstrate that the implementation  
 2252 of CUT-SHORTCUT does not introduce *additional* unsoundness: the set of dynamically recorded reachable methods and  
 2253 call-graph edges recalled by other compared analyses can also be recalled by CUT-SHORTCUT.

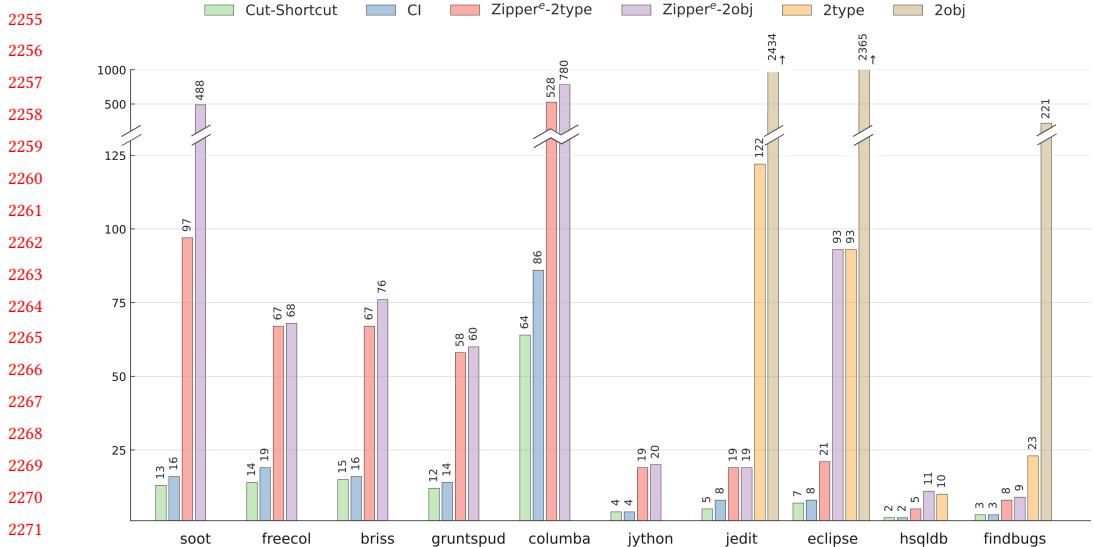


Fig. 31. Analysis time (in seconds) of CUT-SHORTCUT, context-insensitive (CI), ZIPPER<sup>e</sup>-guided selective 2-type and 2-object sensitive, and conventional context-sensitive (2type and 2obj) pointer analyses. CUT-SHORTCUT (green bars) is consistently as fast as or faster than all compared analyses, including CI, across all benchmarks.

CUT-SHORTCUT prevents a substantial amount of spurious points-to information from propagating across the program. Because the cost of suppressing imprecise edges and adding precise ones is negligible compared to the efficiency gains from improved precision, CUT-SHORTCUT can in practice outperform even CI.

*Precision.* As shown in Table 2, CUT-SHORTCUT consistently achieves better precision than CI across all benchmarks and all precision metrics, with particularly notable improvements for #fail-cast, #call-edge, and #var-pts. These gains arise from guiding PFG construction using the proposed patterns, demonstrating that the three patterns underlying CUT-SHORTCUT are effective in improving precision <sup>19</sup>.

## 8.2 RQ2: CUT-SHORTCUT vs. Conventional Context Sensitivity

In this section, we investigate how CUT-SHORTCUT performs compared to mainstream conventional context-sensitive pointer analyses, namely 2obj and 2type.

*Efficiency.* As shown in Table 2, 2obj runs out of 128GB of memory for 7 out of 10 programs. 2type performs slightly better, but still runs out of memory on 6 programs. In contrast, CUT-SHORTCUT successfully completes all benchmarks. On the programs on which 2obj and 2type do succeed, CUT-SHORTCUT is remarkably faster: for the three programs on which 2obj succeeds, CUT-SHORTCUT is 229.7× faster, while CUT-SHORTCUT is 10.6× on the 4 programs on which 2type succeeds. In summary, CUT-SHORTCUT provides dramatically better scalability and efficiency than conventional context sensitivity.

<sup>19</sup>Relative to the results reported earlier in [48]: (1) we now include results for an additional compared analysis, ZIPPER<sup>e</sup>-2type; (2) all analyses show improved efficiency on the updated TAI-E framework, due to its recent performance enhancements; and (3) with the refinements we introduced to the field-access and container-access patterns, CUT-SHORTCUT is now more efficient, albeit with slightly reduced precision on some metrics. We find this trade-off worthwhile—our refinements make both patterns more principled—and therefore make CUT-SHORTCUT easier to adopt, extend, and maintain in practice.

Table 2. Efficiency and precision results of context-insensitive (CI), conventional context-sensitive (2obj and 2type), ZIPPER<sup>e</sup>-guided selective 2-object and 2-type sensitive analyses, and CUT-SHORTCUT across 10 benchmarks. For all precision metrics, smaller values indicate better precision. “OOM” indicates that the analysis terminated with an out-of-memory error (unscalable under our experimental settings).

Program	Analysis	Time(s)	#fail-cast	#reach-mtd	#poly-call	#call-edge	#var-ppts
soot	CI	16	16,640	32,913	16,385	415,710	84,491,342
	2obj	OOM	—	—	—	—	—
	2type	OOM	—	—	—	—	—
	ZIPPER <sup>e</sup> -2obj	488	9,392	32,165	13,837	281,289	14,701,551
	ZIPPER <sup>e</sup> -2type	97	11,683	32,275	14,724	353,082	23,224,281
freecol	CUT-SHORTCUT	13	10,321	32,659	15,188	328,689	30,248,396
	CI	19	9,739	46,915	15,489	322,837	71,774,368
	2obj	OOM	—	—	—	—	—
	2type	OOM	—	—	—	—	—
	ZIPPER <sup>e</sup> -2obj	68	6,885	46,058	13,501	277,883	15,650,971
briss	ZIPPER <sup>e</sup> -2type	67	8,269	46,282	13,735	281,618	19,079,170
	CUT-SHORTCUT	14	6,434	46,448	14,210	293,814	21,925,124
	CI	16	7,788	41,853	12,763	294,125	65,924,173
	2obj	OOM	—	—	—	—	—
	2type	OOM	—	—	—	—	—
gruntspud	ZIPPER <sup>e</sup> -2obj	76	5,998	41,426	11,579	260,792	23,098,275
	ZIPPER <sup>e</sup> -2type	67	6,752	41,468	11,709	263,373	25,024,647
	CUT-SHORTCUT	15	5,531	41,515	11,931	265,592	27,606,506
	CI	14	6,749	39,863	12,348	274,755	49,608,582
	2obj	OOM	—	—	—	—	—
columba	2type	OOM	—	—	—	—	—
	ZIPPER <sup>e</sup> -2obj	60	5,239	39,200	10,981	229,902	16,225,651
	ZIPPER <sup>e</sup> -2type	58	5,774	39,266	11,082	233,617	17,836,307
	CUT-SHORTCUT	12	4,759	39,465	11,699	230,641	18,822,052
	CI	86	10,453	56,737	17,680	426,134	168,420,335
jython	2obj	OOM	—	—	—	—	—
	2type	OOM	—	—	—	—	—
	ZIPPER <sup>e</sup> -2obj	780	8,070	55,809	15,897	341,394	42,395,433
	ZIPPER <sup>e</sup> -2type	528	8,904	55,970	16,205	351,855	48,735,371
	CUT-SHORTCUT	64	7,416	56,151	16,724	363,149	60,650,004
jedit	CI	4	2,402	13,086	2,952	121,489	12,107,477
	2obj	OOM	—	—	—	—	—
	2type	OOM	—	—	—	—	—
	ZIPPER <sup>e</sup> -2obj	20	1,851	12,612	2,598	111,632	8,093,191
	ZIPPER <sup>e</sup> -2type	19	1,987	12,648	2,642	111,964	8,115,225
eclipse	CUT-SHORTCUT	4	1,903	12,956	2,813	117,943	9,731,365
	CI	8	4,090	25,318	6,358	150,742	25,849,267
	2obj	2,434	2,566	24,229	4,944	122,037	1,645,511
	2type	122	3,083	24,278	5,108	122,889	1,939,458
	ZIPPER <sup>e</sup> -2obj	19	2,910	24,331	5,207	123,645	2,420,703
hsqldb	ZIPPER <sup>e</sup> -2type	19	3,272	24,395	5,288	124,632	2,656,387
	CUT-SHORTCUT	5	2,851	24,691	5,869	133,691	5,735,162
	CI	8	5,078	23,973	10,664	184,875	25,095,353
	2obj	2,365	3,621	23,212	9,734	164,462	2,626,243
	2type	93	4,181	23,355	9,920	166,097	3,102,358
findbugs	ZIPPER <sup>e</sup> -2obj	93	4,056	23,593	9,948	170,952	12,066,219
	ZIPPER <sup>e</sup> -2type	21	4,442	23,623	10,108	173,328	12,859,903
	CUT-SHORTCUT	7	3,915	23,783	10,329	175,241	13,545,299
	CI	2	1,653	11,586	1,830	64,398	1,850,474
	2obj	OOM	—	—	—	—	—
findbugs	2type	10	1,066	11,171	1,480	56,624	489,923
	ZIPPER <sup>e</sup> -2obj	11	927	11,215	1,511	56,930	729,197
	ZIPPER <sup>e</sup> -2type	5	1,100	11,258	1,100	57,475	769,173
	CUT-SHORTCUT	2	1,151	11,429	1,689	60,695	1,065,182
	CI	3	3,374	17,352	4,480	107,391	7,326,707
findbugs	2obj	221	1,962	16,821	3,578	89,003	886,778
	2type	23	2,338	16,874	3,770	90,073	994,799
	ZIPPER <sup>e</sup> -2obj	9	2,428	17,093	3,932	95,278	2,592,066
	ZIPPER <sup>e</sup> -2type	8	2,672	17,115	4,102	96,048	2,720,052
	CUT-SHORTCUT	3	2,112	17,182	4,226	97,461	2,123,513

*Precision.* When 2obj succeeds (on 3 out of 10 programs), it achieves the highest precision across all metrics, outperforming CUT-SHORTCUT and all other analyses. When 2type succeeds (on 4 out of 10 programs), it is generally more precise than CUT-SHORTCUT, however, it performs worse on the #fail-cast metric for 3 of the 4 programs. For #var-pts, 2obj and 2type typically yield smaller sets than CUT-SHORTCUT, reflecting their finer-grained context distinctions. On average for the precision improvement over CI on #var-pts, CUT-SHORTCUT retains 70.1% of 2obj’s improvement (over the 3 programs on which 2obj succeeds) and 67.7% of 2type’s improvement (over the 4 programs on which 2type succeeds). These results highlight the high precision of conventional context-sensitive analyses; however, their poor scalability makes them impractical as a general solution for large programs.

### 8.3 RQ3: CUT-SHORTCUT vs. State-of-the-Art selective Context Sensitivity

In this section, we compare CUT-SHORTCUT to ZIPPER<sup>e</sup> [41], a state-of-the-art selective context sensitivity approach. We use the default configuration of ZIPPER<sup>e</sup> as in [41], which is well-tuned and achieves a very good trade-off between efficiency and precision. We compare CUT-SHORTCUT to 2 variants, ZIPPER<sup>e</sup>-guided 2-object-sensitive (ZIPPER<sup>e</sup>-2obj) and ZIPPER<sup>e</sup>-guided 2-type sensitive (ZIPPER<sup>e</sup>-2type) analyses.

*Efficiency.* A ZIPPER<sup>e</sup>-guided analysis consists of three stages: (1) a context-insensitive pre-analysis (a standard CI analysis, necessary to provide information for ZIPPER<sup>e</sup>’s method-selection strategy), (2) running ZIPPER<sup>e</sup> itself (to select a set of methods according to its strategy), and (3) the main analysis (a context-sensitive analysis applied only to the selected methods). Thus, the elapsed time of a ZIPPER<sup>e</sup>-guided analysis is the sum of all three stages. As shown in Table 2, ZIPPER<sup>e</sup> is substantially more scalable (completing on all benchmarks) and faster than conventional 2obj and 2type analyses. Nonetheless, CUT-SHORTCUT remains more efficient, achieving an average speedup of 6.9× over ZIPPER<sup>e</sup>-2obj and 4.3× over ZIPPER<sup>e</sup>-2type. To compare ZIPPER<sup>e</sup> and CUT-SHORTCUT thoroughly, Table 3 breaks down the elapsed time of the ZIPPER<sup>e</sup>-guided analyses. For ZIPPER<sup>e</sup>, the “Total time” column lists the total analysis time (as reported in Figure 31 and Table 2); “Pre-analysis” shows the combined time of the CI pre-analysis and the method-selection phase (accounts for 44.0% of the total time of ZIPPER<sup>e</sup>-2obj and 70.0% for ZIPPER<sup>e</sup>-2type, on average); and “Main analysis” reports the time of ZIPPER<sup>e</sup>-guided context-sensitive analysis.

*Precision.* Overall, CUT-SHORTCUT is comparable in precision to ZIPPER<sup>e</sup>-2obj and ZIPPER<sup>e</sup>-2type. For #var-pts, CUT-SHORTCUT achieves on average 83.7% of the precision improvement of ZIPPER<sup>e</sup>-2obj (over CI) and 87.8% of that of ZIPPER<sup>e</sup>-2type (over CI). Regarding the four client metrics, CUT-SHORTCUT and ZIPPER<sup>e</sup> exhibit different strengths: for #fail-cast, CUT-SHORTCUT outperforms ZIPPER<sup>e</sup>-2obj on all programs except hsqldb, soot, and jython, and outperforms ZIPPER<sup>e</sup>-2type on all programs except hsqldb. For the remaining metrics, the ZIPPER<sup>e</sup>-guided analyses generally outperform CUT-SHORTCUT across most benchmarks, while the results of CUT-SHORTCUT remain competitive—for instance, CUT-SHORTCUT outperforms both ZIPPER<sup>e</sup>-2obj and ZIPPER<sup>e</sup>-2type on #var-pts for findbugs, and outperforms ZIPPER<sup>e</sup>-2type on #call-edge for soot and gruntsput. We anticipate that incorporating additional program patterns following the principle used in CUT-SHORTCUT could further enhance its precision.

Although ZIPPER<sup>e</sup> and CUT-SHORTCUT improve precision in fundamentally different ways, a natural question is whether the methods selected by ZIPPER<sup>e</sup> overlap with those involved in CUT-SHORTCUT’s PFG edge suppression and addition. Table 3 provides a detailed answer. On average, CUT-SHORTCUT considers 5,281 methods (i.e., methods involved in PFG edge suppression and addition), accounting for 17.2% of all reachable methods per program, while ZIPPER<sup>e</sup> selects 3,770 methods. Among them, on average only 34.5% of the methods identified by CUT-SHORTCUT overlap

Table 3. Detailed comparison between ZIPPER<sup>e</sup>-guided selective context sensitivity and CUT-SHORTCUT.

Program	CUT-SHORTCUT		Selective Context-sensitivity					Overlapped methods
	Time	#involved-methods	Analysis	Pre-analysis	Main-analysis	Total time	#selected-methods	
soot	13	5,585	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	23	465 74	488 97	12,102	71.91%
freecol	14	7,818	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	55	13 12	68 67	3,845	27.31%
briss	15	6,891	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	49	27 18	76 67	3,211	26.43%
gruntspud	12	6,922	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	47	13 11	60 58	3,503	26.77%
columba	64	9,415	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	462	318 66	780 528	4,069	21.19%
jython	4	2,270	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	17	3 2	20 19	1,611	35.29%
jedit	5	5,095	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	16	3 3	19 19	2,278	26.09%
eclipse	7	4,160	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	13	80 8	93 21	3,404	46.54%
hsqldb	2	1,754	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	4	7 1	11 5	1,491	47.43%
findbugs	3	2,902	ZIPPER <sup>e</sup> -2obj ZIPPER <sup>e</sup> -2type	6	3 2	9 8	2,182	40.80%

with those selected by ZIPPER<sup>e</sup> (see column “Overlapped methods”). This limited overlap highlights the orthogonality of the two approaches, and suggests that there may be opportunities for future schemes that combine their strengths.

#### 8.4 RQ4: Effect of Each Pattern of CUT-SHORTCUT

We instantiated the CUT-SHORTCUT philosophy via three patterns, so it is natural to ask which pattern contributes most to the precision of the algorithm. To answer this question, we conducted experiments running CUT-SHORTCUT with different combinations of patterns to isolate their impact.

*Precision.* We use CI as a baseline (i.e., no CUT-SHORTCUT pattern enabled). We measure the impact of each pattern by computing the precision improvement over CI for different pattern combinations, and normalize the results with respect to the overall improvement achieved by enabling all three patterns together. Figure 32 shows the results for the metric #var-pts. On average, using only the field-access pattern achieves an improvement of 22.7%, while combining the field-access and container-access patterns yields 95.8% improvement (with all three patterns together achieving 100%). For the other four metrics, the relative contributions vary. For example, on average, the field-access pattern; the field-access and container-access combination; and all three patterns together improve precision by 70.5%, 95.1%, and 100%, respectively, for #reach-mtd; and by 11.4%, 88.3%, and 100%, respectively, for #fail-cast.

*Efficiency.* Table 4 reports the elapsed time of CUT-SHORTCUT under different combinations of the three patterns. Including CI (which corresponds to enabling no patterns), the overall trend is that enabling more patterns makes the analysis run faster. With only the field-access pattern enabled, CUT-SHORTCUT runs faster than CI on 5 out of 10 programs, as fast as CI on 3 programs,

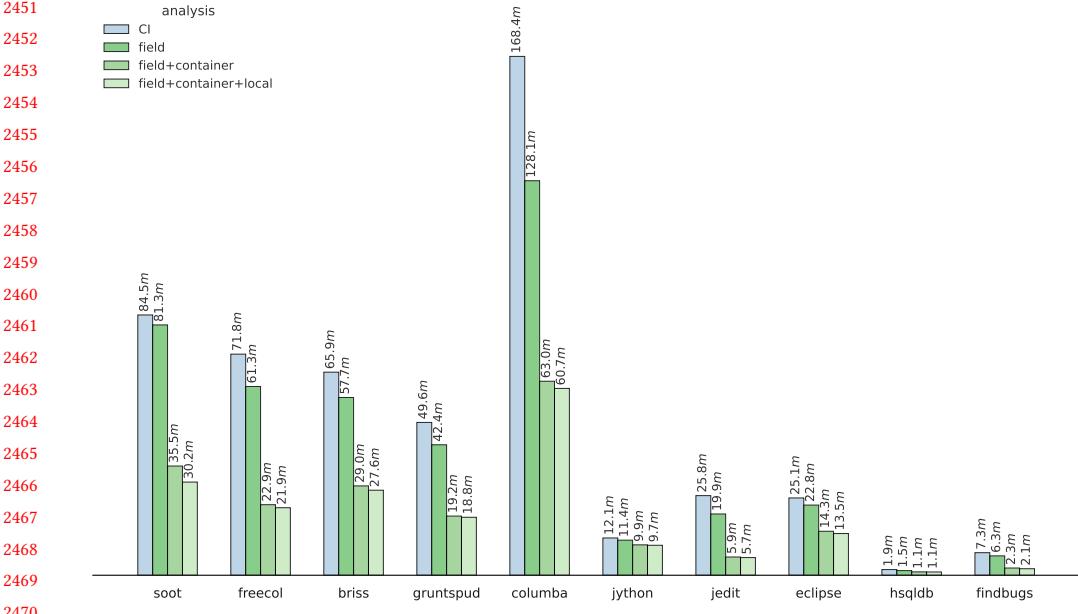


Fig. 32. #var-pts (in millions) for different combinations of the three patterns (the field-access pattern, abbreviated as “field”; the container-access pattern, abbreviated as “container”; and the local-flow pattern, abbreviated as “local”) of CUT-SHORTCUT. Small value indicate better precision.

Table 4. The corresponding elapsed time (in seconds) of the analyses in Figure 32.

	soot	freecol	briss	gruntspud	columba	jython	jedit	eclipse	hsqldb	findbugs
CI	16.5	18.9	15.8	13.9	86.3	4.4	8.0	7.8	1.9	3.0
Field	15.2	18.0	16.8	14.2	67.6	4.4	7.3	7.7	1.9	3.0
Field+Container	15.4	15.4	16.4	13.5	72.4	4.7	5.6	7.5	1.9	2.9
Field+Container+Local	13.3	13.6	15.3	12.3	64.3	4.4	5.5	7.2	1.9	2.7

and slightly slower on 2 programs. Adding the container-access pattern further reduces analysis time (or leave it unchanged) on 8 out of 10 programs, while the remaining 2 are slightly slower than with field-access alone. Finally, enabling all three patterns (field-access, container-access, and local-flow) further reduces analysis time on 9 out of 10 programs, with the last one unchanged compared to field + container. While one might expect that enabling more patterns—thus adding more analysis logic—would slow the analysis down, in practice, the opposite occurs: each additional pattern filters out more spurious points-to information, reducing the propagation of points-to facts during pointer analysis. The overhead of identifying suppressed and added PFG edges is negligible compared to the performance gain from reduced propagation. These results also suggest that incorporating additional patterns in the future could further improve efficiency.

## 8.5 RQ5: Effectiveness of CUT-SHORTCUT<sup>S</sup>

CUT-SHORTCUT<sup>S</sup> extends CUT-SHORTCUT to streams by statically abstracting pipelines, tracking their input elements and output locations, and modeling stream–lambda interactions on the fly, thereby improving precision without relying on context sensitivity. To evaluate CUT-SHORTCUT<sup>S</sup> comprehensively—not only on isolated operations but also in realistic usage—we constructed three

complementary benchmark suites (20 programs in total). These suites, released in our artifact, balance *operation coverage and controllability, representativeness, and realism*:

- (a) *Synthetic* (10 programs): micro-benchmarks that systematically cover combinations of stream operations on small inputs, each designed to isolate and highlight a specific aspect of stream behavior, with high controllability and enabling easy manual verification of analysis results.
- (b) *Workflow-mimicking* (5 programs): enterprise-backend-inspired benchmarks built with hand-crafted in-memory mock databases. This suite adapts and extends the BSS benchmark [82], originally generated by the S2S tool [64], which translates SQL queries into Java stream code. The resulting programs exercise end-to-end data-processing pipelines (e.g., filter-map-aggregate, PO/VO conversions), thereby bridging the gap between toy micro-benchmarks and production code while avoiding external dependencies. We introduce this suite because real enterprise backends rely heavily on streams for such data-processing and transmission tasks. However, analyzing real enterprise backends would additionally require modeling database queries, framework-specific behaviors (e.g., Spring DI & AOP), and other challenges that demand specialized solutions in pointer analysis and are orthogonal to our contribution. Thus, although current pointer analysis techniques do not yet fully address the complexities of enterprise applications, this benchmark suite serves as a proxy: it demonstrates the degree to which our stream-handling technique can improve precision once those obstacles are addressed, offering a forward-looking perspective on its applicability to real-world enterprise software.
- (c) *Real-world* (5 programs): open-source Java projects with intensive stream usage, providing realistic code bases outside the enterprise domain.

Together, the three suites strike a balance: the synthetic suite provides broad operation coverage with high controllability, the workflow-mimicking suite captures realistic backend-style usage, and the real-world suite validates effectiveness on existing projects beyond the enterprise domain. Detailed descriptions of each suite appear in the following subsections. Each of these benchmarks is analyzed with a large Java library, JDK 1.8.0\_312.

*Precision Metrics.* To measure precision while separating the parts of a program affected by stream computations from the rest, we introduce two new metrics tailored for RQ5: (1) The total number of objects pointed to by variables in application code ( $\sum_{v \notin \text{JDK}} |pt(v)|$ ), denoted by #app-vpt. We use this metric instead of #var-pts to focus on the application-level portion of the program, avoiding dilution from pointers in the Java 8 library, which account for a large share of #var-pts. (2) The total number of objects pointed to by stream-output pointers, denoted by #stm-out. This metric includes both the outputs of stream pipelines (i.e., the TARGET pointers of  $h_{\text{Strm}}$ ) and the outputs of collections directly derived from stream pipelines.

*Compared Analyses.* We compare CUT-SHORTCUT<sup>S</sup> (CSC-S) with four baselines: context insensitivity (CI), conventional context sensitivity, selective context sensitivity, and CUT-SHORTCUT without the stream extension (CSC). CI serves as a highly efficient precision baseline. For conventional context sensitivity, we use 1-object sensitivity (1obj), because 2obj does not scale (running out of 128GB memory) even on toy Java 8 programs due to the complexity of the JDK 8 library. For selective context sensitivity, we use selective 4-object sensitivity with a 3-layer context-sensitive heap abstraction targeting stream-related libraries (JDK collections, maps, streams, and the Optional class), denoted by stm-4obj3h. We fix the context depth at 4obj3h, because deeper settings do not yield additional precision improvements for stream-output pointers. We do not report the results for selective call-site sensitivity, because under the maximum depth that such an analysis algorithm can succeed in our experimental setting, it performs worse than stm-4obj3h in both run time and precision across all suites.

Table 5. Efficiency and precision results of context-insensitive analysis (CI), conventional context sensitivity (1obj), selective 4-object sensitivity with 3 context-sensitive heap abstractions on stream-related library (stm-4obj3h), CUT-SHORTCUT (CSC), and CUT-SHORTCUT<sup>S</sup> (CSC-S) across 10 synthetic micro-benchmarks.

Micro-benchmark	Static Analysis			#app-vpt	#stm-out	Micro-benchmark	Static Analysis			#stm-out	#pipe-line	Analysis	Time(s)
	#stm-out	#pipe-line	Analysis				#stm-out	#app-vpt	#stm-out				
simpleStruc	6	6	CI	2	13,905	46	stmReduce	34	8	CI	1	2,987	216
			1obj	4	9,014	46				1obj	4	1,752	216
			stm-4obj3h	24	5,550	46				stm-4obj3h	13	1,979	216
			CSC	2	8,716	46				CSC	1	937	216
			CSC-S	2	3,112	22				CSC-S	1	339	52
buildStm	6	10	CI	1	8,951	66	stmCollect	30	4	CI	1	8,876	128
			1obj	4	257	66				1obj	4	7,236	128
			stm-4obj3h	14	3,263	66				stm-4obj3h	13	434	108
			CSC	1	503	66				CSC	1	512	108
			CSC-S	1	103	13				CSC-S	1	230	32
simpleFunc	12	6	CI	2	10,527	92	stmCollInter	10	6	CI	1	8,656	84
			1obj	4	7,270	92				1obj	4	6,953	84
			stm-4obj3h	20	4,065	92				stm-4obj3h	16	1,956	84
			CSC	2	7,027	92				CSC	1	3,931	60
			CSC-S	2	3,163	18				CSC-S	1	235	42
stmMap	16	7	CI	1	2,707	119	primStm	4	4	CI	1	1,375	8
			1obj	4	502	119				1obj	4	99	8
			stm-4obj3h	13	2,511	119				stm-4obj3h	15	1,263	8
			CSC	1	766	119				CSC	1	263	8
			CSC-S	1	159	18				CSC-S	1	54	6
stmFlatMap	20	4	CI	2	1,700	116	conservStm	9	5	CI	2	11,729	107
			1obj	4	1,454	112				1obj	4	5,192	89
			stm-4obj3h	25	752	96				stm-4obj3h	14	971	102
			CSC	2	684	94				CSC	2	3,308	106
			CSC-S	2	532	22				CSC-S	2	2,236	106

8.5.1 *Synthetic Micro-Benchmarks.* To comprehensively evaluate CUT-SHORTCUT<sup>S</sup> on stream APIs, we constructed a suite of 10 synthetic micro-benchmarks. Each benchmark targets specific aspects of stream usage, covering a wide range of operations and their combinations. For example, `buildStm` focuses on stream builders and concatenation; `simpleStruc` covers simple structural operations (and their combinations); `simpleFunc` addresses simple functional operations (and their combinations); `stmCollInter` captures bidirectional interactions between streams and collections; `primStm` examines primitive-type streams and their conversions to object streams and vice versa; and `conservStm` evaluates conservatively modeled operations. The remaining benchmarks are self-explanatory by their name (see the “Micro-benchmark” column in Table 5).

Each benchmark contains 4–10 stream pipelines (see the “#pipeline” column in Table 5), with the number of input heap objects per pipeline ranging from 1 to 5. To provide a baseline for evaluating precision, we manually recorded the abstract heap objects pointed to by stream-output pointers during dynamic execution (results available in our artifact) and reported their totals for each benchmark (see the “Dyn #stm-out” column). To validate soundness, we verified that every heap object observed dynamically for stream-output pointers is also captured by the static analyses, including both CUT-SHORTCUT and CUT-SHORTCUT<sup>S</sup> (achieving 100% recall). This measurement confirms that the precision results of CUT-SHORTCUT<sup>S</sup> are trustworthy.

*Efficiency.* The red lines in Figure 33 highlight the efficiency advantage of CUT-SHORTCUT<sup>S</sup>: across all 10 micro-benchmarks, its runtime is comparable to CI and CUT-SHORTCUT. Because these benchmarks are relatively small, the efficiency gains of CUT-SHORTCUT<sup>S</sup> (and CUT-SHORTCUT) over

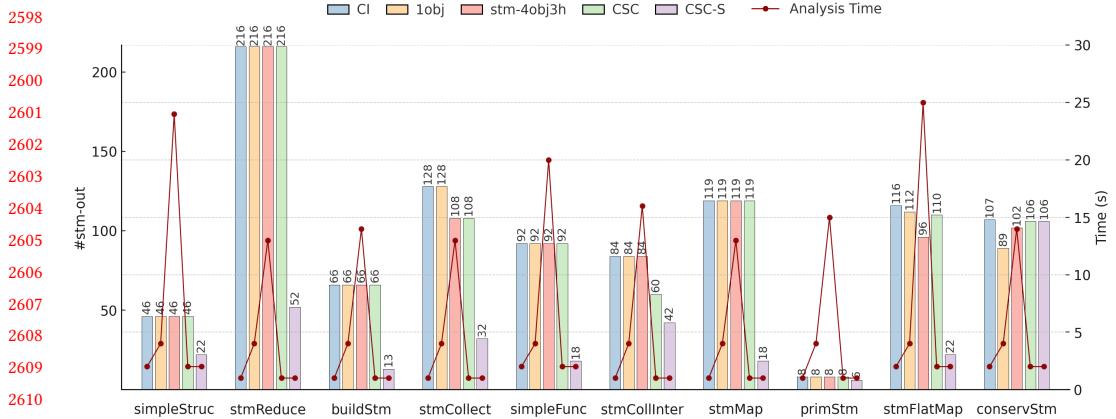


Fig. 33. Evaluation results across 10 micro-benchmarks corresponding to Table 5. Bars (with numbers) indicate #stm-out (left axis, lower is more precise), and red lines show analysis time in seconds (right axis).

CI are not pronounced; more substantial advantages will appear in the evaluation of real-world Java 8 programs (see later subsections). Nonetheless, CUT-SHORTCUT<sup>S</sup> clearly outperforms conventional and selective context-sensitive analyses, achieving an average speedup of 3.0× over 1obj and 12.3× over stm-4obj3h.

*Precision.* The colored bars in Figure 33 present the precision results for #stm-out across all 10 micro-benchmarks. CUT-SHORTCUT<sup>S</sup> consistently achieves the highest precision, except on conservStm, which consists solely of operations conservatively modeled in our approach. The reason Cut-SHORTCUT<sup>S</sup> attains superior precision is because it can distinguish between different stream pipelines—an effect unattainable with conventional or selective context sensitivity in practice. In principle, full call-site sensitivity could yield similar precision, but the deepest selective configuration that succeeded on our machine (3-call-site sensitivity with 2-layer heap abstraction) still failed to match the precision of CUT-SHORTCUT<sup>S</sup> and was far less efficient than stm-4obj3h; thus, it is omitted from Table 5. Overall, across the 10 benchmarks, CUT-SHORTCUT<sup>S</sup> reduces #stm-out by 55.9% on average compared to all four baselines.

While conventional and selective context sensitivity fail to accurately distinguish objects propagated through different streams, they nonetheless yield notable precision gains for other pointers used within or related to pipelines, reflected in improvements in the #app-vpt metric. Summarizing, for #app-vpt, CUT-SHORTCUT<sup>S</sup> achieves average precision improvements of 86.2% over CI, 67.2% over 1obj, 31.9% over stm-4obj3h, and 57.9% over CUT-SHORTCUT.

**8.5.2 Micro-Benchmarks Mimicking Backend Workflows.** In real-world software development, the Java stream API is widely used in backend systems to support declarative, functional-style processing of collections. It enables concise and elegant implementations of common tasks such as filtering, mapping, grouping, and data transformation—operations central to business logic, PO/VO/BO/DTO conversions, and query-result aggregation. However, static analysis of such enterprise backend code remains challenging due to obstacles such as modeling Java/database interactions and identifying application entry points. To enable evaluation, we designed a benchmark suite that mimics backend workflows using a hard-coded in-memory mock database, avoiding the need for actual database connections.

**The workflowX Micro-Benchmarks.** This suite builds on and extends the BSS benchmark [82], originally generated by the S2S tool [64], which translates SQL queries into Java stream code.

Table 6. Efficiency and precision results of context-insensitive analysis (CI), conventional context sensitivity (1obj), selective 4-object sensitivity with 3 context-sensitive heap abstractions on stream-related library code (stm-4obj3h), CUT-SHORTCUT (CSC), and CUT-SHORTCUT<sup>S</sup> (CSC-S) across 5 micro-benchmarks mimicking backend workflows (left half), and their 5 string-outputting variants (right half).

Micro-benchmark	Dyn #stm-out	Static Analysis				Micro-benchmark	Dyn #stm-out	Static Analysis			
		Analysis	Time(s)	#app-vpt	#stm-out			Analysis	Time(s)	#app-vpt	#stm-out
workflow1	65	CI	2	14,433	202	workflow1-str	63	CI	8	89,068	31,326
		1obj	5	12,501	202			1obj	155	67,451	26,544
		stm-4obj3h	16	8,597	202			stm-4obj3h	22	72,620	26,550
		CSC	3	12,297	202			CSC	7	84,066	30,072
		CSC-S	2	3,896	99			CSC-S	7	15,246	85
workflow2	32	CI	3	29,908	357	workflow2-str	12	CI	9	165,794	31,038
		1obj	5	14,210	323			1obj	157	125,906	26,514
		stm-4obj3h	28	10,422	357			stm-4obj3h	58	116,839	26,242
		CSC	3	13,612	357			CSC	7	138,749	29,814
		CSC-S	2	5,021	90			CSC-S	7	11,409	46
workflow3	35	CI	3	47,747	507	workflow3-str	35	CI	10	181,839	31,542
		1obj	5	18,167	507			1obj	170	87,926	26,952
		stm-4obj3h	28	14,496	507			stm-4obj3h	67	84,395	26,712
		CSC	3	18,163	507			CSC	7	95,920	30,228
		CSC-S	3	7,773	228			CSC-S	7	8,251	228
workflow4	44	CI	3	75,642	382	workflow4-str	44	CI	10	276,354	32,052
		1obj	5	36,368	382			1obj	174	162,678	27,474
		stm-4obj3h	115	17,819	382			stm-4obj3h	247	88,176	27,288
		CSC	3	26,970	382			CSC	8	167,770	30,768
		CSC-S	3	11,374	188			CSC-S	8	32,335	188
workflow5	40	CI	4	70,720	290	workflow5-str	82	CI	10	281,520	32,070
		1obj	7	28,580	273			1obj	202	149,626	27,474
		stm-4obj3h	42	18,678	290			stm-4obj3h	102	139,656	27,306
		CSC	3	24,460	290			CSC	8	159,372	30,774
		CSC-S	3	18,783	177			CSC-S	8	90,930	10,184

To adapt BSS for static analysis, we replaced its database connection with an in-memory mock database, modified and expanded the original 8 tasks into 30, then grouped into five benchmarks (workflow-1 through workflow-5, see the left half of Table 6). Each benchmark contains 6 tasks (6 stream pipelines), representing distinct backend business logic tasks. The numbering (1–5) reflects increasing task complexity, from simple querying and mapping to advanced grouping and data transformation. Our mock database consists of 8 tables, each with 5–100 rows, where each row corresponds to one abstract heap object. For each task, rows are queried from the database, used as inputs to stream pipelines, processed according to the specified business logic, and the resulting outputs mimic the objects returned to the application frontend.

For each benchmark, we manually recorded the abstract heap objects pointed to by stream-output variables during dynamic execution (results available in our artifact) and reported their totals in the “Dyn #stm-out” column in Table 6 (the left half). To validate soundness, we confirmed that every heap object observed dynamically at stream-output pointers is also captured by all static analyses, including both CUT-SHORTCUT and CUT-SHORTCUT<sup>S</sup> (achieving 100% recall). This measurement confirms the reliability of the precision improvements obtained from CUT-SHORTCUT<sup>S</sup>.

*The workflowX-str Micro-Benchmarks.* To more closely reflect real-world backend scenarios, we further adapted each benchmark workflow1–workflow5 into a corresponding variant, workflow1-str to workflow5-str, in which the final outputs of the stream pipelines are of type String. This extension is motivated by the ubiquity of string processing in enterprise applications: many

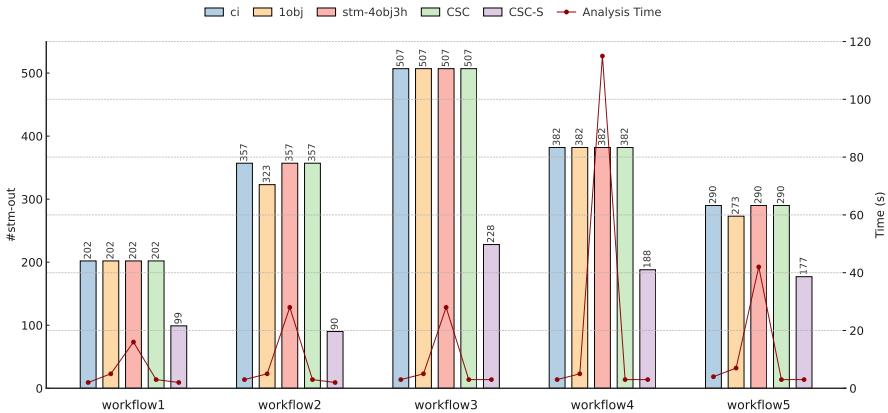


Fig. 34. Evaluation results across the 5 micro-benchmarks mimicking backend workflows (corresponding to the left half of Table 6). Bars (with numbers) indicate #stm-out (left axis, lower is more precise), while red lines show analysis time in seconds (right axis).

database fields are stored as strings, and backend systems frequently convert diverse data types into strings prior to transmission, serialization, or logging. Capturing this scenario is therefore essential for evaluating pointer analyses on realistic workloads.

More importantly, these benchmarks expose a key weakness of conventional analyses. As with other types, when pipelines are String-typed, all compared analyses other than CUT-SHORTCUT<sup>S</sup> merge the input elements of different pipelines, leading to imprecision. For String, however, this problem is especially acute due to the prevalence of string-related pipelines in the JDK 8 library. Critically, this imprecision cannot be mitigated by a trivial post-hoc type filter on stream outputs: once merged, the spurious String objects propagate throughout the analysis, contaminating both library and application-level results. By contrast, CUT-SHORTCUT<sup>S</sup> preserves the separation of pipelines regardless of element type, and thus remains robust in this setting. We note that analogous issues also arise for other generic-typed pipelines, such as those over Object. However, we restrict our evaluation to String-typed pipelines, because they are both pervasive in practice and sufficient to demonstrate the precision degradation of existing approaches. The results for these extended benchmarks are reported in the right half of Table 6.

*Efficiency.* The red lines in Figure 34 illustrate the efficiency advantage of CUT-SHORTCUT<sup>S</sup>. Across the five workflowX benchmarks, its elapsed time is consistently on par with or faster than CI and CUT-SHORTCUT. This efficiency gain arises because CUT-SHORTCUT<sup>S</sup> reduces the amount of spurious points-to information propagated during analysis, while the overhead introduced by its stream-handling extension is negligible. Compared to 1obj, CUT-SHORTCUT<sup>S</sup> achieves an average speedup of 2.1×; compared to stm-4obj3h, the speedup is 14.1×.

For the five workflowX-str benchmarks, the analysis times are shown in the “Time(s)” column in the right half of Table 6. For these benchmarks, we configured the TAI-E framework to distinguish String constants in pointer analysis in order to evaluate precision (whereas the default configuration merges all String constants that are not used in reflection into a single abstract heap object). This configuration substantially increases the analysis time relative to the non-string variants (left half of Table 6). Nevertheless, CUT-SHORTCUT<sup>S</sup> and CUT-SHORTCUT remain the most efficient analyses, achieving on average 1.3× speedup over CI, 23.1× over 1obj, and 10.0× over stm-4obj3h.

*Precision.* The colored bars in Figure 34 show the #stm-out precision results across the five workflowX benchmarks. As expected, CUT-SHORTCUT<sup>S</sup> consistently achieves the highest precision

Table 7. Efficiency and precision results of context-insensitive analysis (CI), conventional context sensitivity (1obj), selective 4-object sensitivity with 3 context-sensitive heap abstractions on stream-related library (stm-4obj3h), CUT-SHORTCUT (CSC), and CUT-SHORTCUT<sup>S</sup> (CSC-S) across 5 real-world stream-intensive programs.

Program	#pipeline	Analysis	Time(s)	#app-vpt	#stm-out
ws4j	23 (25)	ci	6	1,422,873	5,772
		1obj	14	730,580	3,738
		stm-4obj3h	49	860,197	4,301
		CSC	4	374,182	4,209
		CSC-S	3	287,807	1,540
finmath	66 (91)	ci	8	3,181,920	5,148
		1obj	35	923,423	4,495
		stm-4obj3h	116	1,351,276	4,956
		CSC	6	1,127,181	4,678
		CSC-S	6	1,098,442	1,763
jbayes	75 (84)	ci	105	21,828,769	10,873
		1obj	—	—	—
		stm-4obj3h	285	15,812,207	8,774
		CSC	32	14,294,470	10,656
		CSC-S	36	14,205,333	6,440
amazon-sqs	382 (399)	ci	143	45,846,018	404,354
		1obj	730	23,449,094	232,277
		stm-4obj3h	1,157	29,087,035	338,019
		CSC	76	14,673,724	255,103
		CSC-S	70	12,167,567	135,274
mnemonics	14 (24)	ci	5	179,135	14,440
		1obj	122	90,024	12,153
		stm-4obj3h	21	84,788	11,512
		CSC	4	116,763	13,536
		CSC-S	3	108,316	8,939

by distinguishing objects flowing through different stream pipelines, which is beyond the capabilities of the other analyses. Overall, CUT-SHORTCUT<sup>S</sup> improves #stm-out precision by 52.6% on average compared to all four baselines. For #app-vpt, it achieves average improvements of 79.5% over CI, 57.0% over 1obj, 36.0% over stm-4obj3h, and 50.6% over CUT-SHORTCUT.

For the five workflowX-str benchmarks, the precision degradation of the four baseline analyses is substantial, whereas CUT-SHORTCUT<sup>S</sup> retains its high precision. (See the right half of Table 6.) The only exception is workflow5-str, in which 2 of the 6 pipelines include conservatively modeled stream operations; nonetheless, CUT-SHORTCUT<sup>S</sup> still delivers the best precision among all analyses. On average, across these benchmarks, CUT-SHORTCUT<sup>S</sup> improves #stm-out precision by 91.4% compared to the baselines. For #app-vpt, it achieves average improvements of 84.9% over CI, 72.5% over 1obj, 67.7% over stm-4obj3h, and 75.0% over CUT-SHORTCUT.

*Forward Looking.* Recent work has advanced the static analysis of enterprise applications [3, 44, 45], addressing challenges such as identifying application entry points [3], handling database connections [44], and modeling dependency injection, RPC, and message-based communication in microservices [45]. Our workflow-mimicking suite provides evidence of the potential benefits of precise stream handling by emulating enterprise backend workflows, thereby serving as a proxy to demonstrate the extent of precision improvements achievable in real-world scenarios. Nonetheless, integrating CUT-SHORTCUT<sup>S</sup> with these emerging techniques and applying it to real enterprise systems remains non-trivial. We anticipate that such integration would enable highly precise, scalable analyses of enterprise applications in the future.

**8.5.3 Real-World Stream-Intensive Programs.** This benchmark suite consists of 5 real-world Java programs with intensive use of streams. Unlike the workflow-mimicking suite, these are not enterprise applications, making them complementary to the other benchmark suites. The mnemonics

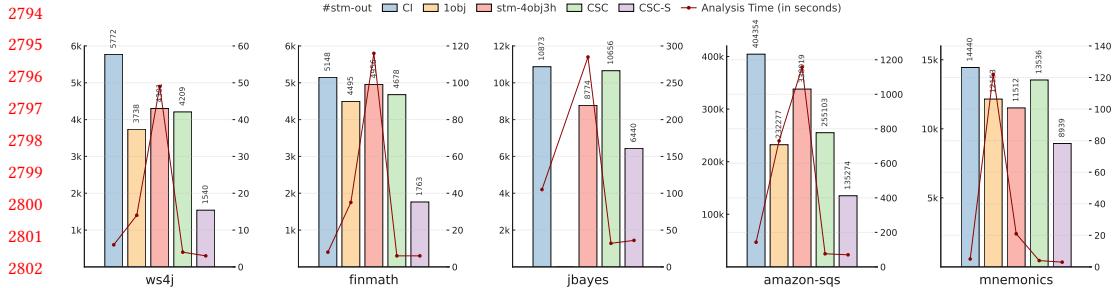


Fig. 35. Evaluation results across 5 real-world programs corresponding to Table 7. Bars (with numbers) indicate #stm-out (left axis, lower is more precise), and red lines show analysis time in seconds (right axis).

benchmark is drawn from the Renaissance suite [58], a widely used benchmark suite targeting advanced Java features. The remaining four benchmarks are selected from Rosales et al. [62], which surveyed Java projects with heavy stream usage. From their collection, we chose four projects that avoid complex framework-specific features (e.g., database queries, Spring AOP) that require additional handling in pointer analysis. Evaluation results on these five benchmarks are reported in Table 7. The “#pipeline” column indicates the number of abstract stream pipelines identified in each program, (numbers in parentheses include conservatively modeled ones).

*Efficiency.* The red lines in Figure 35 illustrate the efficiency advantage of CUT-SHORTCUT<sup>S</sup>. Across all five programs, CUT-SHORTCUT is consistently faster than CI. Compared to CUT-SHORTCUT, CUT-SHORTCUT<sup>S</sup> is even faster for three programs and equally fast in one program; on the remaining one (finmath), it is only slightly slower (by about one second) yet still faster than CI. On average, CUT-SHORTCUT<sup>S</sup> achieves a 1.9× speedup over CI, a 10.4× speedup over 1obj (for the 4 program that 1obj can scale), and an 12.4× speedup over stm-4obj3h.

*Precision.* The colored bars in Figure 35 show the #stm-out precision results across the five programs. As with the previous benchmarks, CUT-SHORTCUT<sup>S</sup> consistently delivers the best precision. While (selective) context sensitivity and CUT-SHORTCUT cannot distinguish objects flowing through different pipelines, they can still yield some precision improvements for #stm-out by reducing the size of the points-to sets of stream inputs. On average, CUT-SHORTCUT<sup>S</sup> improves #stm-out precision by 54.9% over CI, 44.6% over 1obj (on the four programs where 1obj scales), 43.0% over stm-4obj3h, and 47.8% over CUT-SHORTCUT.

For #app-vpt, CUT-SHORTCUT<sup>S</sup> achieves the highest precision for 3 out of 5 programs. Overall, CUT-SHORTCUT<sup>S</sup> improves precision by 55.6% over CI, 11.3% over 1obj (for the four benchmarks on which 1obj succeeded, on two of which 1obj obtained better precision than CUT-SHORTCUT<sup>S</sup>), 20.0% over stm-4obj3h (including one on which stm-4obj3h obtained better precision), and 5.4% over CUT-SHORTCUT. These gains, though substantial, are somewhat smaller than in the synthetic and backend workflow-style benchmarks. The reason is that the five real-world Java programs, while stream-intensive, do not fully capture the backend workflow patterns—such as query transformations or large-scale data aggregation—for which CUT-SHORTCUT<sup>S</sup> offers its greatest advantages.

## 9 Related Work

### 9.1 Context-Sensitivity

Context sensitivity plays an essential role for whole-program Java pointer analysis, and we have discussed some of the related work in earlier sections. The other relevant research is covered below.

In recent years, many selective context-sensitivity approaches have been proposed, which provide different efficiency and precision trade-offs, especially for large and complex Java programs. We can

2843 understand “selective” in two ways: select more effective context elements to distinguish different  
2844 invocations of the same method (*Select context elements*), versus select a set of program methods  
2845 to which a context-sensitive analysis is to be applied (*Select program methods*). Below, we discuss  
2846 these two types of selective context-sensitivity approaches.

2847 *Select Context Elements.* Conventional ( $k$ -limiting) context sensitivity uses  $k$  consecutive context  
2848 elements, e.g.,  $[c_k \dots, c_2, c_1]$  to analyze a method  $m$ , where  $c_1$  is  $m$ ’s call site and  $c_2$  is the call site  
2849 of the method containing  $c_1$ , etc. However, Tan et al. [77] found that this approach may result in  
2850 many context elements that are not useful for improving precision while occupying a portion of the  
2851 limited number of  $k$  slots. Tan et al. [77] presented an approach to recognize such redundant context  
2852 elements by exploiting the so-called object-allocation graph that they proposed. Similarly, Jeon et al.  
2853 [29] developed a machine-learning scheme called context tunneling to select such precision-useless  
2854 context elements for more effective analysis. Furthermore, building on context tunneling, Jeon and  
2855 Oh [31] proposed an interesting approach that transforms object sensitivity to call-site sensitivity,  
2856 and showed that the latter can simulate the former (but not vice versa). In addition, He et al. [24] first  
2857 proposed an IFDS-based [61] approach that identifies “context-independent” objects and “debloats”  
2858 unnecessary context, thereby significantly reducing overhead in object-sensitive pointer analyses  
2859 while retaining almost all precision. Later, they refined this concept by focusing on container-usage  
2860 patterns [21], identifying context-independent objects even more accurately. Recently, Li et al. [37]  
2861 presented generic sensitivity, which preserves generic instantiation sites as key context elements,  
2862 leveraging a type-variable lookup map that is updated during a context-sensitive analysis.  
2863

2864 *Select Program Methods.* Conventional context sensitivity uniformly applies contexts to every  
2865 method in a given program. However, for certain program methods, context sensitivity does not  
2866 help improve precision: it only incurs extra analysis cost. As a result, researchers have proposed  
2867 ways to select a set of methods for which it is (likely to be) beneficial to analyze precisely; the  
2868 analyzer only performs a context-sensitive analysis to those methods, and analyzes the remaining  
2869 methods in a context-insensitive manner. To make a good efficiency/precision trade-off, the overall  
2870 principle is to select the methods that are precision-critical [39] but not scalability-threatening [41].  
2871 To do so, various selection strategies exist: they rely on parameterized heuristics (whose thresholds  
2872 are based on expert experience) [20, 70], machine-learning approaches [28, 30, 32], abstracted  
2873 memory capacity [40], CFL-reachability-based pre-analysis [47], program patterns [39, 41]. Tan  
2874 et al. [76] gave a scheme that supports (i) applying multiple selection strategies—with different  
2875 variants of context sensitivity—to different program parts, combining their precision gains, and  
2876 (ii) iteratively applying multiple strategies to the same program part to jointly enhance precision.  
2877 In addition, context sensitivity can also be applied to selected variables and objects (rather than  
2878 methods) [22, 23, 46], leveraging a pre-analysis in which object reachability is formulated (and  
2879 solved) as a CFL-reachability problem [59] on the so-called pointer-assignment graph.  
2880

2881 *Hybrid Context Sensitivity.* In some pointer analyses, the context elements for the same method  
2882 may vary. Kastrinis and Smaragdakis [34] present a hybrid context-sensitivity approach in which a  
2883 method can be analyzed under both call-site sensitivity and object sensitivity, and in some cases,  
2884 such a combination can lead to more effective results than using a single context type. Thakur and  
2885 Nandivada [80] mix object-sensitivity contexts and level-summarized relevant-value contexts (a  
2886 context abstraction proposed in their earlier work [79]) to analyze each method, and are sometimes  
2887 able to obtain more precise results than with either individual approach alone.

2888 All of the approaches described above rely on the core idea of context sensitivity to replicate  
2889 (a set of) methods and analyze the program elements in them separately under different (types of)  
2890 contexts. Our approach is different in that it simulates the effect of context sensitivity by  
2891

2892 suppressing (i.e., not generating) edges that would cause a loss of precision, and adding shortcut  
 2893 edges directly on the PFG, as explained throughout the paper.

## 2895 9.2 Java Streams

2896 Previous research on Java streams mainly falls into two categories: one for empirical study of  
 2897 stream usage in real-world code [35, 53, 62, 63], and the other for re-writing and optimizing stream  
 2898 performance [5, 51]. In contrast, the static modeling of streams in Cut-SHORTCUT<sup>S</sup> is the first  
 2899 technique that aims at improving the precision of pointer analysis for programs that use Java  
 2900 streams.

## 2902 9.3 Other Related Work

2904 In Section 3.1, we discussed how CUT-SHORTCUT differs from summary-based pointer analysis  
 2905 techniques. However, it is worth noting that CUT-SHORTCUT employ a high-level idea similar to  
 2906 the use of summary edges in interprocedural slicing [27] and the later stages of interprocedural  
 2907 data-flow analysis [61, 65]. In the interprocedural-slicing algorithm proposed by Horwitz et al.  
 2908 [27], when marking vertices that can reach a given input vertex set  $s$ , their algorithm avoids  
 2909 descending into called procedures and instead uses *transitive flow dependence edges* (connecting  
 2910 actual-in to actual-out vertices) to discover paths that reach  $s$  via a procedure call—conceptually  
 2911 similar to how our shortcut edges in the PFG enable precise propagation without following full  
 2912 interprocedural flows. A particularly close parallel is the use of summary relations by Reps et al. [61]  
 2913 and Sharir and Pnueli [65]: once new relations are identified (shortcut edges in CUT-SHORTCUT and  
 2914 CUT-SHORTCUT<sup>S</sup>, summary edges linking invocation arguments to returns in [61], and tabulated  $\phi$   
 2915 functions in [65]), the analysis proceeds using these facts, while method-return facts are effectively  
 2916 ignored.

2917 Although these techniques share the same high-level idea, CUT-SHORTCUT and CUT-SHORTCUT<sup>S</sup>  
 2918 differ in both their target problem (pointer analysis) and concrete methodology. CUT-SHORTCUT is  
 2919 a pattern-based approach, with shortcut-edge generation specifically tailored to each pattern. This  
 2920 design allows CUT-SHORTCUT and CUT-SHORTCUT<sup>S</sup> to operate at a finer granularity by selectively  
 2921 suppressing imprecise flows and introducing precise shortcut flows for a certain method (or  
 2922 sequences of methods) while leaving unrelated parts of that method unchanged. Furthermore, the  
 2923 container-access pattern in CUT-SHORTCUT and the stream modeling in CUT-SHORTCUT<sup>S</sup> enables  
 2924 cross-call-site object-flow routing—i.e., directly connecting the SOURCES of a container/stream  
 2925 instance to its TARGETS—a capability that is difficult to achieve using traditional summarization.  
 2926 Finally, unlike the phased approach of the interprocedural slicing and data-flow analysis, CUT-  
 2927 SHORTCUT and CUT-SHORTCUT<sup>S</sup> are implemented to run on-the-fly alongside the pointer analysis  
 2928 itself, rather than performing a separate stage to compute summary relations.

2929 In addition, Zhang et al. [89] present a hybrid top-down and bottom-up analysis. Top-down and  
 2930 bottom-up approaches are general ideas that many static analyses adopt. Like other mainstream  
 2931 pointer analyses for Java [67], CUT-SHORTCUT also performs in a top-down manner; if we treat  
 2932 the identification of cut and shortcut spots in the PFG for the field access pattern as a separate  
 2933 process, it shares a flavor of bottom-up summarization. However, our pattern-specific imprecise flow  
 2934 identification is on-the-fly, and neither the target problem, nor its underlying concrete methodology  
 2935 of CUT-SHORTCUT is similar to [89].

2936 A different approach to avoiding spurious points-to information during the analysis of Java  
 2937 programs was presented by De and D’Souza [12]. In addition to using context sensitivity, their  
 2938 method provides a flow-sensitive approach to partially performing strong updates to heap objects.  
 2939 They accomplish this task by expressing points-to relations via access paths [16, 33]. Their approach

is different from CUT-SHORTCUT and many other schemes, which express points-to information via PFG-like graphs.

Finally, there are efforts to accelerate context-sensitive Java pointer analysis by merging heap objects [10, 78]; these approaches are orthogonal to ours.

## 10 Conclusion

For the past 20 years, context sensitivity has been considered as virtually the most useful technique for increasing the precision of Java pointer analysis. However, it imposes substantial efficiency costs, especially for large and complex programs. Selective context-sensitivity approaches partially alleviate this issue, but have not solved the efficiency bottleneck: it is hard to select the correct methods that are precision-critical, but do not threaten scalability. As a result, selective context-sensitivity approaches still run the risk of computing and maintaining a large number of contexts to distinguish spurious object flows, which limits their ability to analyze large programs successfully.

To address this dilemma, we developed a fundamentally different approach, called CUT-SHORTCUT, which tries to simulate the effect of context sensitivity without applying contexts. This effect is achieved by suppressing in the pointer flow graph the addition of edges that bring precision loss, and adding precise shortcut edges. We instantiated this high-level principle by exploiting three program patterns, and designing rules based on them in accordance with our principle. Then we formalized it and prove its soundness. We implemented CUT-SHORTCUT on the state-of-the-art Java pointer analysis framework TAI-E. The experimental results are extremely promising: CUT-SHORTCUT consistently matches or outperforms context insensitivity in terms of analysis time across all benchmarks, while obtaining significantly better precision—precision that obtained via context-sensitive approaches in some cases.

Building on this foundation, we extended CUT-SHORTCUT to modern Java language features, focusing on Java streams, which introduce new challenges for pointer analysis. We systematically classified stream operations into categories and developed a static modeling of Java streams that aligns with the principles of CUT-SHORTCUT, yielding the CUT-SHORTCUT<sup>S</sup> extension. This extension improves precision by associating the input sources and output targets of stream pipeline instances, without relying on context sensitivity. Our evaluation on both micro-benchmarks and real-world programs shows that CUT-SHORTCUT<sup>S</sup> enhances precision while preserving scalability and efficiency—consistently outperforming or matching context-insensitive analyses—and successfully handles stream-intensive programs where (selective) context-sensitive techniques fail to provide analyses that are both precise and scalable.

Together, these contributions demonstrate the flexibility and scalability of the CUT-SHORTCUT approach, both for classic Java programs and for modern constructs like streams. Given the encouraging outcomes, more program patterns or insights are expected to be explored on top of our approach, and we hope that our approach could provide new perspectives for developing more effective pointer analysis for Java in the future.

## Acknowledgments

This work was supported, in part, by a gift from Rajiv and Ritu Batra, and by NSF under grants CCF-2211968 and CCF-2212558. Any opinions, findings, and conclusions or recommendations expressed in this publication are those of the authors, and do not necessarily reflect the views of the sponsoring entities.

## References

- [1] 2023. <https://newrelic.com/resources/report/2023-state-of-the-java-ecosystem>

- [2] Lars Ole Andersen. 1994. *Program analysis and specialization for the C programming language*. Ph. D. Dissertation. University of Copenhagen.
- [3] Anastasios Antoniadis, Nikos Filippakis, Paddy Krishnan, Raghavendra Ramesh, Nicholas Allen, and Yannis Smaragdakis. 2020. Static Analysis of Java Enterprise Applications: Frameworks and Caches, the Elephants in the Room. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation* (London, UK) (*PLDI 2020*). Association for Computing Machinery, New York, NY, USA, 794–807. [doi:10.1145/3385412.3386026](https://doi.org/10.1145/3385412.3386026)
- [4] Steven Arzt, Siegfried Rasthofer, Christian Fritz, Eric Bodden, Alexandre Bartel, Jacques Klein, Yves Le Traon, Damien Octeau, and Patrick McDaniel. 2014. FlowDroid: Precise Context, Flow, Field, Object-Sensitive and Lifecycle-Aware Taint Analysis for Android Apps. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation* (Edinburgh, United Kingdom) (*PLDI ’14*). Association for Computing Machinery, New York, NY, USA, 259–269. [doi:10.1145/2594291.2594299](https://doi.org/10.1145/2594291.2594299)
- [5] Matteo Basso, Filippo Schiavio, Andrea Rosà, and Walter Binder. 2022. Optimizing Parallel Java Streams. In *2022 26th International Conference on Engineering of Complex Computer Systems (ICECCS)*. 23–32. [doi:10.1109/ICECCS54210.2022.00012](https://doi.org/10.1109/ICECCS54210.2022.00012)
- [6] S. M. Blackburn, R. Garner, C. Hoffman, A. M. Khan, K. S. McKinley, R. Bentzur, A. Diwan, D. Feinberg, D. Frampton, S. Z. Guyer, M. Hirzel, A. Hosking, M. Jump, H. Lee, J. E. B. Moss, A. Phansalkar, D. Stefanović, T. VanDrunen, D. von Dincklage, and B. Wiedermann. 2006. The DaCapo Benchmarks: Java Benchmarking Development and Analysis. In *OOPSLA ’06: Proceedings of the 21st annual ACM SIGPLAN conference on Object-Oriented Programming, Systems, Languages, and Applications* (Portland, OR, USA). ACM Press, New York, NY, USA, 169–190. [doi:10.1145/1167473.1167488](https://doi.org/10.1145/1167473.1167488)
- [7] Martin Bravenboer and Yannis Smaragdakis. 2009. Strictly declarative specification of sophisticated points-to analyses. In *Proceedings of the 24th ACM SIGPLAN Conference on Object Oriented Programming Systems Languages and Applications* (Orlando, Florida, USA) (*OOPSLA ’09*). Association for Computing Machinery, New York, NY, USA, 243–262. [doi:10.1145/1640089.1640108](https://doi.org/10.1145/1640089.1640108)
- [8] Yuandao Cai, Peisen Yao, and Charles Zhang. 2021. Canary: Practical Static Detection of Inter-Thread Value-Flow Bugs. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation* (Virtual, Canada) (*PLDI 2021*). Association for Computing Machinery, New York, NY, USA, 1126–1140. [doi:10.1145/3453483.3454099](https://doi.org/10.1145/3453483.3454099)
- [9] Satish Chandra, Stephen J. Fink, and Manu Sridharan. 2009. Snufflebug: A Powerful Approach to Weakest Preconditions. In *Proceedings of the 30th ACM SIGPLAN Conference on Programming Language Design and Implementation* (Dublin, Ireland) (*PLDI ’09*). Association for Computing Machinery, New York, NY, USA, 363–374. [doi:10.1145/1542476.1542517](https://doi.org/10.1145/1542476.1542517)
- [10] Yifan Chen, Chenyang Yang, Xin Zhang, Yingfei Xiong, Hao Tang, Xiaoyin Wang, and Lu Zhang. 2021. Accelerating Program Analyses in Datalog by Merging Library Facts. In *Static Analysis: 28th International Symposium, SAS 2021, Chicago, IL, USA, October 17–19, 2021, Proceedings* (Chicago, IL, USA). Springer-Verlag, Berlin, Heidelberg, 77–101. [doi:10.1007/978-3-030-88806-0\\_4](https://doi.org/10.1007/978-3-030-88806-0_4)
- [11] Ben-Chung Cheng and Wen-Mei W. Hwu. 2000. Modular interprocedural pointer analysis using access paths: design, implementation, and evaluation. In *Proceedings of the ACM SIGPLAN 2000 Conference on Programming Language Design and Implementation* (Vancouver, British Columbia, Canada) (*PLDI ’00*). Association for Computing Machinery, New York, NY, USA, 57–69. [doi:10.1145/349299.349311](https://doi.org/10.1145/349299.349311)
- [12] Arnab De and Deepak D’Souza. 2012. Scalable Flow-Sensitive Pointer Analysis for Java with Strong Updates. In *ECOOP 2012 – Object-Oriented Programming*, James Noble (Ed.). Springer Berlin Heidelberg, Berlin, Heidelberg, 665–687.
- [13] Jeffrey Dean, David Grove, and Craig Chambers. 1995. Optimization of Object-Oriented Programs Using Static Class Hierarchy Analysis. In *ECOOP’95 – Object-Oriented Programming, 9th European Conference, Århus, Denmark, August 7–11, 1995*, Mario Tokoro and Remo Pareschi (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 77–101.
- [14] Pratik Fegade and Christian Wimmer. 2020. Scalable Pointer Analysis of Data Structures Using Semantic Models. In *Proceedings of the 29th International Conference on Compiler Construction* (San Diego, CA, USA) (*CC 2020*). Association for Computing Machinery, New York, NY, USA, 39–50. [doi:10.1145/3377555.3377885](https://doi.org/10.1145/3377555.3377885)
- [15] Yu Feng, Xinyu Wang, Isil Dillig, and Thomas Dillig. 2015. Bottom-Up Context-Sensitive Pointer Analysis for Java. In *Programming Languages and Systems*, Xinyu Feng and Sungwoo Park (Eds.). Springer International Publishing, Cham, 465–484.
- [16] Stephen J. Fink, Eran Yahav, Nurit Dor, G. Ramalingam, and Emmanuel Geay. 2008. Effective typestate verification in the presence of aliasing. *ACM Trans. Softw. Eng. Methodol.* 17, 2, Article 9 (May 2008), 34 pages. [doi:10.1145/1348250.1348255](https://doi.org/10.1145/1348250.1348255)
- [17] Philip J. Fleming and John J. Wallace. 1986. How not to lie with statistics: the correct way to summarize benchmark results. *Commun. ACM* 29, 3 (March 1986), 218–221. [doi:10.1145/5666.5673](https://doi.org/10.1145/5666.5673)
- [18] George Fourtounis and Yannis Smaragdakis. 2019. Deep Static Modeling of invokedynamic. In *33rd European Conference on Object-Oriented Programming (ECOOP 2019) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 134)*, Alastair F. Donaldson (Ed.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 15:1–15:28. [doi:10.4230/LIPIcs.ECOOP.2019.15](https://doi.org/10.4230/LIPIcs.ECOOP.2019.15)

- 3039 [19] Neville Grech and Yannis Smaragdakis. 2017. P/Taint: unified points-to and taint analysis. *Proc. ACM Program. Lang.* 1, OOPSLA, Article 102 (Oct. 2017), 28 pages. [doi:10.1145/3133926](https://doi.org/10.1145/3133926)
- 3040 [20] Behnaz Hassanshahi, Raghavendra Kagalavadi Ramesh, Padmanabhan Krishnan, Bernhard Scholz, and Yi Lu. 2017. An  
3041 Efficient Tunable Selective Points-to Analysis for Large Codebases. In *Proceedings of the 6th ACM SIGPLAN International*  
3042 *Workshop on State Of the Art in Program Analysis* (Barcelona, Spain) (SOAP 2017). Association for Computing Machinery,  
3043 New York, NY, USA, 13–18. [doi:10.1145/3088515.3088519](https://doi.org/10.1145/3088515.3088519)
- 3044 [21] Dongjie He, Yujiang Gui, Wei Li, Yonggang Tao, Changwei Zou, Yulei Sui, and Jingling Xue. 2023. A Container-Usage-  
3045 Pattern-Based Context Debloating Approach for Object-Sensitive Pointer Analysis. *Proc. ACM Program. Lang.* 7,  
3046 OOPSLA2, Article 256 (Oct. 2023), 30 pages. [doi:10.1145/3622832](https://doi.org/10.1145/3622832)
- 3047 [22] Dongjie He, Jingbo Lu, Yaoqing Gao, and Jingling Xue. 2021. Accelerating Object-Sensitive Pointer Analysis by  
3048 Exploiting Object Containment and Reachability. In *35th European Conference on Object-Oriented Programming (ECOOP*  
3049 *2021) (Leibniz International Proceedings in Informatics (LIPIcs)*, Vol. 194), Anders Møller and Manu Sridharan (Eds.).  
3050 Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 16:1–16:31. [doi:10.4230/LIPIcs.ECOOP.2021.16](https://doi.org/10.4230/LIPIcs.ECOOP.2021.16)
- 3051 [23] Dongjie He, Jingbo Lu, Yaoqing Gao, and Jingling Xue. 2023. Selecting Context-Sensitivity Modularly for Accelerating  
3052 Object-Sensitive Pointer Analysis. *IEEE Transactions on Software Engineering* 49, 2 (2023), 719–742. [doi:10.1109/TSE.2022.3162236](https://doi.org/10.1109/TSE.2022.3162236)
- 3053 [24] Dongjie He, Jingbo Lu, and Jingling Xue. 2023. IFDS-based Context Debloating for Object-Sensitive Pointer Analysis.  
3054 *ACM Trans. Softw. Eng. Methodol.* 32, 4, Article 101 (May 2023), 44 pages. [doi:10.1145/3579641](https://doi.org/10.1145/3579641)
- 3055 [25] Dongjie He, Jingbo Lu, and Jingling Xue. 2024. A CFL-Reachability Formulation of Callsite-Sensitive Pointer Analysis  
3056 with Built-In On-The-Fly Call Graph Construction. In *38th European Conference on Object-Oriented Programming*  
3057 (*ECOOP 2024) (Leibniz International Proceedings in Informatics (LIPIcs)*, Vol. 313), Jonathan Aldrich and Guido Salvaneschi  
3058 (Eds.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 18:1–18:29. [doi:10.4230/LIPIcs.ECOOP.2024.18](https://doi.org/10.4230/LIPIcs.ECOOP.2024.18)
- 3059 [26] Kihong Heo, Hakjoo Oh, and Kwangkeun Yi. 2017. Selective conjunction of context-sensitivity and octagon domain  
3060 toward scalable and precise global static analysis. *Softw. Pract. Exp.* 47, 11 (2017), 1677–1705. [doi:10.1002/spe.2493](https://doi.org/10.1002/spe.2493)
- 3061 [27] Susan Horwitz, Thomas W. Reps, and David W. Binkley. 1990. Interprocedural Slicing Using Dependence Graphs.  
3062 *ACM Trans. Program. Lang. Syst.* 12, 1 (1990), 26–60. [doi:10.1145/77606.77608](https://doi.org/10.1145/77606.77608)
- 3063 [28] Minseok Jeon, Sehun Jeong, Sungdeok Cha, and Hakjoo Oh. 2019. A Machine-Learning Algorithm with Disjunctive  
3064 Model for Data-Driven Program Analysis. *ACM Trans. Program. Lang. Syst.* 41, 2, Article 13 (jun 2019), 41 pages.  
3065 [doi:10.1145/3293607](https://doi.org/10.1145/3293607)
- 3066 [29] Minseok Jeon, Sehun Jeong, and Hakjoo Oh. 2018. Precise and Scalable Points-to Analysis via Data-Driven Context  
3067 Tunneling. *Proc. ACM Program. Lang.* 2, OOPSLA, Article 140 (oct 2018), 29 pages. [doi:10.1145/3276510](https://doi.org/10.1145/3276510)
- 3068 [30] Minseok Jeon, Myungho Lee, and Hakjoo Oh. 2020. Learning Graph-Based Heuristics for Pointer Analysis without  
3069 Handcrafting Application-Specific Features. *Proc. ACM Program. Lang.* 4, OOPSLA, Article 179 (nov 2020), 30 pages.  
3070 [doi:10.1145/3428247](https://doi.org/10.1145/3428247)
- 3071 [31] Minseok Jeon and Hakjoo Oh. 2022. Return of CFA: Call-Site Sensitivity Can Be Superior to Object Sensitivity Even  
3072 for Object-Oriented Programs. *Proc. ACM Program. Lang.* 6, POPL, Article 58 (jan 2022), 29 pages. [doi:10.1145/3498720](https://doi.org/10.1145/3498720)
- 3073 [32] Sehun Jeong, Minseok Jeon, Sungdeok Cha, and Hakjoo Oh. 2017. Data-driven context-sensitivity for points-to  
3074 analysis. *Proc. ACM Program. Lang.* 1, OOPSLA, Article 100 (Oct. 2017), 28 pages. [doi:10.1145/3133924](https://doi.org/10.1145/3133924)
- 3075 [33] Vini Kanvar and Uday P. Khedker. 2016. Heap Abstractions for Static Analysis. *ACM Comput. Surv.* 49, 2, Article 29  
3076 (June 2016), 47 pages. [doi:10.1145/2931098](https://doi.org/10.1145/2931098)
- 3077 [34] George Kastrinis and Yannis Smaragdakis. 2013. Hybrid context-sensitivity for points-to analysis. In *ACM SIGPLAN*  
3078 *Conference on Programming Language Design and Implementation, PLDI ’13, Seattle, WA, USA, June 16–19, 2013*, Hans  
3079 Juergen Boehm and Cormac Flanagan (Eds.). ACM, 423–434. [doi:10.1145/2491956.2462191](https://doi.org/10.1145/2491956.2462191)
- 3080 [35] Raffi Khatchadourian, Yiming Tang, Mehdi Bagherzadeh, and Baishakhi Ray. 2020. An Empirical Study on the Use and  
3081 Misuse of Java 8 Streams. In *Fundamental Approaches to Software Engineering*, Heike Wehrheim and Jordi Cabot (Eds.),  
3082 Springer International Publishing, Cham, 97–118.
- 3083 [36] Ondřej Lhoták and Laurie Hendren. 2006. Context-Sensitive Points-to Analysis: Is It Worth It?. In *Compiler Construction*,  
3084 Alan Mycroft and Andreas Zeller (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 47–64.
- 3085 [37] Haofeng Li, Jie Lu, Haining Meng, Liqing Cao, Yongheng Huang, Lian Li, and Lin Gao. 2022. Generic sensitivity:  
3086 customizing context-sensitive pointer analysis for generics. In *Proceedings of the 30th ACM Joint European Software*  
3087 *Engineering Conference and Symposium on the Foundations of Software Engineering* (Singapore, Singapore) (ESEC/FSE  
3088 2022). Association for Computing Machinery, New York, NY, USA, 1110–1121. [doi:10.1145/3540250.3549122](https://doi.org/10.1145/3540250.3549122)
- 3089 [38] Haofeng Li, Chenghang Shi, Jie Lu, Lian Li, and Zixuan Zhao. 2025. Module-Aware Context Sensitive Pointer Analysis  
3090 . In *2025 IEEE/ACM 47th International Conference on Software Engineering (ICSE)*. IEEE Computer Society, Los Alamitos,  
3091 CA, USA, 1819–1831. [doi:10.1109/ICSE55347.2025.00227](https://doi.org/10.1109/ICSE55347.2025.00227)

- 3088 [39] Yue Li, Tian Tan, Anders Møller, and Yannis Smaragdakis. 2018. Precision-guided Context Sensitivity for Pointer  
 3089 Analysis. *Proc. ACM Program. Lang.* 2, OOPSLA, Article 141 (Oct. 2018), 29 pages. [doi:10.1145/3276511](https://doi.org/10.1145/3276511)
- 3090 [40] Yue Li, Tian Tan, Anders Møller, and Yannis Smaragdakis. 2018. Scalability-first pointer analysis with self-tuning  
 3091 context-sensitivity. In *Proceedings of the 2018 26th ACM Joint Meeting on European Software Engineering Conference  
 3092 and Symposium on the Foundations of Software Engineering* (Lake Buena Vista, FL, USA) (ESEC/FSE 2018). Association  
 3093 for Computing Machinery, New York, NY, USA, 129–140. [doi:10.1145/3236024.3236041](https://doi.org/10.1145/3236024.3236041)
- 3094 [41] Yue Li, Tian Tan, Anders Møller, and Yannis Smaragdakis. 2020. A Principled Approach to Selective Context Sensitivity  
 3095 for Pointer Analysis. *ACM Trans. Program. Lang. Syst.* 42, 2, Article 10 (may 2020), 40 pages. [doi:10.1145/3381915](https://doi.org/10.1145/3381915)
- 3096 [42] Yue Li, Tian Tan, and Jingling Xue. 2019. Understanding and Analyzing Java Reflection. *ACM Trans. Softw. Eng.  
 3097 Methodol.* 28, 2, Article 7 (feb 2019), 50 pages. [doi:10.1145/3295739](https://doi.org/10.1145/3295739)
- 3098 [43] Yue Li, Tian Tan, Yifei Zhang, and Jingling Xue. 2016. Program Tailoring: Slicing by Sequential Criteria. In *30th  
 3099 European Conference on Object-Oriented Programming (ECOOP 2016) (Leibniz International Proceedings in Informatics  
 3100 (LIPIcs), Vol. 56*, Shriram Krishnamurthi and Benjamin S. Lerner (Eds.). Schloss Dagstuhl–Leibniz-Zentrum fuer  
 3101 Informatik, Dagstuhl, Germany, 15:1–15:27. [doi:10.4230/LIPIcs.ECOOP.2016.15](https://doi.org/10.4230/LIPIcs.ECOOP.2016.15)
- 3102 [44] Yufei Liang, Teng Zhang, Ganlin Li, Tian Tan, Chang Xu, Chun Cao, Xiaoxing Ma, and Yue Li. 2025. Pointer Analysis for  
 3103 Database-Backed Applications. *Proc. ACM Program. Lang.* 9, PLDI, Article 204 (June 2025), 25 pages. [doi:10.1145/3729307](https://doi.org/10.1145/3729307)
- 3104 [45] Yufei Liang, Teng Zhang, Ganlin Li, Tian Tan, Chang Xu, Chun Cao, Xiaoxing Ma, and Yue Li. 2025. Pointer Analysis for  
 3105 Database-Backed Applications. *Proc. ACM Program. Lang.* 9, PLDI, Article 204 (June 2025), 25 pages. [doi:10.1145/3729307](https://doi.org/10.1145/3729307)
- 3106 [46] Jingbo Lu, Dongjie He, and Jingling Xue. 2021. Eagle: CFL-Reachability-Based Precision-Preserving Acceleration of  
 3107 Object-Sensitive Pointer Analysis with Partial Context Sensitivity. *ACM Trans. Softw. Eng. Methodol.* 30, 4, Article 46  
 3108 (jul 2021), 46 pages. [doi:10.1145/3450492](https://doi.org/10.1145/3450492)
- 3109 [47] Jingbo Lu, Dongjie He, and Jingling Xue. 2021. Selective Context-Sensitivity for k-CFA with CFL-Reachability. In  
 3110 *Static Analysis: 28th International Symposium, SAS 2021, Chicago, IL, USA, October 17–19, 2021, Proceedings* (Chicago, IL,  
 3111 USA). Springer-Verlag, Berlin, Heidelberg, 261–285. [doi:10.1007/978-3-030-88806-0\\_13](https://doi.org/10.1007/978-3-030-88806-0_13)
- 3112 [48] Wenjie Ma, Shengyuan Yang, Tian Tan, Xiaoxing Ma, Chang Xu, and Yue Li. 2023. Context Sensitivity without  
 3113 Contexts: A Cut-Shortcut Approach to Fast and Precise Pointer Analysis. *Proc. ACM Program. Lang.* 7, PLDI, Article  
 3114 128 (June 2023), 26 pages. [doi:10.1145/3591242](https://doi.org/10.1145/3591242)
- 3115 [49] Ana Milanova, Atanas Rountev, and Barbara G. Ryder. 2002. Parameterized Object Sensitivity for Points-to and  
 3116 Side-Effect Analyses for Java. In *Proceedings of the 2002 ACM SIGSOFT International Symposium on Software Testing  
 3117 and Analysis* (Roma, Italy) (ISSTA ’02). Association for Computing Machinery, New York, NY, USA, 1–11. [doi:10.1145/566172.566174](https://doi.org/10.1145/566172.566174)
- 3118 [50] Ana Milanova, Atanas Rountev, and Barbara G. Ryder. 2005. Parameterized Object Sensitivity for Points-to Analysis  
 3119 for Java. *ACM Trans. Softw. Eng. Methodol.* 14, 1 (jan 2005), 1–41. [doi:10.1145/1044834.1044835](https://doi.org/10.1145/1044834.1044835)
- 3120 [51] Anders Møller and Oskar Haarklou Veileborg. 2020. Eliminating abstraction overhead of Java stream pipelines  
 3121 using ahead-of-time program optimization. *Proc. ACM Program. Lang.* 4, OOPSLA, Article 168 (Nov. 2020), 29 pages.  
 3122 [doi:10.1145/3428236](https://doi.org/10.1145/3428236)
- 3123 [52] Mayur Naik, Alex Aiken, and John Whaley. 2006. Effective Static Race Detection for Java. In *Proceedings of the 27th  
 3124 ACM SIGPLAN Conference on Programming Language Design and Implementation* (Ottawa, Ontario, Canada) (PLDI ’06).  
 3125 Association for Computing Machinery, New York, NY, USA, 308–319. [doi:10.1145/1133981.1134018](https://doi.org/10.1145/1133981.1134018)
- 3126 [53] Joshua Nostas, Juan Pablo Sandoval Alcocer, Diego Elias Costa, and Alexandre Bergel. 2021. How Do Developers  
 3127 Use the Java Stream API?. In *Computational Science and Its Applications – ICCSA 2021*, Osvaldo Gervasi, Beniamino  
 3128 Murgante, Sanjay Misra, Chiara Garau, Ivan Blećić, David Taniar, Bernady O. Apduhan, Ana Maria A. C. Rocha,  
 3129 Eufemia Tarantino, and Carmelo Maria Torre (Eds.). Springer International Publishing, Cham, 323–335.
- 3130 [54] Erik M. Nystrom, Hong-Seok Kim, and Wen-mei W. Hwu. 2004. Bottom-Up and Top-Down Context-Sensitive Summary-  
 3131 Based Pointer Analysis. In *Static Analysis*, Roberto Giacobazzi (Ed.). Springer Berlin Heidelberg, Berlin, Heidelberg,  
 3132 165–180.
- 3133 [55] Hakjoo Oh, Wonchan Lee, Kihong Heo, Hongseok Yang, and Kwangkeun Yi. 2014. Selective Context-Sensitivity  
 3134 Guided by Impact Pre-Analysis. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design  
 3135 and Implementation* (Edinburgh, United Kingdom) (PLDI ’14). Association for Computing Machinery, New York, NY,  
 3136 USA, 475–484. [doi:10.1145/2594291.2594318](https://doi.org/10.1145/2594291.2594318)
- 3137 [56] Hakjoo Oh, Wonchan Lee, Kihong Heo, Hongseok Yang, and Kwangkeun Yi. 2015. Selective X-Sensitive Analysis Guided  
 3138 by Impact Pre-Analysis. *ACM Trans. Program. Lang. Syst.* 38, 2, Article 6 (dec 2015), 45 pages. [doi:10.1145/2821504](https://doi.org/10.1145/2821504)
- 3139 [57] Michael Pradel, Ciera Jaspan, Jonathan Aldrich, and Thomas R. Gross. 2012. Statically checking API protocol  
 3140 conformance with mined multi-object specifications. In *2012 34th International Conference on Software Engineering  
 3141 (ICSE)*. 925–935. [doi:10.1109/ICSE.2012.6227127](https://doi.org/10.1109/ICSE.2012.6227127)
- 3142 [58] Aleksandar Prokopec, Andrea Rosà, David Leopoldseder, Gilles Duboscq, Petr Tůma, Martin Studener, Lubomír Bulej,  
 3143 Yudi Zheng, Alex Villazón, Doug Simon, Thomas Würthinger, and Walter Binder. 2019. Renaissance: benchmarking  
 3144

- 3137 suite for parallel applications on the JVM. In *Proceedings of the 40th ACM SIGPLAN Conference on Programming*  
3138 *Language Design and Implementation* (Phoenix, AZ, USA) (PLDI 2019). Association for Computing Machinery, New  
3139 York, NY, USA, 31–47. doi:[10.1145/3314221.3314637](https://doi.org/10.1145/3314221.3314637)
- 3140 [59] Thomas Reps. 1998. Program analysis via graph reachability<sup>1</sup>An abbreviated version of this paper appeared as an  
3141 invited paper in the Proceedings of the 1997 International Symposium on Logic Programming [84].1. *Information and*  
3142 *Software Technology* 40, 11 (1998), 701–726. doi:[10.1016/S0950-5849\(98\)00093-7](https://doi.org/10.1016/S0950-5849(98)00093-7)
- 3143 [60] Thomas Reps. 2000. Undecidability of context-sensitive data-dependence analysis. *ACM Trans. Program. Lang. Syst.*  
3144 22, 1 (Jan. 2000), 162–186. doi:[10.1145/345099.345137](https://doi.org/10.1145/345099.345137)
- 3145 [61] Thomas Reps, Susan Horwitz, and Mooly Sagiv. 1995. Precise interprocedural dataflow analysis via graph reachability.  
3146 In *Proceedings of the 22nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (San Francisco,  
California, USA) (POPL '95). Association for Computing Machinery, New York, NY, USA, 49–61. doi:[10.1145/199448.199462](https://doi.org/10.1145/199448.199462)
- 3147 [62] Eduardo Rosales, Matteo Bassو, Andrea Rosà, and Walter Binder. 2023. Large-scale characterization of Java streams.  
3148 *Software: Practice and Experience* 53, 9 (2023), 1763–1792.
- 3149 [63] Eduardo Rosales, Andrea Rosà, Matteo Bassо, Alex Villazón, Adriana Orellana, Ángel Zenteno, Jhon Rivero, and Walter  
3150 Binder. 2022. Characterizing Java Streams in the Wild. In *2022 26th International Conference on Engineering of Complex*  
3151 *Computer Systems* (ICECCS). 143–152. doi:[10.1109/ICECCS54210.2022.00025](https://doi.org/10.1109/ICECCS54210.2022.00025)
- 3152 [64] Filippo Schiavio, Andrea Rosà, and Walter Binder. 2022. SQL to Stream with S2S: An Automatic Benchmark Generator  
3153 for the Java Stream API. In *Proceedings of the 21st ACM SIGPLAN International Conference on Generative Programming:*  
3154 *Concepts and Experiences* (Auckland, New Zealand) (GPCE 2022). Association for Computing Machinery, New York,  
3155 NY, USA, 179–186. doi:[10.1145/3564719.3568699](https://doi.org/10.1145/3564719.3568699)
- 3156 [65] Micha Sharir and Amir Pnueli. 1981. *Two approaches to interprocedural data flow analysis*. Prentice-Hall, Chapter 7,  
3157 189–234.
- 3158 [66] Olin Grigsby Shivers. 1991. *Control-flow analysis of higher-order languages or taming lambda*. Carnegie Mellon  
3159 University.
- 3160 [67] Yannis Smaragdakis and George Balatsouras. 2015. Pointer Analysis. *Found. Trends Program. Lang.* 2, 1 (April 2015),  
3161 1–69. doi:[10.1561/2500000014](https://doi.org/10.1561/2500000014)
- 3162 [68] Yannis Smaragdakis, George Balatsouras, George Kastrinis, and Martin Bravenboer. 2015. More Sound Static Handling  
3163 of Java Reflection. In *Programming Languages and Systems*, Xinyu Feng and Sungwoo Park (Eds.). Springer International  
3164 Publishing, Cham, 485–503.
- 3165 [69] Yannis Smaragdakis, Martin Bravenboer, and Ondrej Lhoták. 2011. Pick Your Contexts Well: Understanding Object-  
3166 Sensitivity. In *Proceedings of the 38th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming*  
3167 *Languages* (Austin, Texas, USA) (POPL '11). Association for Computing Machinery, New York, NY, USA, 17–30.  
3168 doi:[10.1145/1926385.1926390](https://doi.org/10.1145/1926385.1926390)
- 3169 [70] Yannis Smaragdakis, George Kastrinis, and George Balatsouras. 2014. Introspective Analysis: Context-Sensitivity, across  
3170 the Board. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*  
3171 (Edinburgh, United Kingdom) (PLDI '14). Association for Computing Machinery, New York, NY, USA, 485–495.  
3172 doi:[10.1145/2594291.2594320](https://doi.org/10.1145/2594291.2594320)
- 3173 [71] Manu Sridharan and Rastislav Bodík. 2006. Refinement-based context-sensitive points-to analysis for Java. In *Proceed-  
3174 ings of the 27th ACM SIGPLAN Conference on Programming Language Design and Implementation* (Ottawa, Ontario,  
3175 Canada) (PLDI '06). Association for Computing Machinery, New York, NY, USA, 387–400. doi:[10.1145/1133981.1134027](https://doi.org/10.1145/1133981.1134027)
- 3176 [72] Manu Sridharan, Satish Chandra, Julian Dolby, Stephen J. Fink, and Eran Yahav. 2013. *Alias Analysis for Object-Oriented*  
3177 *Programs*. Springer Berlin Heidelberg, Berlin, Heidelberg, 196–232. doi:[10.1007/978-3-642-36946-9\\_8](https://doi.org/10.1007/978-3-642-36946-9_8)
- 3178 [73] Manu Sridharan, Stephen J. Fink, and Rastislav Bodík. 2007. Thin Slicing. In *Proceedings of the 28th ACM SIGPLAN*  
3179 *Conference on Programming Language Design and Implementation* (San Diego, California, USA) (PLDI '07). Association  
3180 for Computing Machinery, New York, NY, USA, 112–122. doi:[10.1145/1250734.1250748](https://doi.org/10.1145/1250734.1250748)
- 3181 [74] Manu Sridharan, Denis Gopan, Lexin Shan, and Rastislav Bodík. 2005. Demand-driven points-to analysis for Java. In *Proceed-  
3182 ings of the 20th Annual ACM SIGPLAN Conference on Object-Oriented Programming, Systems, Languages, and*  
3183 *Applications* (San Diego, CA, USA) (OOPSLA '05). Association for Computing Machinery, New York, NY, USA, 59–76.  
3184 doi:[10.1145/1094811.1094817](https://doi.org/10.1145/1094811.1094817)
- 3185 [75] Tian Tan and Yue Li. 2023. Tai-e: A Developer-Friendly Static Analysis Framework for Java by Harnessing the  
3186 Good Designs of Classics. In *Proceedings of the 32nd ACM SIGSOFT International Symposium on Software Testing and*  
3187 *Analysis* (Seattle, WA, USA) (ISSTA 2023). Association for Computing Machinery, New York, NY, USA, 1093–1105.  
3188 doi:[10.1145/3597926.3598120](https://doi.org/10.1145/3597926.3598120)
- 3189 [76] Tian Tan, Yue Li, Xiaoxing Ma, Chang Xu, and Yannis Smaragdakis. 2021. Making Pointer Analysis More Precise by  
3190 Unleashing the Power of Selective Context Sensitivity. *Proc. ACM Program. Lang.* 5, OOPSLA, Article 147 (oct 2021),  
3191 27 pages. doi:[10.1145/3485524](https://doi.org/10.1145/3485524)

- 3186 [77] Tian Tan, Yue Li, and Jingling Xue. 2016. Making k-Object-Sensitive Pointer Analysis More Precise with Still k-Limiting.  
3187 In *Static Analysis - 23rd International Symposium, SAS 2016, Edinburgh, UK, September 8–10, 2016, Proceedings (Lecture  
3188 Notes in Computer Science, Vol. 9837)*, Xavier Rival (Ed.). Springer, 489–510. doi:[10.1007/978-3-662-53413-7\\_24](https://doi.org/10.1007/978-3-662-53413-7_24)
- 3189 [78] Tian Tan, Yue Li, and Jingling Xue. 2017. Efficient and precise points-to analysis: modeling the heap by merging  
3190 equivalent automata. In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and  
3191 Implementation* (Barcelona, Spain) (PLDI 2017). Association for Computing Machinery, New York, NY, USA, 278–291.  
3192 doi:[10.1145/3062341.3062360](https://doi.org/10.1145/3062341.3062360)
- 3193 [79] Manas Thakur and V. Krishna Nandivada. 2019. Compare Less, Defer More: Scaling Value-Contexts Based Whole-  
3194 Program Heap Analyses. In *Proceedings of the 28th International Conference on Compiler Construction* (Washington, DC,  
3195 USA) (CC 2019). Association for Computing Machinery, New York, NY, USA, 135–146. doi:[10.1145/3302516.3307359](https://doi.org/10.1145/3302516.3307359)
- 3196 [80] Manas Thakur and V. Krishna Nandivada. 2020. Mix Your Contexts Well: Opportunities Unleashed by Recent Advances  
3197 in Scaling Context-Sensitivity. In *Proceedings of the 29th International Conference on Compiler Construction* (San Diego,  
3198 CA, USA) (CC 2020). Association for Computing Machinery, New York, NY, USA, 27–38. doi:[10.1145/3377555.3377902](https://doi.org/10.1145/3377555.3377902)
- 3199 [81] Paolo Tonella and Alessandra Potrich. 2005. *Reverse Engineering of Object Oriented Code*. Springer. doi:[10.1007/b102522](https://doi.org/10.1007/b102522)
- 3200 [82] USI-DAG. 2022. BSS: Benchmark Suite for the Stream API. <https://github.com/usi-dag/BSS>.
- 3201 [83] WALA. 2006. Watson Libraries for Analysis. <http://wala.sf.net>.
- 3202 [84] Shiyi Wei and Barbara G. Ryder. 2015. Adaptive Context-sensitive Analysis for JavaScript. In *29th European Conference  
3203 on Object-Oriented Programming (ECOOP 2015) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 37)*,  
3204 John Tang Boyland (Ed.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 712–734. doi:[10.4230/LIPIcs.ECOOP.2015.712](https://doi.org/10.4230/LIPIcs.ECOOP.2015.712)
- 3205 [85] Robert P. Wilson and Monica S. Lam. 1995. Efficient context-sensitive pointer analysis for C programs. In *Proceedings  
3206 of the ACM SIGPLAN 1995 Conference on Programming Language Design and Implementation* (La Jolla, California, USA)  
3207 (PLDI '95). Association for Computing Machinery, New York, NY, USA, 1–12. doi:[10.1145/207110.207111](https://doi.org/10.1145/207110.207111)
- 3208 [86] YangShengYuan. 2025. Cut-Shortcut-Artifact. <https://github.com/YangShengYuan/cut-shortcut-artifact>.
- 3209 [87] Chenxi Zhang, Yufei Liang, Tian Tan, Chang Xu, Shuangxiang Kan, Yulei Sui, and Yue Li. 2025. Interactive Cross-  
3210 Language Pointer Analysis For Resolving Native Code in Java Programs . In *2025 IEEE/ACM 47th International  
3211 Conference on Software Engineering (ICSE)*. IEEE Computer Society, Los Alamitos, CA, USA, 612–612. doi:[10.1109/ICSE55347.2025.00075](https://doi.org/10.1109/ICSE55347.2025.00075)
- 3212 [88] Qirun Zhang, Michael R. Lyu, Hao Yuan, and Zhendong Su. 2013. Fast Algorithms for Dyck-CFL-Reachability with  
3213 Applications to Alias Analysis. In *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design  
3214 and Implementation* (Seattle, Washington, USA) (PLDI '13). Association for Computing Machinery, New York, NY, USA,  
3215 435–446. doi:[10.1145/2491956.2462159](https://doi.org/10.1145/2491956.2462159)
- 3216 [89] Xin Zhang, Ravi Mangal, Mayur Naik, and Hongseok Yang. 2014. Hybrid top-down and bottom-up interprocedural  
3217 analysis. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*  
3218 (Edinburgh, United Kingdom) (PLDI '14). Association for Computing Machinery, New York, NY, USA, 249–258.  
3219 doi:[10.1145/2594291.2594328](https://doi.org/10.1145/2594291.2594328)
- 3220  
3221  
3222  
3223  
3224  
3225  
3226  
3227  
3228  
3229  
3230  
3231  
3232  
3233  
3234

### 3235 A Conservative Fallback Mechanisms for Container-Access Pattern

3236 In this appendix, we introduce additional mechanisms for conservative fallbacks for the container-  
 3237 access pattern, including sound handling of custom containers (i.e., containers not defined in the JDK  
 3238 and therefore lacking specified APIs) and the treatment of hosts for this variables. We first extend  
 3239 the domain of hosts to  $\mathbb{H} = (\mathbb{L} \cup \hat{\mathbb{L}}) \times \mathbb{C}$ , where  $\hat{\mathbb{L}}$  is a special mock label, introduced to handle the  
 3240 this variable in the pointer-host mapping ( $pt_H$ ). For clarity of formalism, we define container types  
 3241 as  $\tau_C = \text{Collection} \mid \text{Map}$ , representing the standard JDK container interfaces. Host-relevant types  
 3242 are then defined as  $\tau_{\mathcal{H}} = \tau_C \mid \text{Iterator} \mid \text{ListIterator} \mid \text{Enumeration} \mid \text{Map.Entry}$ , thereby extending  
 3243  $\tau_C$  with container-dependent types.

$$\begin{array}{c} \frac{\mathcal{T}(o_i) = \tau \quad \tau \notin \text{JDK}}{h_c^i \in \text{conservH}} \text{ [CUSTOMCONT]} \quad \frac{h' \triangleleft h \quad h' \in \text{conservH}}{h \in \text{conservH}} \text{ [HOSTDEP-II]} \\ \\ \frac{i \rightarrow_{\text{call}} m \quad m \in \tau \quad \exists \tau_{\mathcal{H}}. \tau <: \tau_{\mathcal{H}} \quad c \in \mathbb{K}}{h_c^i \in pt_H(\text{this}^m) \quad h_c^{\hat{l}} \in \text{conservH}} \text{ [MockHost]} \quad \frac{i \rightarrow_{\text{call}} m \quad \text{ret}^m \in \text{cutRet} \quad h \in pt_H(\text{arg}_0^i) \quad h \in \text{conservH}}{\text{ret}^m \longrightarrow \text{LHS}^i} \text{ [CONSERVHOST]} \end{array}$$

3244 Fig. 36. Rules for handling container access pattern.

3245 The four rules above formalize mechanisms employed to ensure soundness. The central idea  
 3246 is to define a set  $\text{conservH}$ , which records certain hosts—those whose underlying container  
 3247 instances require conservative treatment during element retrieval. Rule [CONSERVHOST] governs this  
 3248 conservative behavior: while our analysis normally suppresses imprecise edges from EXIT methods’  
 3249 return variables to the TARGETS of a host, this suppression is deliberately disabled for hosts in  
 3250  $\text{conservH}$ . Specifically, the method return edges that would otherwise be suppressed (as per the  
 3251 boxed condition in rule [RETURN], Figure 7) are added back to the PFG (such inference takes place  
 3252 during the main pointer analysis, because it is part of the on-the-fly construction of the PFG),  
 3253 preserving soundness for those container usages. The next three rules, [CUSTOMCONT], [HOSTDEP-II],  
 3254 and [MockHost], describe how hosts are added to  $\text{conservH}$ :

- [CUSTOMCONT] adds any host created from a custom container class (i.e., a class of  $\tau$  such that  $\tau$  is not  
 3255 defined in the JDK) to  $\text{conservH}$ . Because such custom classes lack pre-specified ENTRANCE/EXIT  
 3256 specifications, the analysis cannot reason about them and must fall back on conservative handling.
- [HOSTDEP-II] propagates conservativeness via the host dependency preorder  $\triangleleft$ . If host  $h'$  is in  
 3257  $\text{conservH}$  and  $h' \triangleleft h$ , then  $h$  is also added to  $\text{conservH}$ . For example, if a JDK container receives  
 3258 elements via `addAll()` from a custom container, then the JDK container must be conservatively  
 3259 treated, because its SOURCES may now include elements that the analysis cannot track.
- [MockHost] addresses a subtle case: the analysis does not propagate hosts from a receiver variable  
 3260 to the `this` variable during method calls. This omission is deliberate and replaced by conservative  
 3261 handling. For any call site  $i$  that invokes a method  $m$  declared in a container or a container-  
 3262 dependent class  $\tau$  (i.e.,  $\exists \tau_{\mathcal{H}}$  such that  $\tau <: \tau_{\mathcal{H}}^{20}$ ), we construct a set of mock hosts  $\{h_c^{\hat{l}} \mid c \in \mathbb{K}\}$ ,  
 3263 assign them to  $pt_H(\text{this}^m)$ , and also record them in  $\text{conservH}$ . These mock hosts—defined using  
 3264 the special mock-instruction label  $\hat{l}$ —are shared across all this variables of relevant methods.  
 3265 This design avoids the cost of propagating hosts from receivers to `this` at each call site, while  
 3266 maintaining soundness. It also has little impact on precision, because most EXIT invocations on  
 3267 `this` do not cause merged flows to leak beyond the container.

3268  
 3269  
 3270  
 3271  
 3272  
 3273  
 3274  
 3275  
 3276  
 3277  
 3278  
 3279  
 3280  
 3281  
 3282  
 3283  
 3284  
 3285  
 3286  
 3287  
 3288  
 3289  
 3290  
 3291  
 3292  
 3293  
 3294  
 3295  
 3296  
 3297  
 3298  
 3299  
 3300  
 3301  
 3302  
 3303  
 3304  
 3305  
 3306  
 3307  
 3308  
 3309  
 3310  
 3311  
 3312  
 3313  
 3314  
 3315  
 3316  
 3317  
 3318  
 3319  
 3320  
 3321  
 3322  
 3323  
 3324  
 3325  
 3326  
 3327  
 3328  
 3329  
 3330  
 3331  
 3332  
 3333  
 3334  
 3335  
 3336  
 3337  
 3338  
 3339  
 3340  
 3341  
 3342  
 3343  
 3344  
 3345  
 3346  
 3347  
 3348  
 3349  
 3350  
 3351  
 3352  
 3353  
 3354  
 3355  
 3356  
 3357  
 3358  
 3359  
 3360  
 3361  
 3362  
 3363  
 3364  
 3365  
 3366  
 3367  
 3368  
 3369  
 3370  
 3371  
 3372  
 3373  
 3374  
 3375  
 3376  
 3377  
 3378  
 3379  
 3380  
 3381  
 3382  
 3383  
 3384  
 3385  
 3386  
 3387  
 3388  
 3389  
 3390  
 3391  
 3392  
 3393  
 3394  
 3395  
 3396  
 3397  
 3398  
 3399  
 3400  
 3401  
 3402  
 3403  
 3404  
 3405  
 3406  
 3407  
 3408  
 3409  
 3410  
 3411  
 3412  
 3413  
 3414  
 3415  
 3416  
 3417  
 3418  
 3419  
 3420  
 3421  
 3422  
 3423  
 3424  
 3425  
 3426  
 3427  
 3428  
 3429  
 3430  
 3431  
 3432  
 3433  
 3434  
 3435  
 3436  
 3437  
 3438  
 3439  
 3440  
 3441  
 3442  
 3443  
 3444  
 3445  
 3446  
 3447  
 3448  
 3449  
 3450  
 3451  
 3452  
 3453  
 3454  
 3455  
 3456  
 3457  
 3458  
 3459  
 3460  
 3461  
 3462  
 3463  
 3464  
 3465  
 3466  
 3467  
 3468  
 3469  
 3470  
 3471  
 3472  
 3473  
 3474  
 3475  
 3476  
 3477  
 3478  
 3479  
 3480  
 3481  
 3482  
 3483  
 3484  
 3485  
 3486  
 3487  
 3488  
 3489  
 3490  
 3491  
 3492  
 3493  
 3494  
 3495  
 3496  
 3497  
 3498  
 3499  
 3500  
 3501  
 3502  
 3503  
 3504  
 3505  
 3506  
 3507  
 3508  
 3509  
 3510  
 3511  
 3512  
 3513  
 3514  
 3515  
 3516  
 3517  
 3518  
 3519  
 3520  
 3521  
 3522  
 3523  
 3524  
 3525  
 3526  
 3527  
 3528  
 3529  
 3530  
 3531  
 3532  
 3533  
 3534  
 3535  
 3536  
 3537  
 3538  
 3539  
 3540  
 3541  
 3542  
 3543  
 3544  
 3545  
 3546  
 3547  
 3548  
 3549  
 3550  
 3551  
 3552  
 3553  
 3554  
 3555  
 3556  
 3557  
 3558  
 3559  
 3560  
 3561  
 3562  
 3563  
 3564  
 3565  
 3566  
 3567  
 3568  
 3569  
 3570  
 3571  
 3572  
 3573  
 3574  
 3575  
 3576  
 3577  
 3578  
 3579  
 3580  
 3581  
 3582  
 3583  
 3584  
 3585  
 3586  
 3587  
 3588  
 3589  
 3590  
 3591  
 3592  
 3593  
 3594  
 3595  
 3596  
 3597  
 3598  
 3599  
 3600  
 3601  
 3602  
 3603  
 3604  
 3605  
 3606  
 3607  
 3608  
 3609  
 3610  
 3611  
 3612  
 3613  
 3614  
 3615  
 3616  
 3617  
 3618  
 3619  
 3620  
 3621  
 3622  
 3623  
 3624  
 3625  
 3626  
 3627  
 3628  
 3629  
 3630  
 3631  
 3632  
 3633  
 3634  
 3635  
 3636  
 3637  
 3638  
 3639  
 3640  
 3641  
 3642  
 3643  
 3644  
 3645  
 3646  
 3647  
 3648  
 3649  
 3650  
 3651  
 3652  
 3653  
 3654  
 3655  
 3656  
 3657  
 3658  
 3659  
 3660  
 3661  
 3662  
 3663  
 3664  
 3665  
 3666  
 3667  
 3668  
 3669  
 3670  
 3671  
 3672  
 3673  
 3674  
 3675  
 3676  
 3677  
 3678  
 3679  
 3680  
 3681  
 3682  
 3683  
 3684  
 3685  
 3686  
 3687  
 3688  
 3689  
 3690  
 3691  
 3692  
 3693  
 3694  
 3695  
 3696  
 3697  
 3698  
 3699  
 3700  
 3701  
 3702  
 3703  
 3704  
 3705  
 3706  
 3707  
 3708  
 3709  
 3710  
 3711  
 3712  
 3713  
 3714  
 3715  
 3716  
 3717  
 3718  
 3719  
 3720  
 3721  
 3722  
 3723  
 3724  
 3725  
 3726  
 3727  
 3728  
 3729  
 3730  
 3731  
 3732  
 3733  
 3734  
 3735  
 3736  
 3737  
 3738  
 3739  
 3740  
 3741  
 3742  
 3743  
 3744  
 3745  
 3746  
 3747  
 3748  
 3749  
 3750  
 3751  
 3752  
 3753  
 3754  
 3755  
 3756  
 3757  
 3758  
 3759  
 3760  
 3761  
 3762  
 3763  
 3764  
 3765  
 3766  
 3767  
 3768  
 3769  
 3770  
 3771  
 3772  
 3773  
 3774  
 3775  
 3776  
 3777  
 3778  
 3779  
 3780  
 3781  
 3782  
 3783  
 3784  
 3785  
 3786  
 3787  
 3788  
 3789  
 3790  
 3791  
 3792  
 3793  
 3794  
 3795  
 3796  
 3797  
 3798  
 3799  
 3800  
 3801  
 3802  
 3803  
 3804  
 3805  
 3806  
 3807  
 3808  
 3809  
 3810  
 3811  
 3812  
 3813  
 3814  
 3815  
 3816  
 3817  
 3818  
 3819  
 3820  
 3821  
 3822  
 3823  
 3824  
 3825  
 3826  
 3827  
 3828  
 3829  
 3830  
 3831  
 3832  
 3833  
 3834  
 3835  
 3836  
 3837  
 3838  
 3839  
 3840  
 3841  
 3842  
 3843  
 3844  
 3845  
 3846  
 3847  
 3848  
 3849  
 3850  
 3851  
 3852  
 3853  
 3854  
 3855  
 3856  
 3857  
 3858  
 3859  
 3860  
 3861  
 3862  
 3863  
 3864  
 3865  
 3866  
 3867  
 3868  
 3869  
 3870  
 3871  
 3872  
 3873  
 3874  
 3875  
 3876  
 3877  
 3878  
 3879  
 3880  
 3881  
 3882  
 3883  
 3884  
 3885  
 3886  
 3887  
 3888  
 3889  
 3890  
 3891  
 3892  
 3893  
 3894  
 3895  
 3896  
 3897  
 3898  
 3899  
 3900  
 3901  
 3902  
 3903  
 3904  
 3905  
 3906  
 3907  
 3908  
 3909  
 3910  
 3911  
 3912  
 3913  
 3914  
 3915  
 3916  
 3917  
 3918  
 3919  
 3920  
 3921  
 3922  
 3923  
 3924  
 3925  
 3926  
 3927  
 3928  
 3929  
 3930  
 3931  
 3932  
 3933  
 3934  
 3935  
 3936  
 3937  
 3938  
 3939  
 3940  
 3941  
 3942  
 3943  
 3944  
 3945  
 3946  
 3947  
 3948  
 3949  
 3950  
 3951  
 3952  
 3953  
 3954  
 3955  
 3956  
 3957  
 3958  
 3959  
 3960  
 3961  
 3962  
 3963  
 3964  
 3965  
 3966  
 3967  
 3968  
 3969  
 3970  
 3971  
 3972  
 3973  
 3974  
 3975  
 3976  
 3977  
 3978  
 3979  
 3980  
 3981  
 3982  
 3983  
 3984  
 3985  
 3986  
 3987  
 3988  
 3989  
 3990  
 3991  
 3992  
 3993  
 3994  
 3995  
 3996  
 3997  
 3998  
 3999  
 4000  
 4001  
 4002  
 4003  
 4004  
 4005  
 4006  
 4007  
 4008  
 4009  
 4010  
 4011  
 4012  
 4013  
 4014  
 4015  
 4016  
 4017  
 4018  
 4019  
 4020  
 4021  
 4022  
 4023  
 4024  
 4025  
 4026  
 4027  
 4028  
 4029  
 4030  
 4031  
 4032  
 4033  
 4034  
 4035  
 4036  
 4037  
 4038  
 4039  
 4040  
 4041  
 4042  
 4043  
 4044  
 4045  
 4046  
 4047  
 4048  
 4049  
 4050  
 4051  
 4052  
 4053  
 4054  
 4055  
 4056  
 4057  
 4058  
 4059  
 4060  
 4061  
 4062  
 4063  
 4064  
 4065  
 4066  
 4067  
 4068  
 4069  
 4070  
 4071  
 4072  
 4073  
 4074  
 4075  
 4076  
 4077  
 4078  
 4079  
 4080  
 4081  
 4082  
 4083  
 4084  
 4085  
 4086  
 4087  
 4088  
 4089  
 4090  
 4091  
 4092  
 4093  
 4094  
 4095  
 4096  
 4097  
 4098  
 4099  
 4100  
 4101  
 4102  
 4103  
 4104  
 4105  
 4106  
 4107  
 4108  
 4109  
 4110  
 4111  
 4112  
 4113  
 4114  
 4115  
 4116  
 4117  
 4118  
 4119  
 4120  
 4121  
 4122  
 4123  
 4124  
 4125  
 4126  
 4127  
 4128  
 4129  
 4130  
 4131  
 4132  
 4133  
 4134  
 4135  
 4136  
 4137  
 4138  
 4139  
 4140  
 4141  
 4142  
 4143  
 4144  
 4145  
 4146  
 4147  
 4148  
 4149  
 4150  
 4151  
 4152  
 4153  
 4154  
 4155  
 4156  
 4157  
 4158  
 4159  
 4160  
 4161  
 4162  
 4163  
 4164  
 4165  
 4166  
 4167  
 4168  
 4169  
 4170  
 4171  
 4172  
 4173  
 4174  
 4175  
 4176  
 4177  
 4178  
 4179  
 4180  
 4181  
 4182  
 4183  
 4184  
 4185  
 4186  
 4187  
 4188  
 4189  
 4190  
 4191  
 4192  
 4193  
 4194  
 4195  
 4196  
 4197  
 4198  
 4199  
 4200  
 4201  
 4202  
 4203  
 4204  
 4205  
 4206  
 4207  
 4208  
 4209  
 4210  
 4211  
 4212  
 4213  
 4214  
 4215  
 4216  
 4217  
 4218  
 4219  
 4220  
 4221  
 4222  
 4223  
 4224  
 4225  
 4226  
 4227  
 4228  
 4229  
 4230  
 4231  
 4232  
 4233  
 4234  
 4235  
 4236  
 4237  
 4238  
 4239  
 4240  
 4241  
 4242  
 4243  
 4244  
 4245  
 4246  
 4247  
 4248  
 4249  
 4250  
 4251  
 4252  
 4253  
 4254  
 4255  
 4256  
 4257  
 4258  
 4259  
 4260  
 4261  
 4262  
 4263  
 4264  
 4265  
 4266  
 4267  
 4268  
 4269  
 4270  
 4271  
 4272  
 4273  
 4274  
 4275  
 4276  
 4277  
 4278  
 4279  
 4280  
 4281  
 4282  
 4283  
 4284  
 4285  
 4286  
 4287  
 4288  
 4289  
 4290  
 4291  
 4292  
 4293  
 4294  
 4295  
 4296  
 4297  
 4298  
 4299  
 4300  
 4301  
 4302  
 4303  
 4304  
 4305  
 4306  
 4307  
 4308  
 4309  
 4310  
 4311  
 4312  
 4313  
 4314  
 4315  
 4316  
 4317  
 4318  
 4319  
 4320  
 4321  
 4322  
 4323  
 4324  
 4325  
 4326  
 4327  
 4328  
 4329  
 4330  
 4331  
 4332  
 4333  
 4334  
 4335  
 4336  
 4337  
 4338  
 4339  
 4340  
 4341  
 4342  
 4343  
 4344  
 4345  
 4346  
 4347  
 4348  
 4349  
 4350  
 4351  
 4352  
 4353  
 4354  
 4355  
 4356  
 4357  
 4358  
 4359  
 4360  
 4361  
 4362  
 4363  
 4364  
 4365  
 4366  
 4367  
 4368  
 4369  
 4370  
 4371  
 4372  
 4373  
 4374  
 4375  
 4376  
 4377  
 4378  
 4379  
 4380  
 4381  
 4382  
 4383  
 4384  
 4385  
 4386  
 4387  
 4388  
 4389  
 4390  
 4391  
 4392  
 4393  
 4394  
 4395  
 4396  
 4397  
 4398  
 4399  
 4400  
 4401  
 4402  
 4403  
 4404  
 4405  
 4406  
 4407  
 4408  
 4409  
 4410  
 4411  
 4412  
 4413  
 4414  
 4415  
 4416  
 4417  
 4418  
 4419  
 4420  
 4421  
 4422  
 4423  
 4424  
 4425  
 4426  
 4427  
 4428  
 4429  
 4430  
 4431  
 4432  
 4433  
 4434  
 4435  
 4436  
 4437  
 4438  
 4439  
 4440  
 4441  
 4442  
 4443  
 4444  
 4445  
 4446  
 4447  
 4448  
 4449  
 4450  
 4451  
 4452  
 4453  
 4454  
 4455  
 4456  
 4457  
 4458  
 4459  
 4460  
 4461  
 4462  
 4463  
 4464  
 4465  
 4466  
 4467  
 4468  
 4469  
 4470  
 4471  
 4472  
 4473  
 4474  
 4475  
 4476  
 4477  
 4478  
 4479  
 4480  
 4481  
 4482  
 4483  
 4484  
 4485  
 4486  
 4487  
 4488  
 4489  
 4490  
 4491  
 4492  
 4493  
 4494  
 4495  
 4496  
 4497  
 4498  
 4499  
 4500  
 4501  
 4502  
 4503  
 4504  
 4505  
 4506  
 4507  
 4508  
 4509  
 4510  
 4511  
 4512  
 4513  
 4514  
 4515  
 4516  
 4517  
 4518  
 4519  
 4520  
 4521  
 4522  
 4523  
 4524  
 4525  
 4526  
 4527  
 4528  
 4529  
 4530  
 4531  
 4532  
 4533  
 4534  
 4535  
 4536  
 4537  
 4538  
 4539  
 4540  
 4541  
 4542  
 4543  
 4544  
 4545  
 4546  
 4547  
 4548  
 4549  
 4550  
 4551  
 4552  
 4553  
 4554  
 4555  
 4556  
 4557  
 4558  
 4559  
 4560  
 4561  
 4562  
 4563  
 4564  
 4565  
 4566  
 4567  
 4568  
 4569  
 4570  
 4571  
 4572  
 4573  
 4574  
 4575  
 4576  
 4577  
 4578  
 4579  
 4580  
 4581  
 4582  
 4583  
 4584  
 4585

3284    **B Stream Operations**

3285    A complete list of stream operations modeled in our analysis is shown in Figure 37. (The signatures  
 3286    of methods inherited by subclassing? are omitted.) Operations marked with circled letters indicate  
 3287    cases where, beyond the sub-category-specific rules introduced in Sections 7.3–7.5, the following  
 3288    special-case handling is used:

- 3289    **A**: requires special handling to record array elements as host SOURCES;
- 3290    **B**: Stream.concat concatenates two stream pipelines into a new one. Handling this requires  
 3291    propagating SOURCES from multiple Strm hosts to another, which is easily achieved using the  
 3292    host-dependency preorder  $\triangleleft$  (defined in Section 4.3).
- 3293    **C**: requires augmenting the SOURCES of an existing Strm host, which can be handled straightfor-  
 3294    wardly by registering these operations in the CONTEENTR set (defined in Section 4.3).

3295    The detailed rules for such special-case handling are omitted for brevity, because they are not  
 3296    central to the core design of our stream modeling. The complete inference rules for all *functional*  
 3297    stream operations are provided in Appendix C.

3298    **STRUCTURE CREATION** — handled by [S-CREATE]

```
Stream.of: (Object) -> Stream
Stream.of: (Object[]) -> Stream ①
Stream.builder: () -> Stream$Builder
Stream.empty: () -> Stream
Stream.concat: (Stream,Stream) -> Stream ②
Arrays.stream: (Object[]) -> Stream ③
Arrays.stream: (Object[],int,int) -> Stream ④
Optional.of: (Object) -> Stream
Optional.empty: () -> Optional
Optional.of: (Object) -> Optional
Optional.ofNullable: (Object) -> Optional
```

3299    **STRUCTURE PROCESS** — handled by [S-PROCESS]

```
Stream.distinct: () -> Stream
Stream.filter: (Predicate) -> Stream
Stream.findAny: (Predicate) -> Optional
Stream.findFirst: (Predicate) -> Optional
Stream.limit: long -> Stream
Stream.min: (Comparator) -> Optional
Stream.max: (Comparator) -> Optional
Stream.sorted: () -> Stream
Stream.sorted: (Comparatpr) -> Stream
Stream.skip: long -> Stream
BaseStream.sequential: () -> BaseStream
BaseStream.parallel: () -> BaseStream
BaseStream.onClose: (Runnable) -> BaseStream
BaseStream.unordered: () -> BaseStream
Stream$Builder.build: () -> Stream
Stream$Builder.add: (Object) -> Stream$Builder ⑤
Stream$Builder.accept: (Object) -> Stream$Builder ⑥
Optional.filter: (Predicate) -> Optional
```

3300    **STRUCTURE OUTPUT** — handled by [S-OUTPUT]

```
Optional.get: () -> Object
Optional.orElse: (Object) -> Object ⑦
Optional.orElseThrow: (Supplier) -> Object
```

3301    handled by operation-specific rules (in Appendix B)

3302    **FUNCTIONAL CREATION**

```
Stream.generate: (Supplier) -> Stream
Stream.iterate: (Object,UnaryOperator) -> Stream
IntStream.boxed: () -> Stream
IntStream.mapToObj: (IntFunction) -> Stream
LongStream.boxed: () -> Stream
LongStream.mapToObj: (LongFunction) -> Stream
DoubleStream.boxed: () -> Stream
DoubleStream.mapToObj: (DoubleFunction) -> Stream
```

3303    **FUNCTIONAL PROCESS**

```
Stream.flatMap: (Function) -> Stream
Stream.map: (Function) -> Stream
Stream.reduce: (BinaryOperator) -> Optional
Optional.flatMap: (Function) -> Optional
Optional.map: (Function) -> Optional
Stream.peek: (Consumer) -> Stream
```

3304    **FUNCTIONAL OUTPUT**

```
Stream.forEach: (Consumer) -> void
Stream.forEachOrdered: (Consumer) -> void
Stream.collect: (Supplier,BiConsumer,BiConsumer) -> Object
Stream.reduce: (Object,BinaryOperator) -> Object
Stream.reduce: (Object,BiFunction,BinaryOperator) -> Object
Optional.orElseGet: (Supplier) -> Object
Optional.isPresent: (Consumer) -> void
```

3305    **COLLECTION CREATION** — handled by [C-CREATE]

```
Collection.stream: () -> Stream
Collection.parallelStream: () -> Stream
StreamSupport.stream: (Spliterator,int,boolean) -> Stream
```

3306    **COLLECTION OUTPUT** — handled by [C-OUTPUT]

```
Stream.collect: (Collector) -> Object
Stream.toList: () -> List
BaseStreamspliterator: () -> Spliterator
BaseStream.iterator: () -> Iterator
```

Fig. 37. The list of modeled stream operations.

### 3333 C Handling Functional Stream Operations

3334 In Section 7.4, we presented a selected subset of rules (in Figure 25) from the lambda analysis  
 3335 proposed by Fourtounis and Smaragdakis [18]. The complete rule set used to handle lambda call  
 3336 edges is shown in Figure 38. For clarity and consistency, we re-formalize these rules in our own  
 3337 notation.

3338

$$\begin{array}{c}
 \frac{i : r = b.m(a_1, \dots, a_n) \quad \lambda \in pt(b)}{\langle i, \lambda \rangle \in LCallSites} \text{ [L-CALLSITE]} \\
 \\ 
 \frac{\langle i, \lambda \rangle \in LCallSites \quad \#cap^\lambda = 0 \quad \text{Target}(\lambda) = m \quad m \notin Static}{\langle i, \lambda, Larg_1^{i,\lambda} \rangle \in LRecvs} \text{ [L-RECV-I]} \quad \frac{\langle i, \lambda \rangle \in LCallSites \quad \#cap^\lambda \neq 0 \quad \text{Target}(\lambda) = m \quad m \notin Static}{\langle i, \lambda, cap_1^\lambda \rangle \in LRecvs} \text{ [L-RECV-II]} \\
 \\ 
 \frac{\langle i, \lambda \rangle \in LCallSites \quad \text{Target}(\lambda) = m \quad m \notin Static \quad \langle i, \lambda, x \rangle \in LRecvs \quad o_j \in pt(x) \quad m' = \text{dispatch}(m, o_j)}{o_j \in pt(\text{param}_0^{m'}) \quad i \rightarrow_\lambda m'} \text{ [L-INS CALL]} \\
 \\ 
 \frac{i \rightarrow_\lambda m \quad m \notin Static \quad \#cap^\lambda > 1}{\forall k \geq 2 : cap_k^\lambda \rightarrow \text{param}_{k-1}^m} \text{ [L-INS CAP]} \\
 \\ 
 \frac{i \rightarrow_\lambda m \quad m \notin Static \quad \#cap^\lambda = 0}{\forall k \geq 2 : Larg_k^{i,\lambda} \rightarrow \text{param}_{k-1}^m} \text{ [L-INS LARG-I]} \quad \frac{i \rightarrow_\lambda m \quad m \notin Static \quad \#cap^\lambda = n > 0}{\forall k \geq 1 : Larg_k^{i,\lambda} \rightarrow \text{param}_{k+n-1}^m} \text{ [L-INS LARG-II]} \\
 \\ 
 \frac{\langle i, \lambda \rangle \in LCallSites \quad \text{Target}(\lambda) = m \quad m \in Static}{i \rightarrow_\lambda m} \text{ [L-STACALL]} \quad \frac{i \rightarrow_\lambda m \quad m \in Static \quad \#cap^\lambda \neq 0}{\forall k \geq 1 : cap_k^\lambda \rightarrow \text{param}_k^m} \text{ [L-STACAP]} \\
 \\ 
 \frac{i \rightarrow_\lambda m \quad m \in Static \quad \#cap^\lambda = n \geq 0}{\forall k \geq 1 : Larg_k^{i,\lambda} \rightarrow \text{param}_{n+k}^m} \text{ [L-STALARG]} \quad \frac{i \rightarrow_\lambda m}{\text{ret}^m \rightarrow \text{LHS}^i} \text{ [L-RET]} \\
 \\ 
 \frac{i \rightarrow_\lambda m \quad \boxed{i \notin \text{STM CALL SITE}}}{\forall k \geq 1 : \text{arg}_k^i \rightarrow Larg_k^{i,\lambda}} \text{ [ARG2LARG]}
 \end{array}$$

3344 Fig. 38. Rules for handling functional call-graph edges. The boxed premise in [ARG2LARG] is added when  
 3345 integrated with stream modeling.

3346

3347 Compared to its original formulation in [18], this version introduces slight modifications to  
 3348 support more precise handling of dynamic dispatch and receiver passing for *target methods* of  
 3349 lambda expressions and method references. that are instance methods (see [L-RECV-I], [L-RECV-II], and  
 3350 [L-INS CALL]). The primary complexity of this rule set lies in uniformly modeling receiver passing and  
 3351 environment-captured variable passing for both method references and lambdas. While method  
 3352 references have a constrained syntax that makes their target methods straightforward to identify,  
 3353 lambdas can capture local variables from the surrounding scope. To handle this issue, the compiler  
 3354 generates a synthetic method that encapsulates the lambda body, with any captured variables  
 3355 passed as parameters—this generated method becomes the target method of the functional object  
 3356 created by the lambda. Notably, the rules for handling a *target method* being an instance method  
 3357 require shifting the argument list to accommodate the receiver. This issue is reflected in the "-1" shift  
 3358 in [L-INS CAP], [L-INS LARG-I], and [L-INS LARG-II], where captured variables (if any) are followed by the  
 3359 lambda's actual arguments when passed to the parameters of the target method. In contrast, static  
 3360 methods do not require this shift ([L-STACAP] and [L-STAACT]). These distinctions ensure accurate  
 3361

3362

3382 modeling of function-invocation semantics in the presence of lambda expressions and method  
 3383 references.

3384

3385  $m \in \{\text{Stream.reduce(Object, BinaryOperator)}, \text{Stream.reduce(Object, BiFunction, BinaryOperator)},$   
 $\text{Stream.collect(Supplier, BiConsumer, BiConsumer)}, \text{Optional.orElseGet(Supplier)}\}$

[CUT-FSTM]

3387

3388  $i \rightarrow_{call} m \quad \lambda \in pt(\text{arg}_2^i) \quad \text{Target}(\lambda) = m'$   $i \rightarrow_{call} m \quad \lambda \in pt(\text{arg}_1^i) \quad \text{Target}(\lambda) = m'$   
 3389  $m = \text{Stream.iterate(Object, UnaryOperator)}$   $m \in \{\text{Stream.generate}, \text{IntStream.mapToObj},$   
 $\text{LongStream.mapToObj}, \text{DoubleStream.mapToObj}\}$  [GENERATE]  
 3390  $\text{arg}_1^i \rightarrow \text{Larg}_1^{*,\lambda} \quad ret^{m'} \rightarrow \text{Larg}_1^{*,\lambda}$  [ITERATE]  
 $h_{\text{Strm}}^i \in pt_H(\text{LHS}^i) \quad h_{\text{Strm}}^i \Leftarrow ret^{m'}$   
 3391  $h_{\text{Strm}}^i \Leftarrow \text{arg}_1^i \quad h_{\text{Strm}}^i \Leftarrow ret^{m'}$

3392

3393  $i \rightarrow_{call} m \quad m \in \{\text{Stream.map}, \text{Optional.map}\}$   $i \rightarrow_{call} m \quad m \in \{\text{Stream.map}, \text{Optional.map}\}$   
 3394  $\lambda \in pt(\text{arg}_1^i) \quad \text{Target}(\lambda) = m' \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i)$  [MAP]  
 $h_{\text{Strm}}^j \Rightarrow \text{Larg}_1^{*,\lambda} \quad h_{\text{Strm}}^i \in pt_H(\text{LHS}^i) \quad h_{\text{Strm}}^i \Leftarrow ret^{m'}$

3395

3396  $i \rightarrow_{call} m \quad m \in \{\text{Stream.flatMap}, \text{Optional.flatMap}\}$   $i \rightarrow_{call} m \quad m \in \{\text{Stream.flatMap}, \text{Optional.flatMap}\}$   
 3397  $\lambda \in pt(\text{arg}_1^i) \quad \text{Target}(\lambda) = m' \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_1^i) \quad h_{\text{Strm}}^k \in pt_H(ret^{m'})$  [FLATMAP]  
 $h_{\text{Strm}}^j \Rightarrow \text{Larg}_1^{*,\lambda} \quad h_{\text{Strm}}^i \in pt_H(\text{LHS}^i) \quad h_{\text{Strm}}^k \triangleleft h_{\text{Strm}}^i$

3398

3399  $i \rightarrow_{call} m \quad m = \text{Stream.reduce(BinaryOperator)}$   $i \rightarrow_{call} m \quad m = \text{Stream.reduce(BinaryOperator)}$   
 3400  $\lambda \in pt(\text{arg}_1^i) \quad \text{Target}(\lambda) m' \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i)$  [REDUCE-I]  
 $h_{\text{Strm}}^j \Rightarrow \text{Larg}_1^{*,\lambda} \quad h_{\text{Strm}}^j \Rightarrow \text{Larg}_2^{*,\lambda} \quad h_{\text{Strm}}^i \in pt_H(\text{LHS}^i)$   
 $h_{\text{Strm}}^j \triangleleft h_{\text{Strm}}^i \quad h_{\text{Strm}}^i \Leftarrow ret^{m'} \quad ret^{m'} \rightarrow \text{Larg}_1^{*,\lambda}$

3401

3402  $i \rightarrow_{call} m \quad m = \text{Stream.reduce(Object, BinaryOperator)}$   $i \rightarrow_{call} m \quad m = \text{Stream.reduce(Object, BinaryOperator)}$   
 3403  $\lambda \in pt(\text{arg}_2^i) \quad \text{Target}(\lambda) = m' \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i)$  [REDUCE-II]  
 $h_{\text{Strm}}^j \Rightarrow \text{Larg}_2^{*,\lambda} \quad arg_1^i \rightarrow LHS^i$   
 $arg_1^i \rightarrow \text{Larg}_1^{*,\lambda} \quad ret^{m'} \rightarrow \text{Larg}_1^{*,\lambda} \quad ret^{m'} \rightarrow LHS^i$

3404

3405  $i \rightarrow_{call} m \quad \lambda_1 \in pt(\text{arg}_2^i) \quad \text{Target}(\lambda_1) = m'$   $i \rightarrow_{call} m \quad \lambda_1 \in pt(\text{arg}_2^i) \quad \text{Target}(\lambda_1) = m'$   
 3406  $\lambda_2 \in pt(\text{arg}_3^i) \quad \text{Target}(\lambda_2) = m'' \quad h \in pt_H(\text{arg}_0^i)$  [REDUCE-III]  
 $m = \text{Stream.reduce(Object, BiFunction, BinaryOperator)}$

3407

3408  $arg_1^i \rightarrow \text{Larg}_1^{*,\lambda_1} \quad ret^{m'} \rightarrow \text{Larg}_1^{*,\lambda_1} \quad ret^{m''} \rightarrow \text{Larg}_1^{*,\lambda_1}$   
 $arg_1^i \rightarrow \text{Larg}_1^{*,\lambda_2} \quad ret^{m'} \rightarrow \text{Larg}_1^{*,\lambda_2} \quad ret^{m''} \rightarrow \text{Larg}_1^{*,\lambda_2}$   
 $arg_1^i \rightarrow \text{Larg}_2^{*,\lambda_2} \quad ret^{m'} \rightarrow \text{Larg}_2^{*,\lambda_2} \quad ret^{m''} \rightarrow \text{Larg}_2^{*,\lambda_2}$

3409

3410  $i \rightarrow_{call} m \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i) \quad \lambda_1 \in pt(\text{arg}_1^i) \quad \text{Target}(\lambda_1) = m'$   $i \rightarrow_{call} m \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i) \quad \lambda_1 \in pt(\text{arg}_1^i) \quad \text{Target}(\lambda_1) = m'$   
 3411  $\lambda_2 \in pt(\text{arg}_2^i) \quad \text{Target}(\lambda_2) = m'' \quad \lambda_3 \in pt(\text{arg}_3^i) \quad \text{Target}(\lambda_3) = m'''$   
 $m = \text{Stream.collect(Supplier, BiConsumer, BiConsumer)}$  [COLLECT]  
 $h_{\text{Strm}}^j \Rightarrow \text{Larg}_2^{*,\lambda_2} \quad ret^{m'} \rightarrow \text{Larg}_1^{*,\lambda_1} \quad ret^{m'} \rightarrow \text{Larg}_1^{*,\lambda_2}$   
 $ret^{m'} \rightarrow \text{Larg}_1^{*,\lambda_3} \quad ret^{m'} \rightarrow \text{Larg}_2^{*,\lambda_3}$

3412

3413  $i \rightarrow_{call} m \quad m = \text{Optional.orElseGet(Supplier)}$   $i \rightarrow_{call} m \quad \lambda \in pt(\text{arg}_1^i) \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i)$   
 3414  $\lambda \in pt(\text{arg}_1^i) \quad \text{Target}(\lambda) = m' \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i)$  [ORELSE]  
 $ret^{m'} \rightarrow \text{LHS}^i \quad h_{\text{Strm}}^j \Rightarrow \text{LHS}^i$  [FOREACH]  
 $ret^{m'} \rightarrow \text{LHS}^i \quad h_{\text{Strm}}^j \Rightarrow \text{LHS}^i$

3415

3416  $i \rightarrow_{call} m \quad \lambda \in pt(\text{arg}_1^i) \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i)$  [FOREACH]  
 $m \in \{\text{Stream.forEach}, \text{Stream.forEachOrdered},$   
 $\text{Stream.peek}, \text{Optional.ifPresent}\}$

3417

3418  $h_{\text{Strm}}^j \Rightarrow \text{Larg}_1^{*,\lambda} \quad h_{\text{Strm}}^j \Rightarrow \text{Larg}_2^{*,\lambda}$

3419

3420  $i \rightarrow_{call} m \quad m = \text{Optional.orElseGet(Supplier)}$   $i \rightarrow_{call} m \quad \lambda \in pt(\text{arg}_1^i) \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i)$   
 $\lambda \in pt(\text{arg}_1^i) \quad \text{Target}(\lambda) = m' \quad h_{\text{Strm}}^j \in pt_H(\text{arg}_0^i)$  [FOREACH]

3421

3422  $ret^{m'} \rightarrow \text{LHS}^i \quad h_{\text{Strm}}^j \Rightarrow \text{LHS}^i$

3423

Fig. 39. Rules for *functional stream operations*, where "\*" indicates a wildcard.

3424

3425 The complete rules for each *functional stream operation* are provided in Figure 39. The rule [CUT-FSTM] plays a role similar to [CUTCONT] in Figure 11, both should be performed prior to the pointer analysis to determine the set *cutRet*. The remaining rules in Figure 39 handle each *functional stream operation*. Some operations can be covered by a single rule because, in terms of our analysis, their semantics are equivalent. Notably, IntStream.boxed, LongStream.boxed, and DoubleStream.boxed are

3426

3431 implicitly handled by the rule [GENERATE], because these methods internally invoke the corresponding  
 3432 `mapToObj` method. The most intricate cases are the rules for the three variants of `Stream.reduce`,  
 3433 which differ in their argument counts and semantics. These require careful modeling to capture  
 3434 how the arguments interact with each other, and with the host representing the stream pipeline  
 3435 instance. For example, the rule [REDUCE-III] models the most complex variant, `Stream.reduce(Object, BiFunction, BinaryOperator)`, which takes three arguments:

- *identity*: This argument is the initial value of the reduction. We model its flow by connecting it to actual arguments of the two lambda functions that take in partial-reduction results ( $\lambda_1$  and  $\lambda_2$ ), i.e.,  $\text{arg}_1^i \rightarrow \text{Larg}_1^{*,\lambda_1}$ ,  $\text{arg}_1^i \rightarrow \text{Larg}_1^{*,\lambda_2}$ , and  $\text{arg}_1^i \rightarrow \text{Larg}_2^{*,\lambda_2}$ . It also serves as the default result if the stream is empty, so  $\text{arg}_1^i \rightarrow \text{LHS}^i$ .
- *accumulator*: This function takes two arguments—a partial-reduction result and the next stream element. The stream element is passed via  $h \Rightarrow \text{Larg}_2^{*,\lambda_1}$ . The function returns a new partial result, which is connected back to wherever a partial-reduction result is needed: to *accumulator* in subsequent reduction steps ( $\text{ret}^{m'} \rightarrow \text{Larg}_1^{*,\lambda_1}$ ), to the *combiner*'s arguments ( $\text{ret}^{m'} \rightarrow \text{Larg}_1^{*,\lambda_2}$  and  $\text{ret}^{m'} \rightarrow \text{Larg}_2^{*,\lambda_2}$ ), and to the overall result ( $\text{ret}^{m'} \rightarrow \text{LHS}^i$ ).
- *combiner*: This function merges two partial results when the reduction is parallelized. The return value of it is also a partial result, and flows into wherever a partial result is needed: it is connected back to the final result ( $\text{ret}^{m''} \rightarrow \text{LHS}^i$ ), to the *accumulator* ( $\text{ret}^{m''} \rightarrow \text{Larg}_1^{*,\lambda_1}$ ), and recursively to the *combiner* itself ( $\text{ret}^{m''} \rightarrow \text{Larg}_1^{*,\lambda_2}$  and  $\text{ret}^{m''} \rightarrow \text{Larg}_2^{*,\lambda_2}$ ).

3450 Among the remaining rules, [FLATMAP] is the most interesting. Here, the return value of the lambda's  
 3451 target method is itself a stream, and we need to identify all hosts of kind `Strm` that are associated  
 3452 with that return value ( $h_{\text{Strm}}^k \in pt_H(\text{ret}^{m'})$ ). We then add a preorder  $\lhd$  between each of these hosts  
 3453 and the respective result host associated with the LHS variable ( $h_{\text{Strm}}^i \in pt_H(\text{LHS}^i)$ ), indicating that  
 3454 we propagate the SOURCES of the former to the newly generated host representing the result stream,  
 3455 so that the “flattened” semantics is accurately modeled.

## 3457 D Ensuring Soundness for Handling Streams

3458 In this appendix, we introduce additional mechanisms and conservative fallbacks to ensure sound-  
 3459 ness for our handling of Java streams in CUT-SHORTCUT<sup>S</sup>.

3460 *Unmodeled Operations.* Although our implementation fully models all relevant stream operations  
 3461 up to Java 8, newer APIs (e.g., `Stream.mapMulti` introduced in Java 16) may not be covered. To  
 3462 maintain soundness, we conservatively handle such cases using rule [UNMODELED] in Figure 40. If a  
 3463 method is not included in the set of modeled operations (MODELEDSTMOPS), but its return type is  
 3464 a subtype of `Stream` or `Optional`, we generate a placeholder host  $h_{\text{Strm}}^i$  and record it in *conservH*,  
 3465 indicating that it must be treated conservatively. Rule [CONSERVHOST] in Figure 36 then ensures those  
 3466 otherwise-suppressed return edges for the relevant stream pipelines are preserved, and soundness  
 3467 is maintained.

$$\frac{i \rightarrow_{\text{call}} m \quad m \notin \text{MODELEDSTMOPS} \quad \mathcal{T}(\text{ret}^{m'}) \lhd: \text{Stream} \vee \text{Optional}}{h_{\text{Strm}}^i \in pt_H(\text{LHS}^i) \quad h_{\text{Strm}}^i \in \text{conservH}} \text{ [UNMODELED]}$$

3473 Fig. 40. The rule for handling unmodeled stream operations.  
 3474

3475 *Functional Operations.* To ensure that the integration of our handling of streams with lambda  
 3476 analysis [18] is sound, we introduce the rule [CONSERVFUNC] in Figure 41 to disable PFG edge  
 3477 suppression for *functional* stream operations when the associated host is conservatively handled.  
 3478 Specifically, [CONSERVFUNC] add edges for the actual arguments of lambdas involved in *functional*  
 3479

stream operations, only when the invocation receiver is associated with a host recorded in  $\text{conservH}$ . These edges connect call-site arguments ( $\text{arg}_k^{i'}$ ) to the mock actual arguments ( $\text{Larg}_k^{i',\lambda}$ ), and would otherwise be suppressed due to the boxed premise in [ARG2LARG] in Figure 25.

$$\frac{i \rightarrow_{call} m \quad \lambda \in pt(\text{arg}_k^i) \quad i' \rightarrow_{\lambda} m' \quad h \in pt_H(\text{arg}_0^i) \quad h \in \text{conservH}}{\forall k \geq 1 : \text{arg}_k^{i'} \longrightarrow \text{Larg}_k^{i',\lambda}} \text{[CONSERVFUNC]}$$

Fig. 41. The rule to disable PFG edge suppression for *functional* stream operations.

*Transformation through Complex Collectors.* Finally, we conservatively handle cases where Stream.collect(Collector) is invoked with a *complex* collector. Specifically, if the collector argument is not constructed via Collector.toSet, Collector.toList, or Collector.toCollection (determined by checking whether its allocation site belongs to SIMPLECOLLECTORALLOC, a set of instruction labels gathered by a naive pre-analysis over the Collector class), we conservatively generate one host per  $c \in \mathbb{K}$  and associate them with the LHS variable. This association signals that suppression of the corresponding PFG edges must be disabled.

$$\frac{i \rightarrow_{call} m \quad m = \text{Stream.collect}(\text{Collector}) \quad o_j \in \text{arg}_1^i \quad j \notin \text{SIMPLECOLLECTORALLOC} \quad c \in \mathbb{K}}{h_c^i \in \text{conservH}} \text{[CONSERVCOLLECT]}$$

Fig. 42. Rules for conservatively handling Stream.collect through complex collectors.

## E The Recall Experiment Results for RQ1

We summarized the recall results in Section 8.1; here, Table 8 presents the detailed experimental data. All results can be reproduced using our accompanying artifact.

As mentioned in Section 8.1, we conducted a recall experiment to validate the soundness of CUT-SHORTCUT. For each benchmark, we recorded the reachable methods and call-graph edges observed during dynamic execution, and then measured how many of these were recalled by CUT-SHORTCUT and by other analyses. As shown in Table 8, CUT-SHORTCUT’s recall rate is comparable to, and in most cases no lower than, that of 2obj, 2type, ZIPPER<sup>e</sup>-2obj, and ZIPPER<sup>e</sup>-2type. The only exception is the call-edge metric for jython, where CUT-SHORTCUT’s recall rate is slightly lower than all baselines. Note that a static analysis cannot achieve 100% recall because, in general, dynamic features—like reflection—cannot be soundly analyzed by a static analysis.

We further examined the reachable methods and call-graph edges that CUT-SHORTCUT missed (relative to the dynamic results) but that were not missed by any of the compared analyses. Items missed by both CUT-SHORTCUT and at least one baseline are excluded from consideration, because such misses cannot be attributed to CUT-SHORTCUT. Across all benchmarks, CUT-SHORTCUT missed only 1 reachable method and 8 call-graph edges. Manual inspection revealed that all of these were incorrectly recorded by the instrumentation tool rather than being true misses. Therefore, we conclude that the set of dynamically recorded reachable methods and call-graph edges recalled by existing analyses can also be recalled by CUT-SHORTCUT. This result indicates that our implementation of CUT-SHORTCUT does not introduce additional unsoundness compared to the analyses we compares.

Table 8. Recall results against dynamic ground truth on 10 benchmarks.

Program	Dynamic		Analysis	#reach-mtd		#call-edge	
	#reach-mtd	#call-edge		Recall	Rate	Recall	Rate
eclipse	8093	20916	CI	7083	87.52%	18068	86.38%
			2obj	7078	87.46%	18053	86.31%
			2type	7078	87.46%	18054	86.32%
			ZIPPER <sup>e</sup> -2obj	7078	87.46%	18056	86.33%
			ZIPPER <sup>e</sup> -2type	7078	87.46%	18059	86.34%
freecol	19387	57455	CUT-SHORTCUT	7080	87.48%	18057	86.33%
			CI	19126	98.65%	56612	98.53%
			2obj	—	—	—	—
			2type	—	—	—	—
			ZIPPER <sup>e</sup> -2obj	19124	98.64%	56605	98.52%
briss	14244	39575	ZIPPER <sup>e</sup> -2type	19126	98.65%	56606	98.52%
			CUT-SHORTCUT	19126	98.65%	56606	98.52%
			CI	14009	98.35%	38746	97.91%
			2obj	—	—	—	—
			2type	—	—	—	—
hsqldb	2733	6295	ZIPPER <sup>e</sup> -2obj	14001	98.29%	38702	97.79%
			ZIPPER <sup>e</sup> -2type	14001	98.29%	38706	97.80%
			CUT-SHORTCUT	14001	98.29%	38705	97.80%
			CI	2608	95.43%	5856	93.03%
			2obj	—	—	—	—
jedit	6028	13203	2type	2607	95.39%	5846	92.87%
			ZIPPER <sup>e</sup> -2obj	2607	95.39%	5846	92.88%
			ZIPPER <sup>e</sup> -2type	2607	95.39%	5850	92.93%
			CUT-SHORTCUT	2607	95.39%	5850	92.93%
			CI	5859	97.20%	12714	96.30%
gruntspud	14543	42099	2obj	5853	97.10%	12695	96.15%
			2type	5853	97.10%	12695	96.15%
			ZIPPER <sup>e</sup> -2obj	5853	97.10%	12695	96.15%
			ZIPPER <sup>e</sup> -2type	5853	97.10%	12698	96.18%
			CUT-SHORTCUT	5853	97.10%	12695	96.15%
soot	4372	16452	CI	4288	98.08%	16232	98.66%
			2obj	—	—	—	—
			2type	—	—	—	—
			ZIPPER <sup>e</sup> -2obj	4287	98.06%	16229	98.64%
			ZIPPER <sup>e</sup> -2type	4288	98.08%	16229	98.64%
columba	6757	14689	CUT-SHORTCUT	4287	98.06%	16229	98.64%
			CI	6674	98.77%	14369	97.82%
			2obj	—	—	—	—
			2type	—	—	—	—
			ZIPPER <sup>e</sup> -2obj	6674	98.77%	14369	97.82%
jython	4835	35970	ZIPPER <sup>e</sup> -2type	6674	98.77%	14369	97.82%
			CUT-SHORTCUT	6674	98.77%	14369	97.82%
			CI	4243	87.76%	11642	32.37%
			2obj	—	—	—	—
			2type	—	—	—	—
findbugs	5857	14075	ZIPPER <sup>e</sup> -2obj	4241	87.71%	11636	32.35%
			ZIPPER <sup>e</sup> -2type	4241	87.71%	11636	32.35%
			CUT-SHORTCUT	4241	87.71%	11629	32.33%
			CI	5808	99.16%	13901	98.76%
			2obj	5808	99.16%	13900	98.76%
			2type	5808	99.16%	13900	98.76%
			ZIPPER <sup>e</sup> -2obj	5808	99.16%	13901	98.76%
			ZIPPER <sup>e</sup> -2type	5808	99.16%	13901	98.76%
			CUT-SHORTCUT	5808	99.16%	13900	98.76%
			CI	5808	99.16%	13900	98.76%